



University of Colorado **Boulder**

Randomization methods for big-data (other than SGD)

Stephen Becker, CU Applied Math dept

Applied Math and Statistics colloquium, Colorado School of Mines

Fri Sept 19 2025

Joint work with several authors, mostly former PhD students:

(and colleagues at CU & the Colorado School of Mines)



Farhad Pourkamali-Anaraki
(Asst. Prof at UC Denver)



David Kozak
(private sector)



Osman Asif Malik
(was Alvarez postdoc at Berkeley Labs,
now private sector)

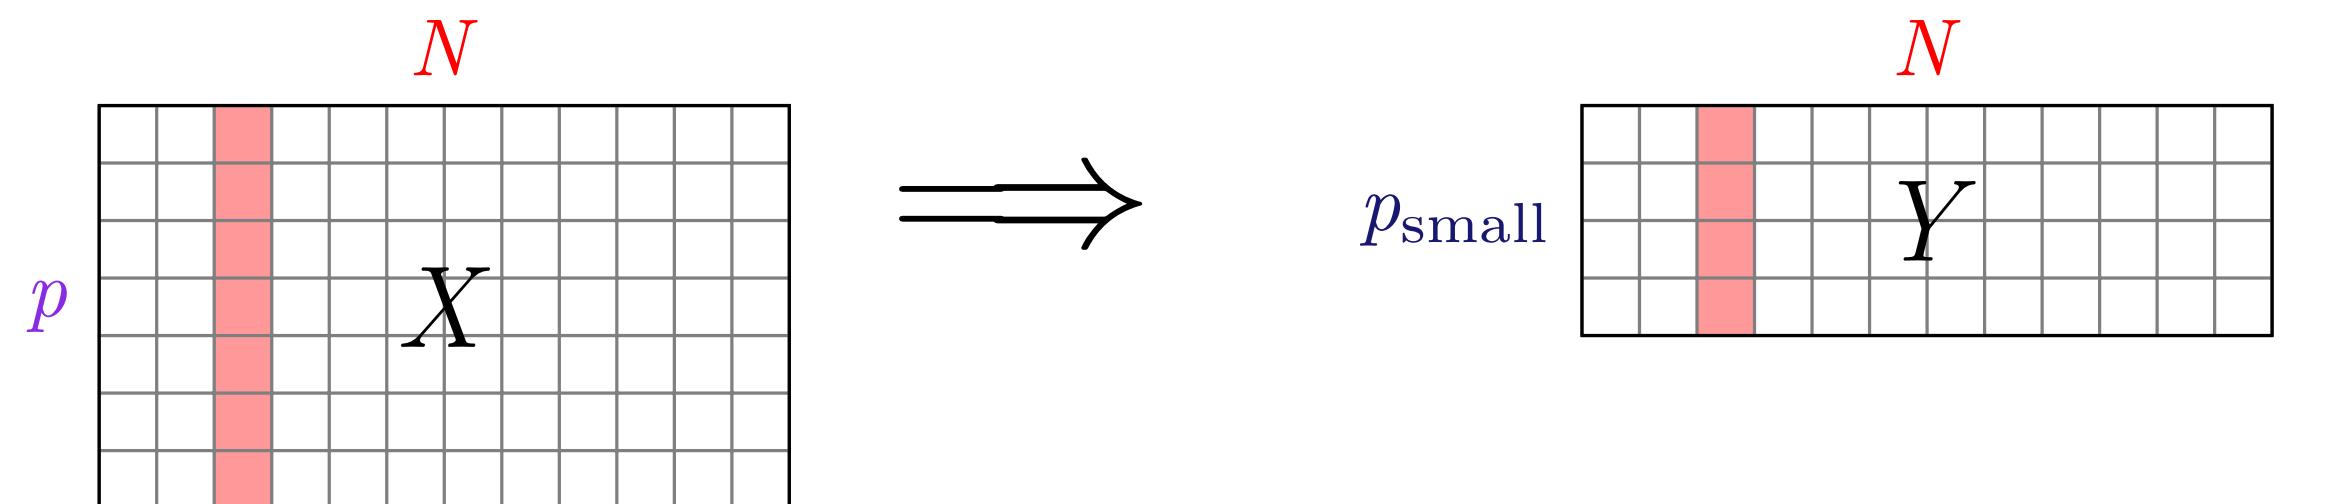
OUTLINE

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

- a. Warmup: PCA
- b. Classical sketches
- c. Structured sketches

- a. Warmup: linear algebra
- b. K-means clustering
- c. Tensor factorizations
- d. Gradient-free optimization

Reducing the dimensionality of data

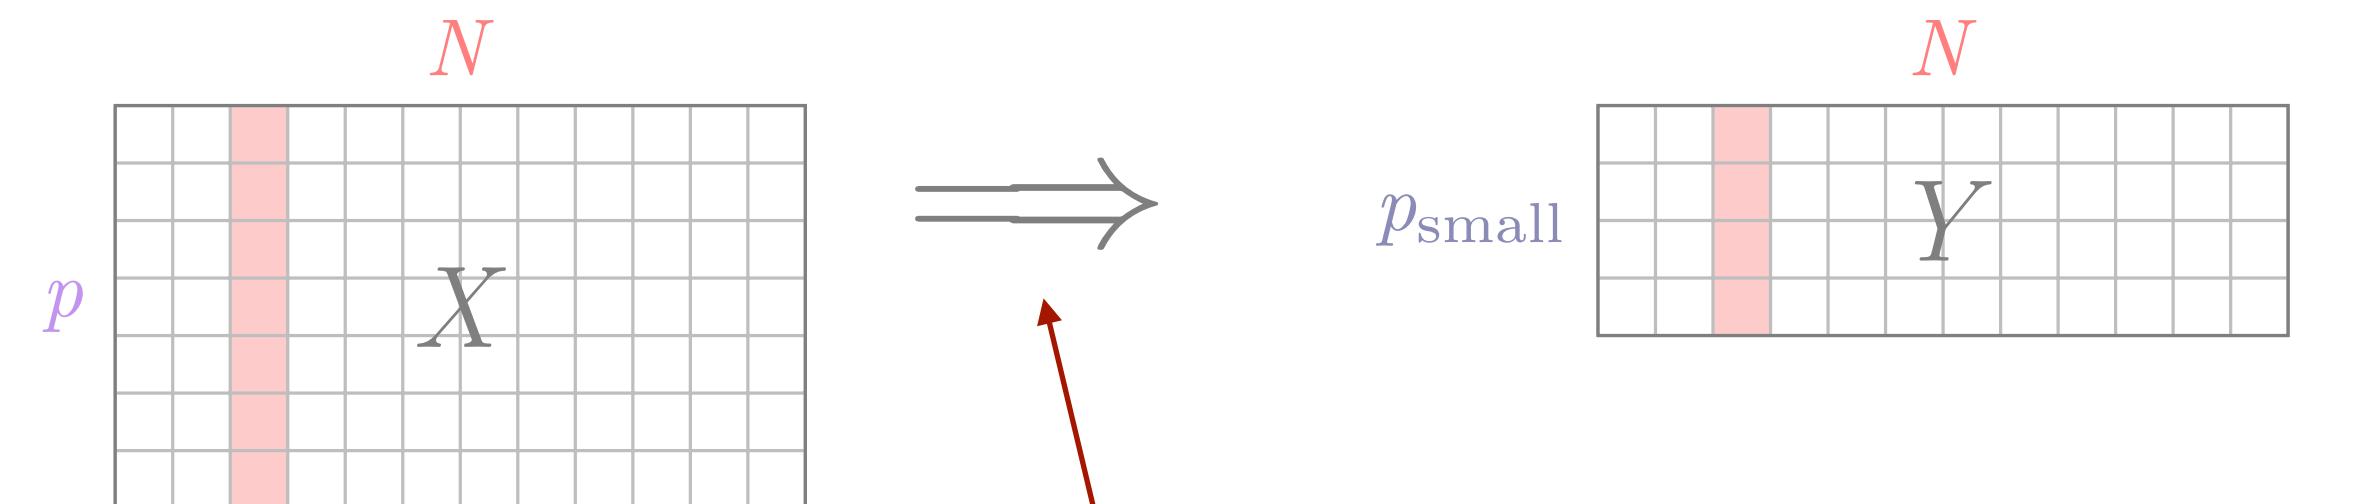


$x_i \in \mathbb{R}^p, \quad i = 1, \dots, N$ is converted to $y_i \in \mathbb{R}^{p_{\text{small}}}, \quad i = 1, \dots, N$

- a. Warmup: PCA
- b. Classical sketches
- c. Structured sketches

- a. Warmup: linear algebra
- b. K-means clustering
- c. Tensor factorizations
- d. Gradient-free optimization

Reducing the dimensionality of data



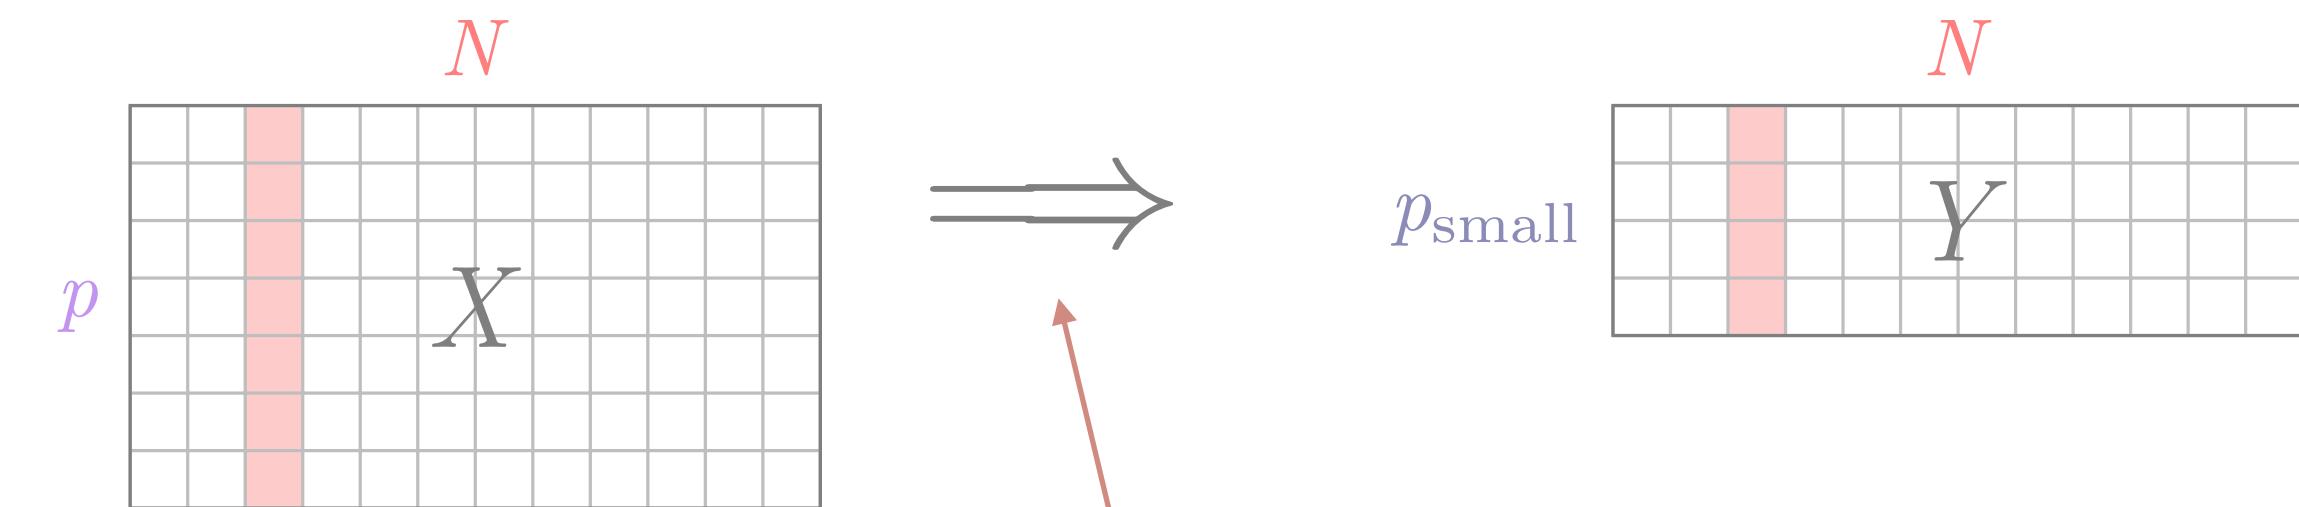
$x_i \in \mathbb{R}^p, \quad i = 1, \dots, N$ is converted to $y_i \in \mathbb{R}^{p_{\text{small}}}, \quad i = 1, \dots, N$

this reduction is often **linear**, in which case can write as

$$p_{\text{small}} \begin{matrix} N \\ Y \end{matrix} = p_{\text{small}} \begin{matrix} p \\ \Phi \end{matrix} \begin{matrix} N \\ X \end{matrix}$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Reducing the dimensionality of data



$x_i \in \mathbb{R}^p, i = 1, \dots, N$ is converted to $y_i \in \mathbb{R}^{p_{\text{small}}}, i = 1, \dots, N$

this reduction is often **linear**, in which case can write as

$$p_{\text{small}} \begin{matrix} N \\ Y \end{matrix} = p_{\text{small}} \begin{matrix} p \\ \Phi \end{matrix} \begin{matrix} N \\ X \end{matrix}$$

Why?

- ▶ Faster computation
 - ▶ Especially if the complexity of subsequent processing is $\mathcal{O}(p^2 N)$
 - ▶ If $p_{\text{small}} = 0.05p$ then speedup is $400\times$... or for storage reasons
- ▶ Denoising (remove irrelevant components) (especially to create **one-pass methods**)
- ▶ Fewer degrees-of-freedom reduces chance of overfitting models
- ▶ Visual interpretation (e.g., $p_{\text{small}} = \{2, 3\}$) ← This is “multidimensional scaling”. Specialized techniques (tSNE, UMAP) are best

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 1: PCA (non-linear)

Principal Component Analysis (aka Hotelling or Karhunen-Loeve transform)

Take SVD

$$X = [\mathbf{U}_1 \mathbf{U}_2] \begin{bmatrix} \Sigma_1 & 0 \\ 0 & \Sigma_2 \end{bmatrix} [\mathbf{V}_1^T \mathbf{V}_2^T]$$

where $\|\Sigma_2\| \approx$ small. Then PCA takes the form of

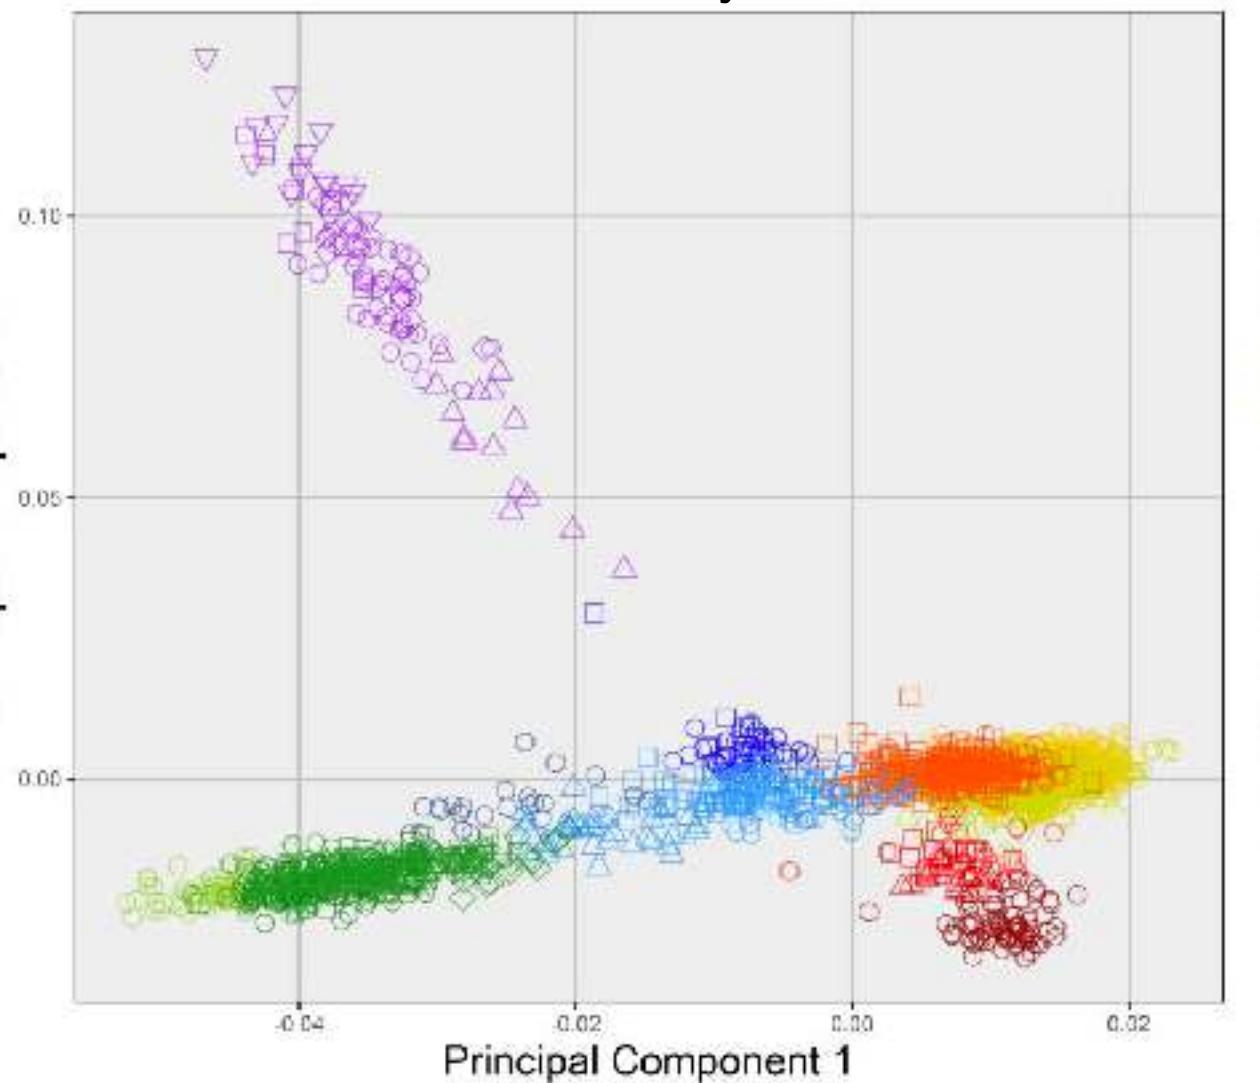
$$Y \leftarrow \Phi_{\text{PCA}} \cdot X$$

where $\Phi_{\text{PCA}} = \mathbf{U}_1^T$.

Note:

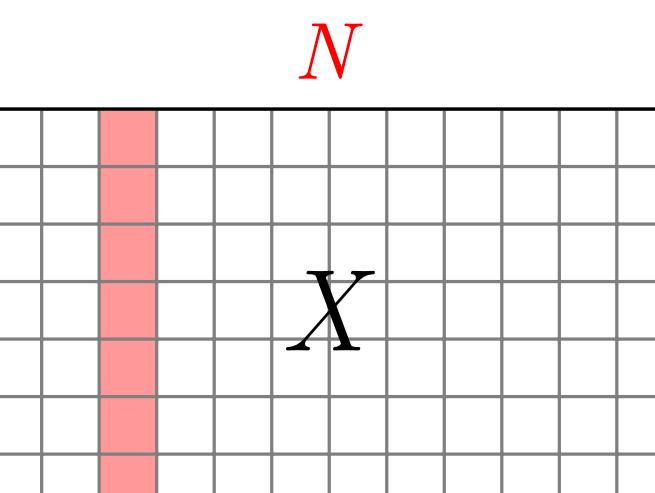
- ▶ Costly but linear in N : $(\mathcal{O}(p^2 N))$ direct, or about $\mathcal{O}(p_{\text{small}} p N)$ Krylov)
- ▶ Non-linear, since Φ_{PCA} is determined by X
- ▶ Need to process all of X before we can apply $y_i \leftarrow \Phi_{\text{PCA}} \cdot x_i$

Principal component analysis of the fineStructure co-ancestry matrix



https://commons.wikimedia.org/wiki/File:PCA_of_British_people.png

Gilbert, E., O'Reilly, S., Merrigan, M. et al. The Irish DNA Atlas: Revealing Fine-Scale Population Structure and History within Ireland. Sci Rep 7, 17199 (2017)



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 2: sub-sample dimensions (i.e., rows)

R randomly samples rows

$$\Phi_{\text{rows}} = p_{\text{small}} \begin{matrix} & p \\ & \boxed{R} \\ \boxed{\quad} & \end{matrix}$$

$$p_{\text{small}} \begin{matrix} N \\ Y \\ \boxed{\quad} \end{matrix} = p_{\text{small}} \begin{matrix} p \\ \boxed{\Phi} \\ \boxed{\quad} \end{matrix} \begin{matrix} N \\ X \\ \boxed{\quad} \end{matrix}$$

Method 2: sub-sample dimensions (i.e., rows)

R randomly samples rows

$$\Phi_{\text{rows}} = p_{\text{small}} \begin{matrix} & p \\ & | \\ \textcolor{blue}{\Phi}_{\text{rows}} = & p_{\text{small}} \end{matrix} \begin{matrix} & R \\ & | \\ & \textcolor{red}{R} \\ \hline & | \\ & \textcolor{red}{R} \end{matrix}$$

$$p_{\text{small}} \begin{matrix} & N \\ & | \\ & Y \\ \hline & | \\ & \textcolor{purple}{Y} \end{matrix} = p_{\text{small}} \begin{matrix} & p \\ & | \\ & \Phi \\ \hline & | \\ & \textcolor{red}{\Phi} \end{matrix} \begin{matrix} & N \\ & | \\ & X \\ \hline & | \\ & \textcolor{purple}{X} \end{matrix}$$

How to choose R ?

- ▶ Uniformly at random (... poor performance) “oblivious”
- ▶ According to row-norm (... better, not great, and slower) “non-oblivious”
- ▶ According to *leverage scores* based on partial SVD (... better, slow) (like PCA)

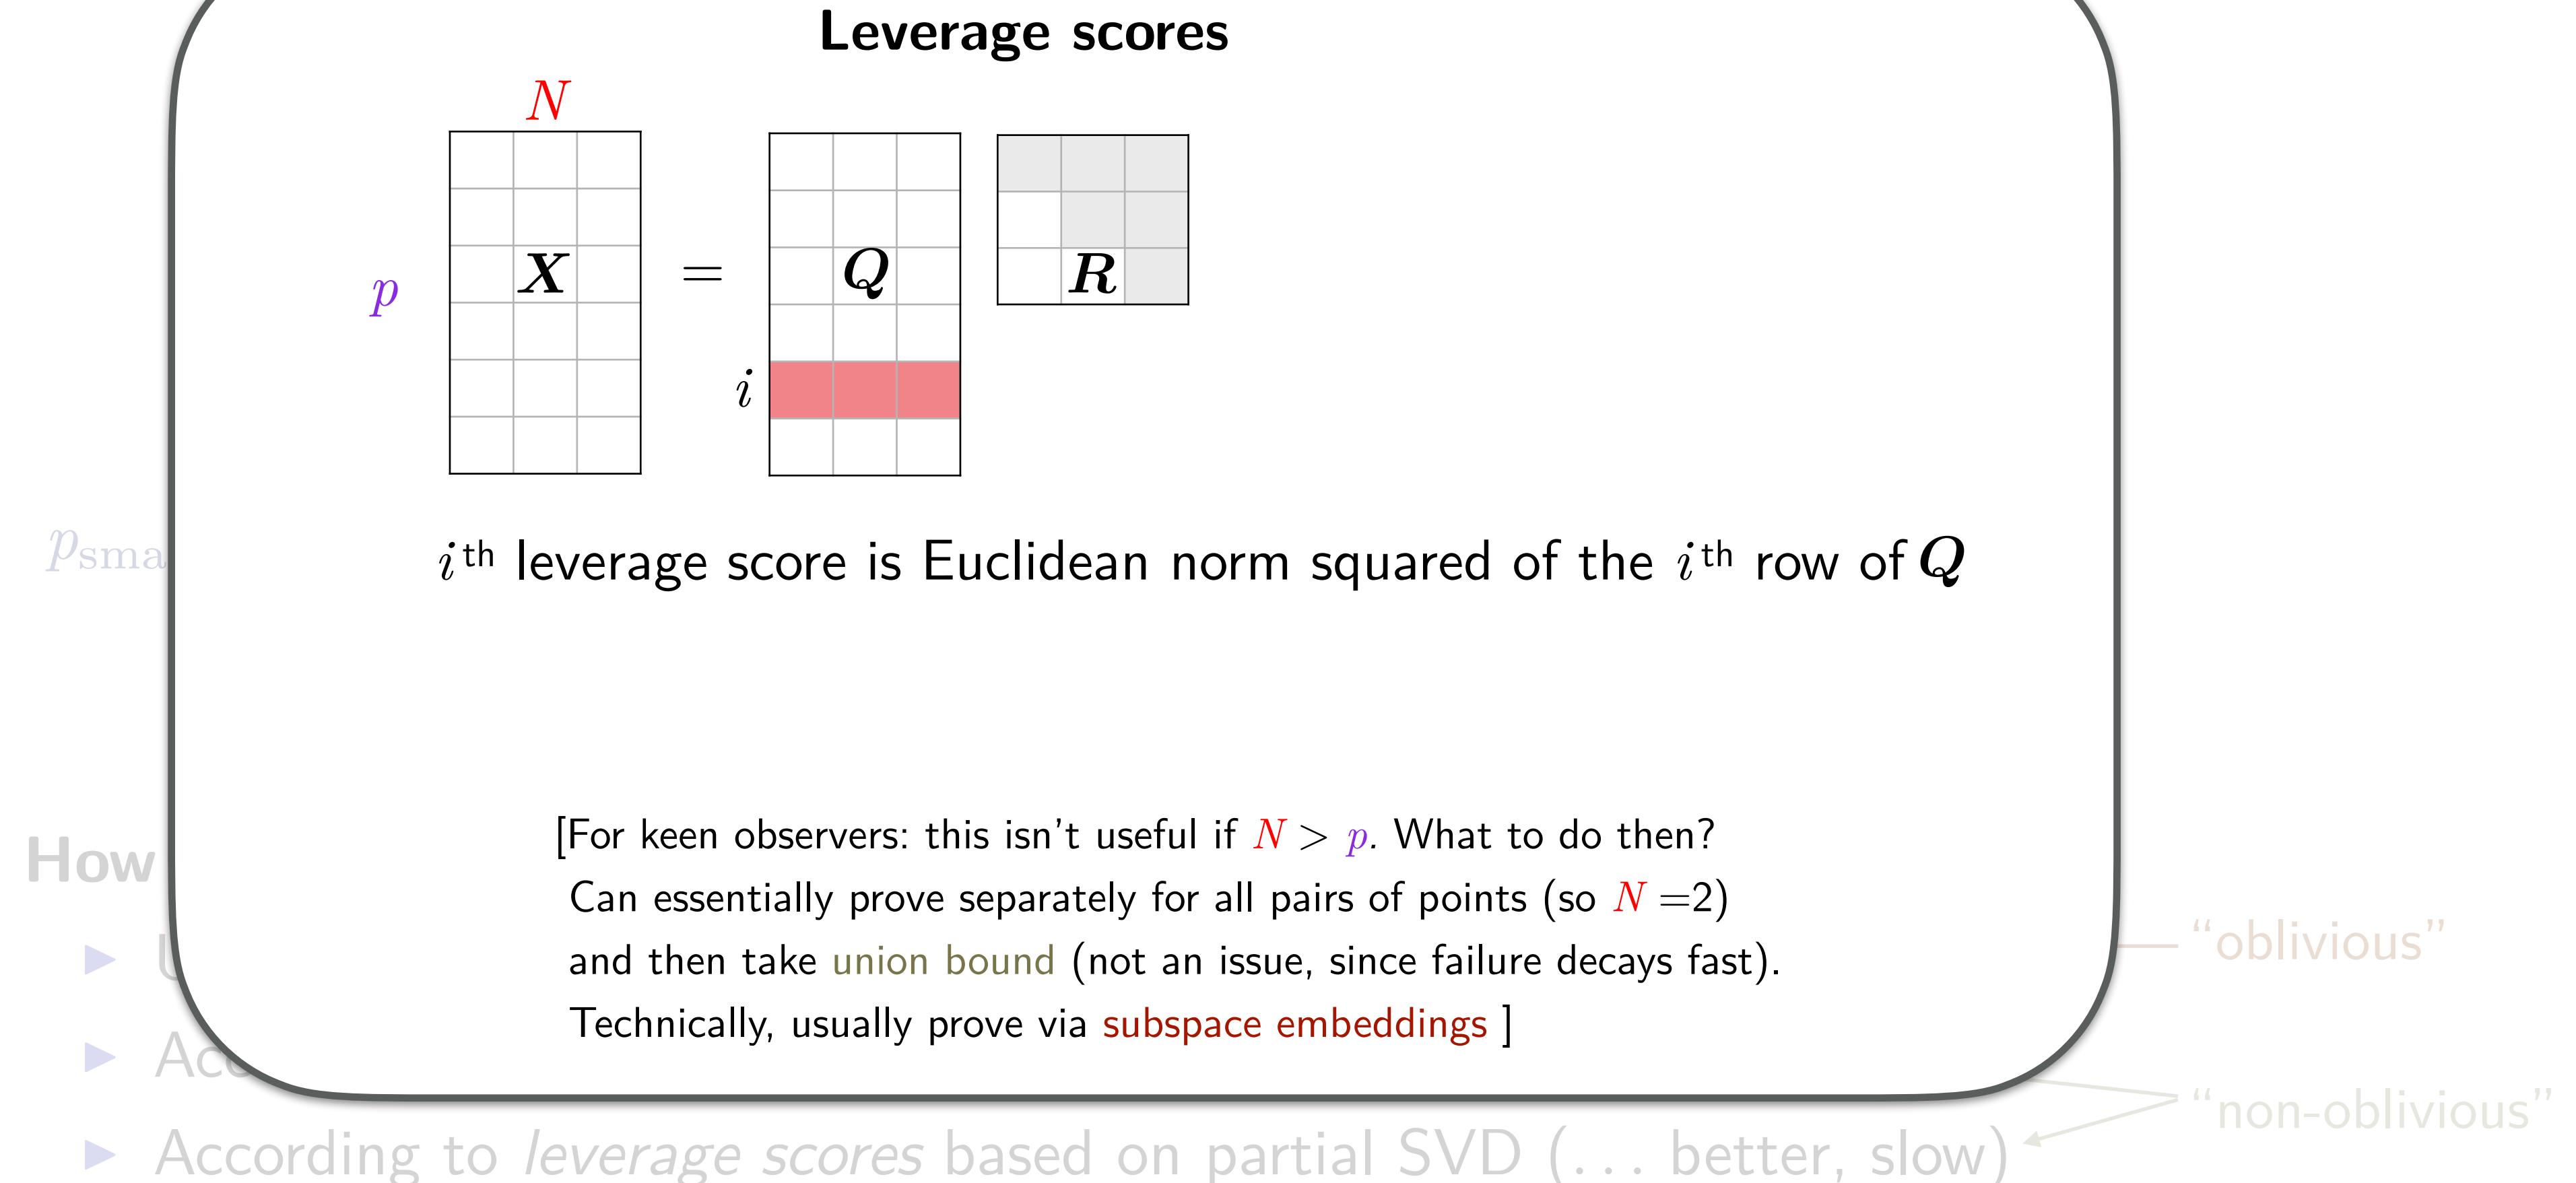
One benefit: for many applications, this avoids computing all entries of the RHS in the first place!

That's sometimes the bottleneck...

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 2: sub-sample dimensions (i.e., rows)

R randomly samples rows



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 3: dense random matrix

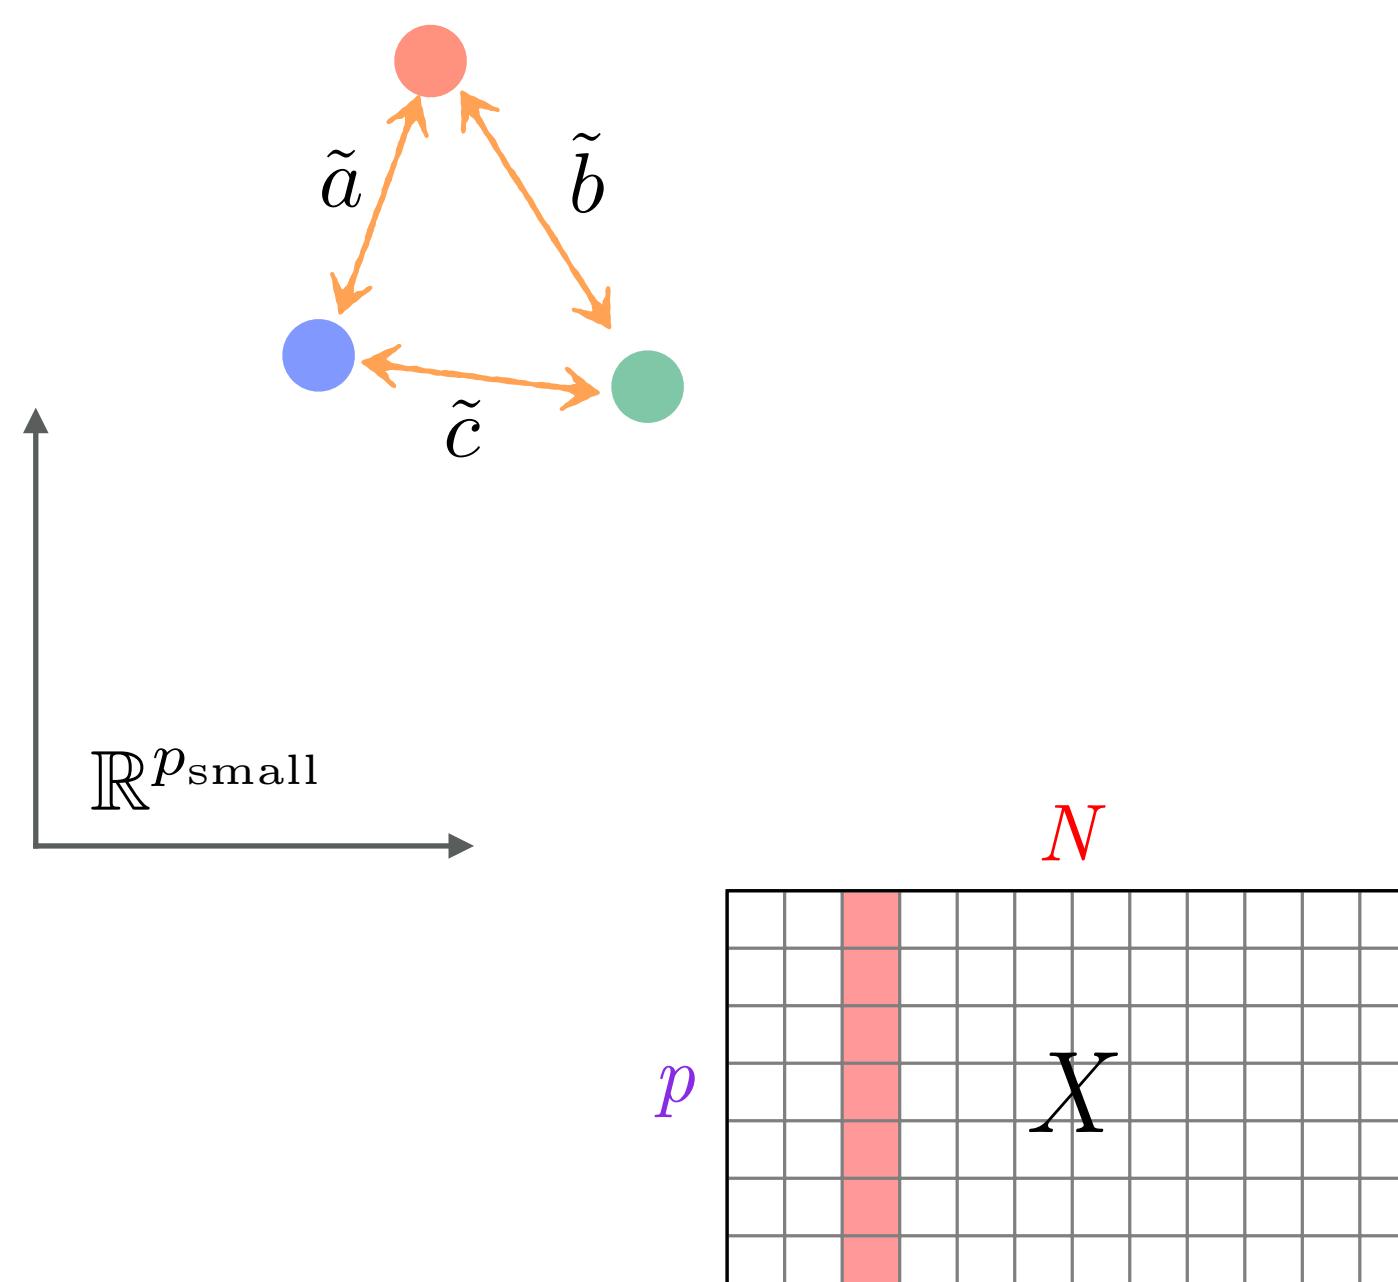
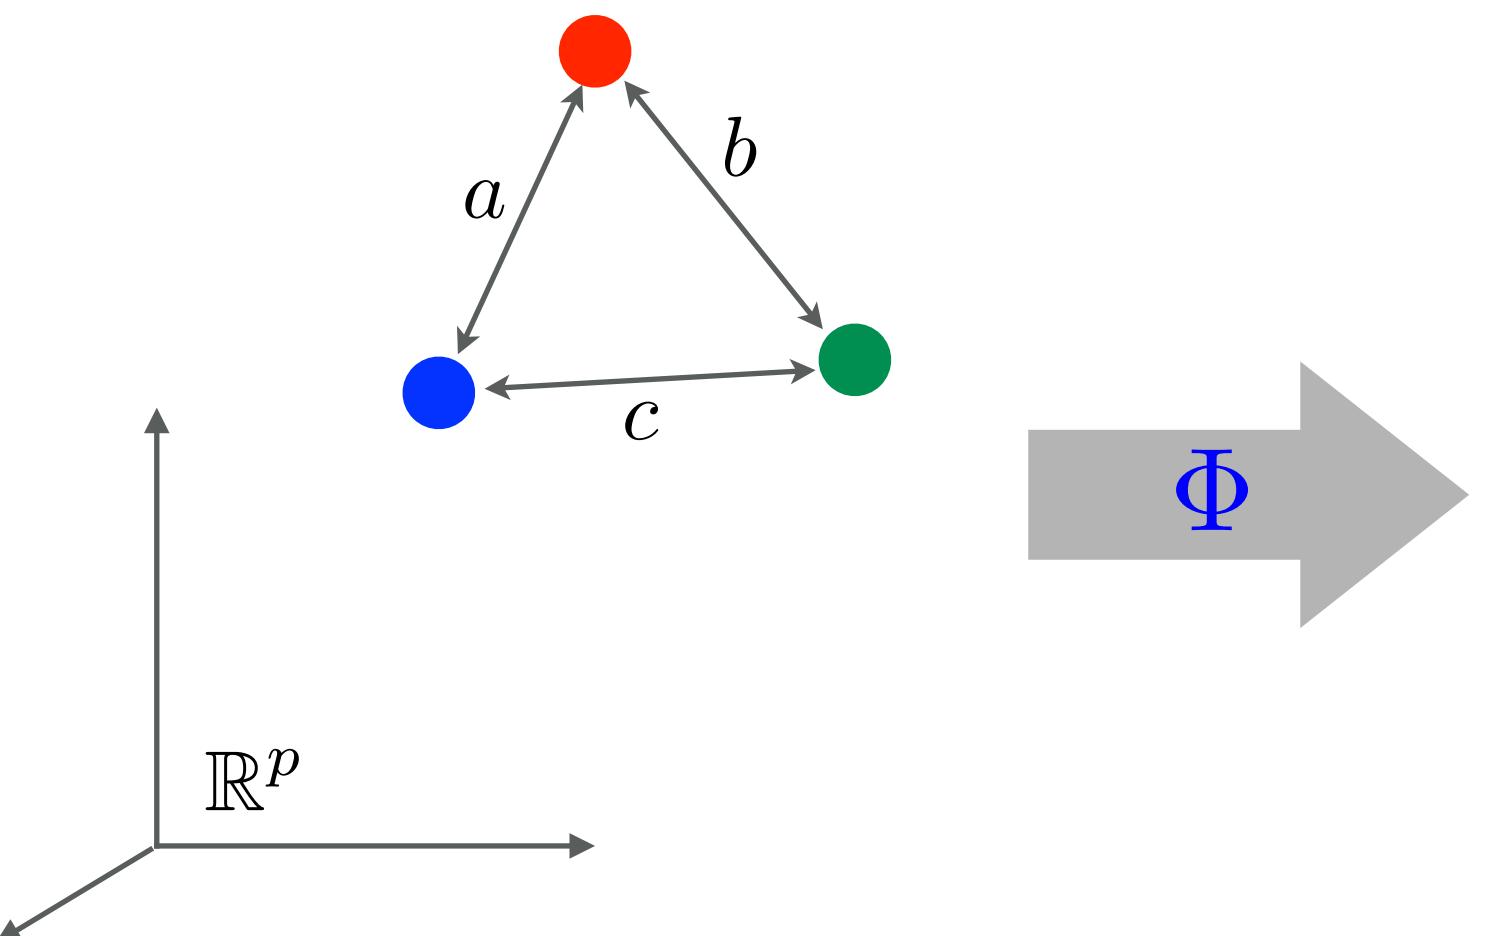
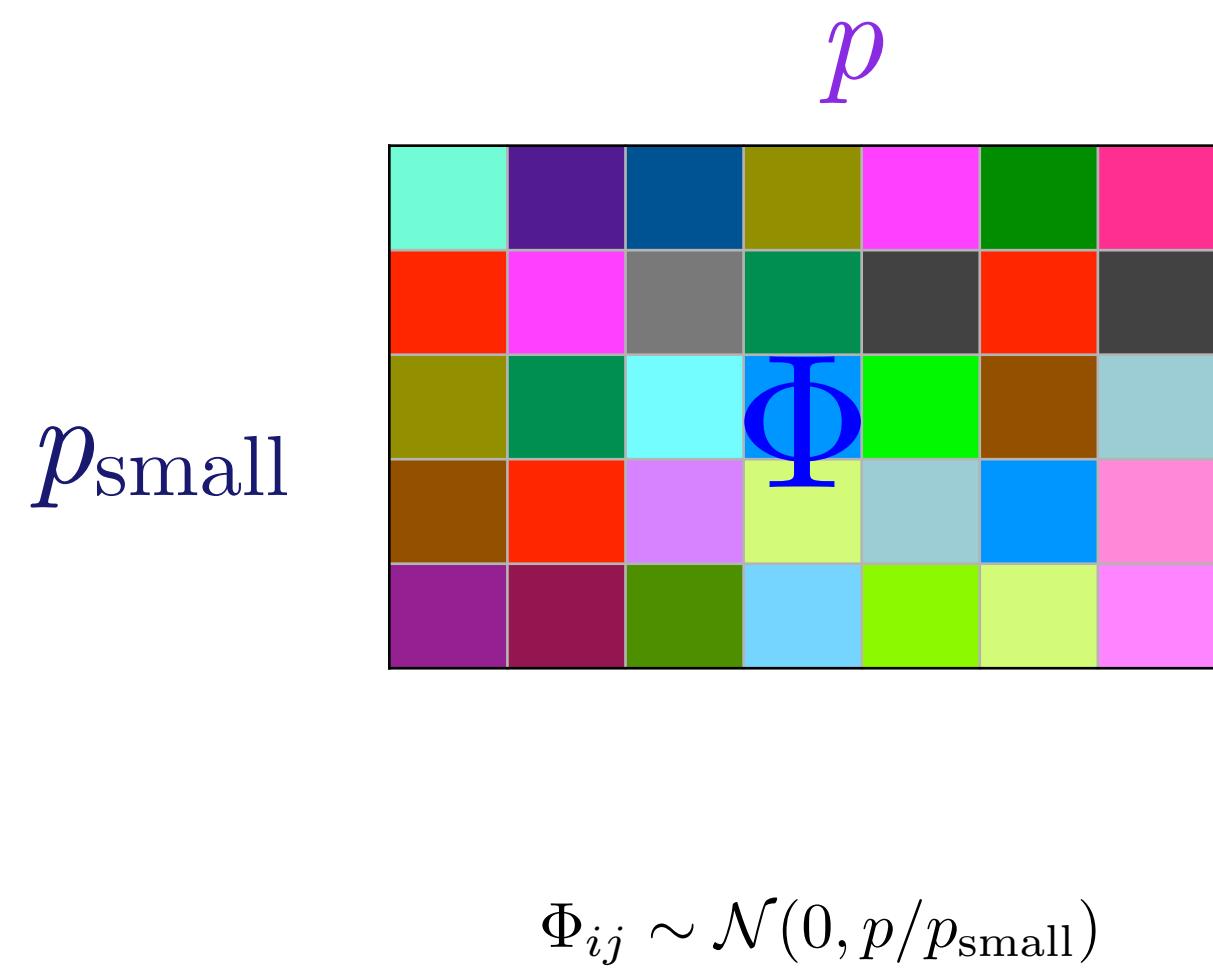
Theorem (Johnson-Lindenstrauss, 1984 (and Indyk-Motwani, 1998))

Choose $\Phi = \Phi_{randn}$ with $p_{small} \propto \varepsilon^{-2} \log N$ iid rows each $\mathcal{N}(0, p/p_{small})$, then for all $x_i, x_j \in \{x_1, \dots, x_N\} \subset \mathbb{R}^p$,

$$1 - \varepsilon \leq \frac{\|\Phi x_i - \Phi x_j\|_2}{\|x_i - x_j\|_2} \leq 1 + \varepsilon$$

with constant probability.

“gold standard” for accuracy among randomized data-oblivious sketches
(for any Euclidean norm results)



Method 3: dense random matrix

Theorem (Johnson-Lindenstrauss, 1984 (and Indyk-Motwani, 1998))

Choose $\Phi = \Phi_{randn}$ with $p_{small} \propto \varepsilon^{-2} \log N$ iid rows each $\mathcal{N}(0, p/p_{small})$, then for all $x_i, x_j \in \{x_1, \dots, x_N\} \subset \mathbb{R}^p$,

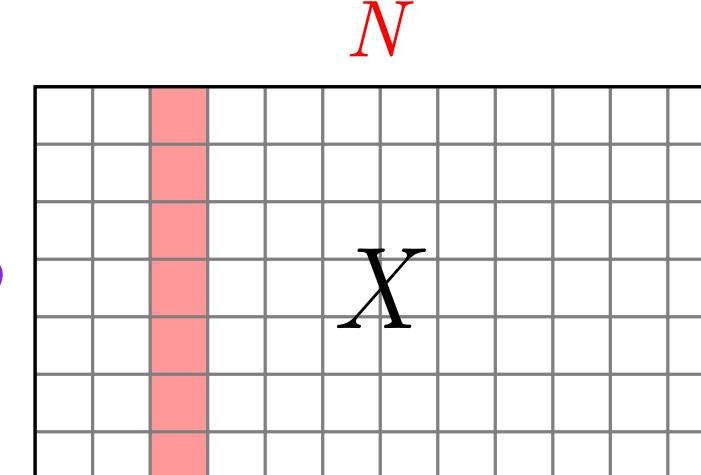
$$1 - \varepsilon \leq \frac{\|\Phi x_i - \Phi x_j\|_2}{\|x_i - x_j\|_2} \leq 1 + \varepsilon$$

with constant probability.

- ▶ Φ is linear and data-independent (unlike PCA)
- ▶ but Φ is not orthogonal (unlike PCA) (though variants can be made to be orthogonal)
- ▶ Independent of original dimension p
- ▶ Independent of conditioning of X (unlike PCA)
- ▶ Probabilistic (unlike PCA)

Usually not practical since too costly since computing ΦX is $\mathcal{O}(pp_{small}N)$

Depends on what its purpose is...



Method 3: dense random matrix

Theorem (Johnson-Lindenstrauss, 1984 (and Indyk-Motwani, 1998))

Choose $\Phi = \Phi_{randn}$ with $p_{small} \propto \varepsilon^{-2} \log N$ iid rows each $\mathcal{N}(0, p/p_{small})$, then for all $x_i, x_j \in \{x_1, \dots, x_N\} \subset \mathbb{R}^p$,

$$1 - \varepsilon \leq \frac{\|\Phi x_i - \Phi x_j\|_2}{\|x_i - x_j\|_2} \leq 1 + \varepsilon$$

with constant probability.

Bipartite Expander
graphs

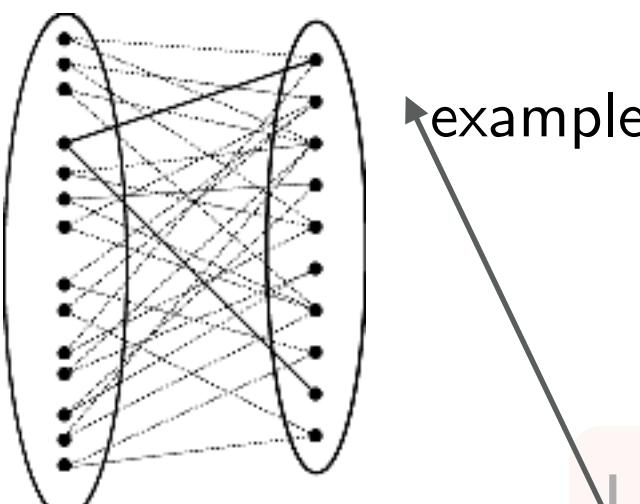


image: De Castro,
IEEE TIT 2010

- ▶ Φ is linear and data-independent (unlike PCA)
- ▶ but Φ is not orthogonal (unlike PCA)
- ▶ Independent of original dimension p
- ▶ Independent of conditioning of X (unlike PCA)
- ▶ Probabilistic (unlike PCA)

Usually not practical since too costly since computing ΦX is $\mathcal{O}(pp_{small}N)$

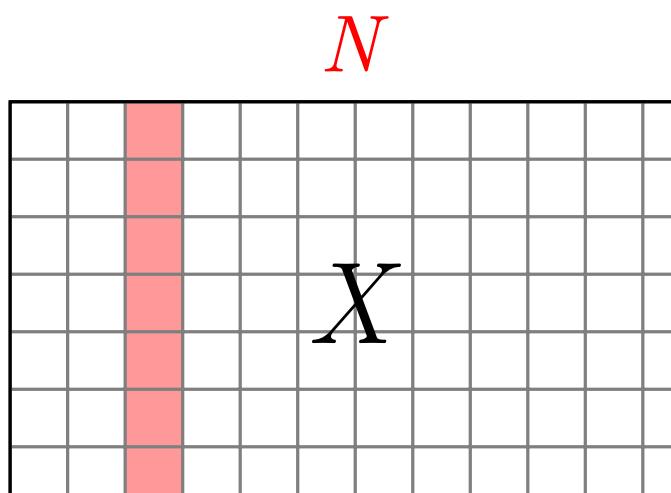
CountSketch

example

Many **extensions**, for example:

- sub-Gaussian entries, and/or dependent rows (e.g., columns uniform on sphere)
- dependent columns, e.g., Haar measure on orthogonal matrices
- **sparse matrix** (worse performance though)

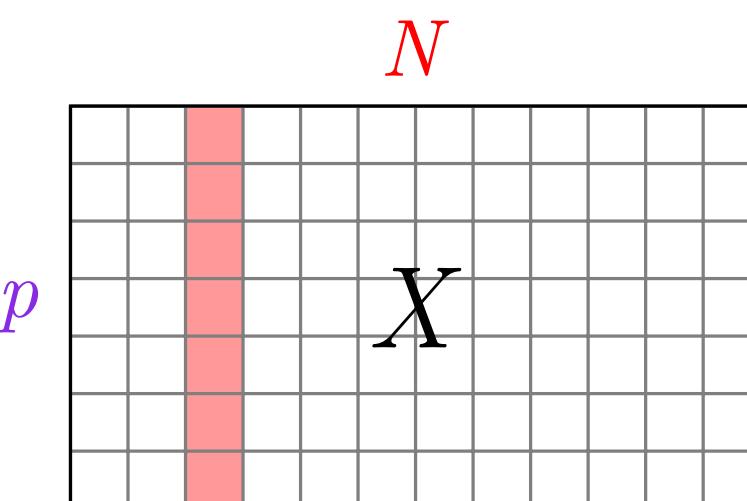
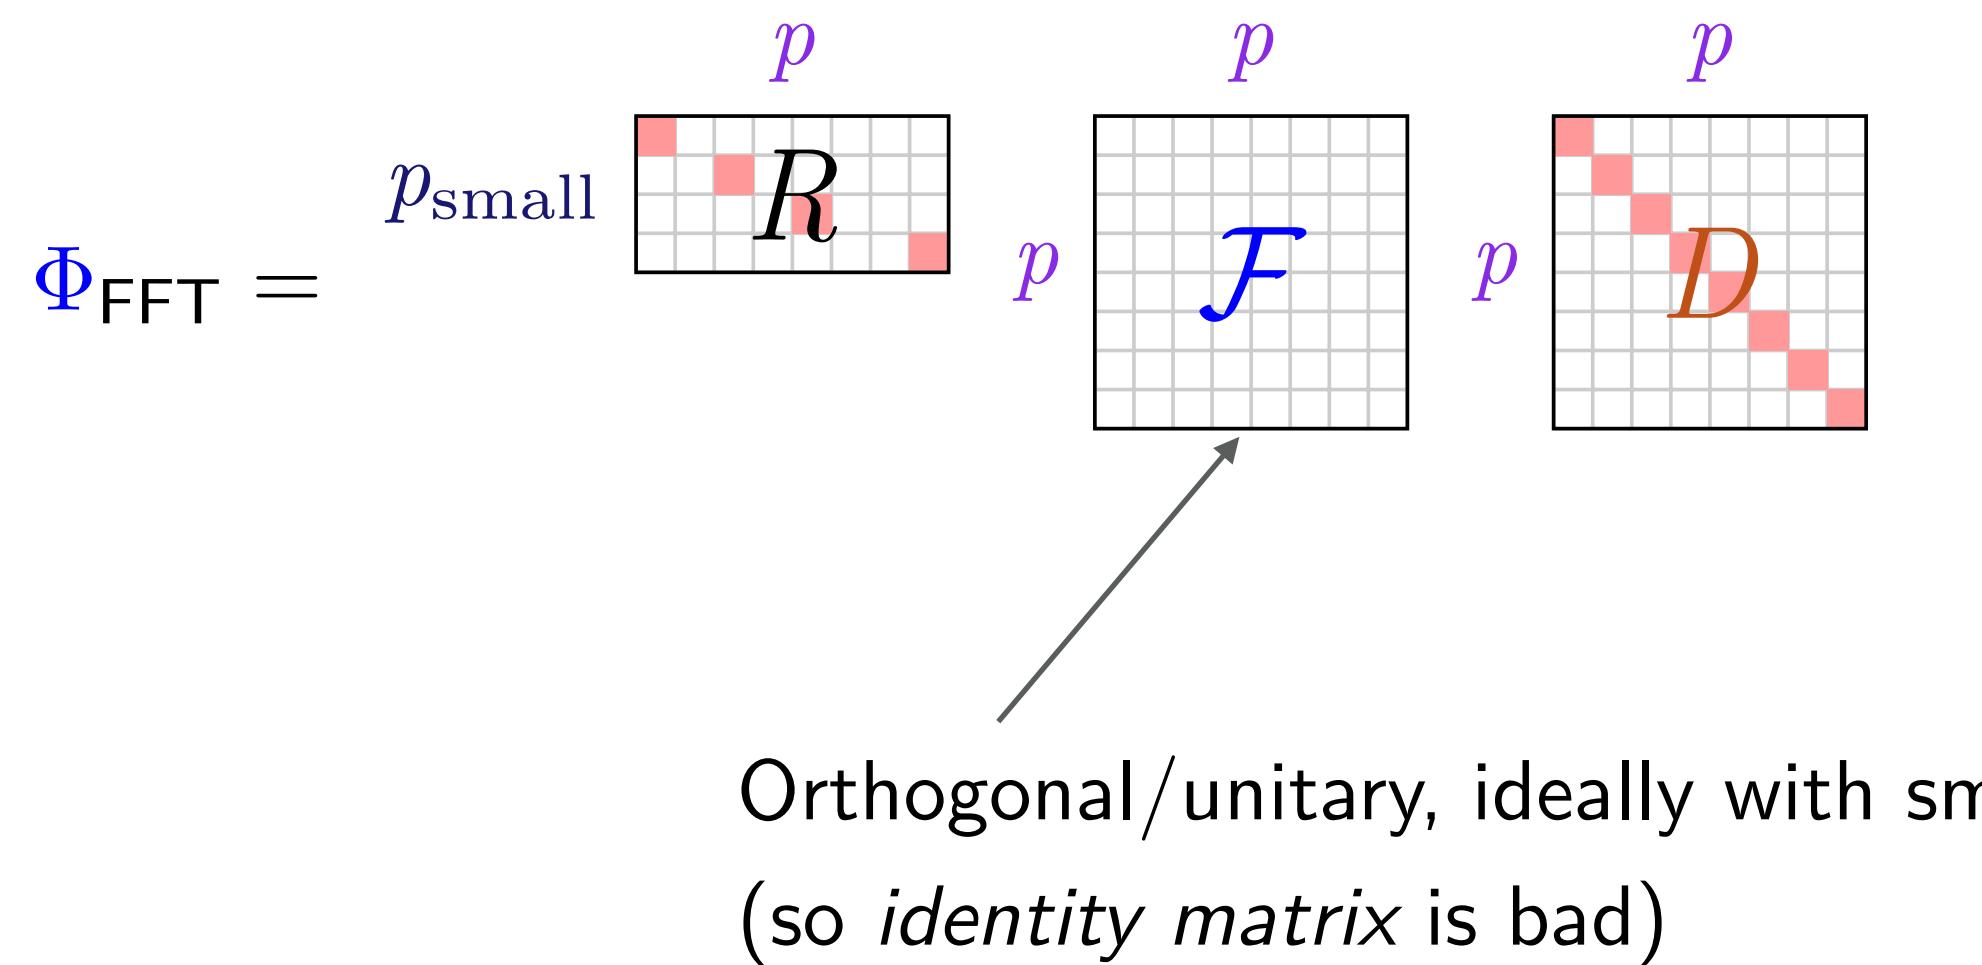
Extensions of **applicability**, for example:
 ○ entire **subspaces** and **manifolds**
 (using covering-number arguments)



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 4: Fast Johnson-Lindenstrauss Transforms

R samples rows; \mathcal{F} Fourier-like; D diag. w/ random ± 1 entries (“Rademacher”) (uniformly)



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 4: Fast Johnson-Lindenstrauss Transforms

R samples rows; \mathcal{F} Fourier-like; D diag. w/ random ± 1 entries (“Rademacher”) (uniformly)

$$\Phi_{\text{FFT}} = p_{\text{small}} \begin{array}{c} p \\ R \\ \hline \end{array} \quad \begin{array}{c} p \\ \mathcal{F} \\ \hline \end{array} \quad \begin{array}{c} p \\ D \\ \hline \end{array}$$

Cost:

- ▶ $\mathcal{F}(x)$ costs $p \log p$ (vs p^2 naively)
- ▶ hence $\mathcal{O}(pN \log p)$ to compute ΦX

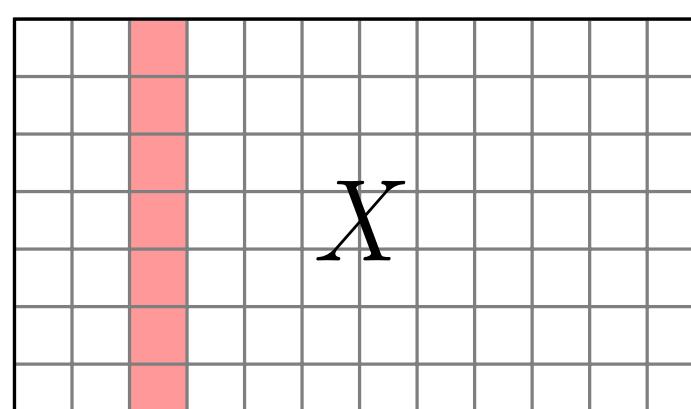
Guarantees: almost as good as classical Johnson-Lindenstrauss
First guarantees in the **Fast Johnson-Lindenstrauss** paper

Ailon and Chazelle, “Approximate nearest neighbors and the fast Johnson-Lindenstrauss transform”, STOC 2006

Think of \mathcal{F} as FFT or DCT or Hadamard transform

also known under many names, e.g., Random Orthogonal System (**ROS**), FJLT

N



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 4: Fast Johnson-Lindenstrauss Transforms

R samples rows; \mathcal{F} Fourier-like; D diag. w/ random ± 1 entries ("Rademacher")

$$\Phi_{\text{FFT}} = p_{\text{small}} \begin{matrix} & p & p & p \\ & R & \mathcal{F} & D \end{matrix}$$

How to get guarantees?

Cost:

- ▶ $\mathcal{F}(x)$ costs $p \log p$
- ▶ hence $\mathcal{O}(pN \log p)$

Basic idea is that after applying the first two steps, resulting matrix has (almost) uniform leverage scores...
... so uniform subsampling is (almost) leverage score sampling.

Guarantees: almost as good as PCA

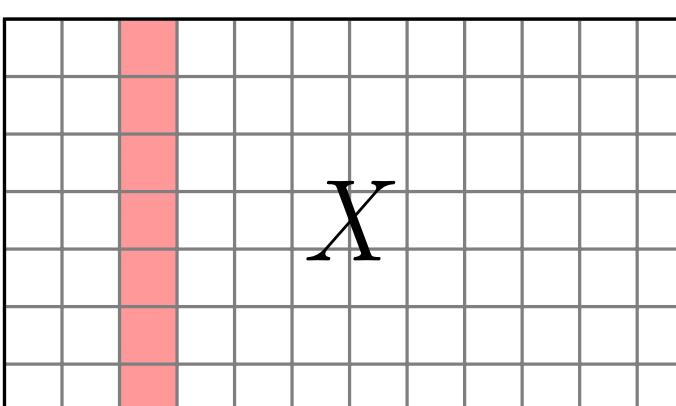
First guarantees in the [Fast Johnson-Lindenstrauss paper](#)

Ailon and Chazelle, "Approximate nearest neighbors and the fast Johnson-Lindenstrauss transform", STOC 2006

Think of \mathcal{F} as FFT or DCT or Hadamard transform

also known under many names, e.g., Random Orthogonal System, FJLT

N



Method 4: Fast Johnson-Lindenstrauss Transforms

R samples rows; \mathcal{F} Fourier-like; \mathcal{D} diag. w/ random ± 1 entries ("Rademacher")

More intuition

Suppose we want to estimate $\|\mathbf{x}\|_2^2$ for $\mathbf{x} \in \mathbb{R}^p$ by sampling just p_{small} entries (and multiplying by p/p_{small}). Wlog, let $\|\mathbf{x}\|_2 = 1$.

- ▶ High variance if \mathbf{x} has a few large components
- ▶ **Intuition:** the matrix $\mathcal{F}\mathcal{D}$ flattens out input vectors
 - ▶ Components of \mathbf{x} can be as large as 1
 - ▶ **Average component** of \mathbf{x} is $p^{-1/2}$
- ▶ Define $\mathbf{y} = \mathcal{F}\mathcal{D}\mathbf{x}$, so $\|\mathbf{y}\|_2 = \|\mathbf{x}\|_2$. Then $y_1 = \sum_{j=1}^p \mathcal{F}_{1,j} \varepsilon_j x_j$
 - ▶ $\mathbb{E} y_1 = 0$ since $\mathbb{E} \varepsilon_j = 0$
 - ▶ $\text{Var}(y_1) = p^{-1}$ since $\mathcal{F}_{1,j}^2 \leq \eta = p^{-1}$ and $\|\mathbf{x}\|_2 = 1$ and ε_j independent
- ▶ Hoeffding inequality:

$$\mathbb{P}(|y_1| \geq t) \leq 2e^{-pt^2/2}$$

- ▶ So chance that one of y_j for $j = 1, \dots, p$ is $\geq t$ is less than $p2e^{-pt^2/2}$
- ▶ So with constant probability, can choose $t \approx \sqrt{p^{-1} \log(p)}$
- ▶ Conclusion: the **max** component of \mathbf{y} is about the same as its **average** component. Subsampling \mathbf{y} will have much less variance than subsampling \mathbf{x} .

from "Improved analysis of the subsampled randomized Hadamard transform" (Tropp, 2010); see also Ailon and Chazelle

also known as

N

X

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 5: subsample entries of a matrix

Keep $x'_i \in \mathbb{R}^p$ but with only p_{small} nonzeros

Idea goes back to landmark Achlioptas/McSherry paper (2001)

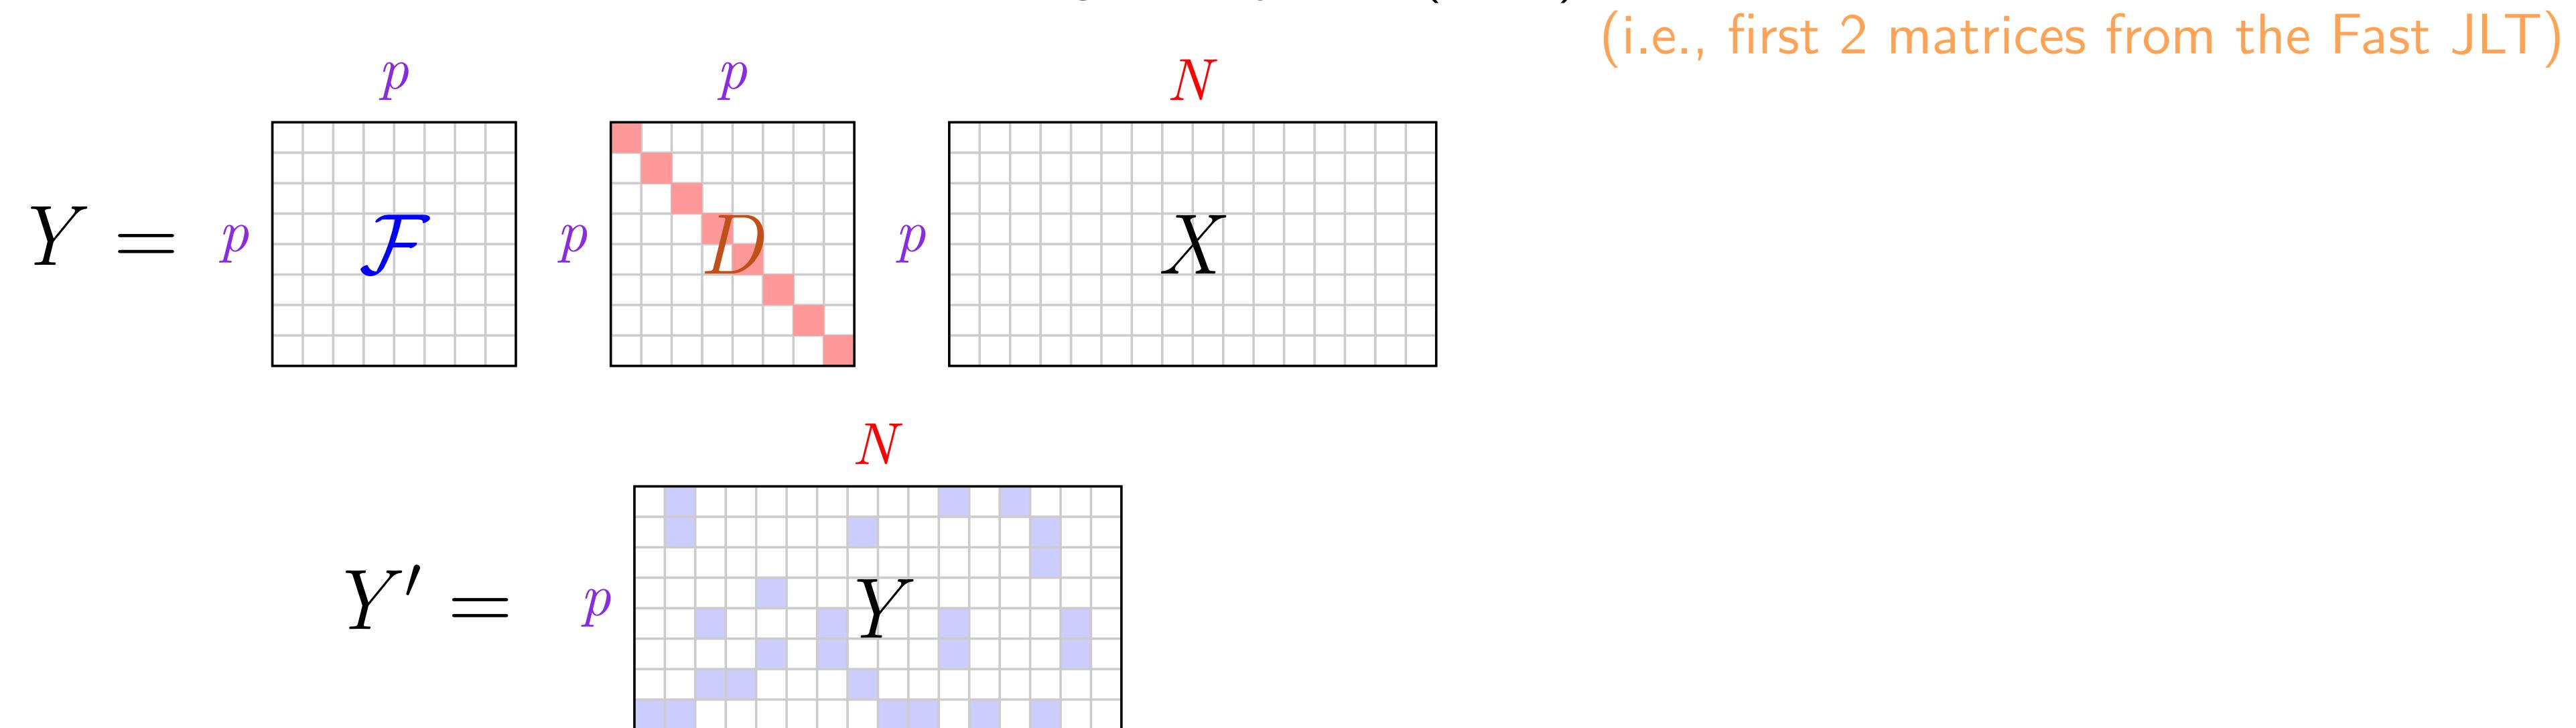
$$X' = p \begin{matrix} & N \\ & | \\ X & \end{matrix}$$

This is the **odd-one-out** so far, because it's not the same linear operator applied to every column

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 6: “precondition,” then subsample entries of a matrix

Our new twist: first **precondition** with Random Orthogonal System (ROS)



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Why precondition?

Easy to see that better than uniform sampling of X is **non-uniform** sampling, with probability proportional to magnitude of entry

cf. Achlioptas, Z. Karnin, and E. Liberty 2013

Disadvantage of weighted sampling is the extra pass through data

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Why precondition?

Easy to see that better than uniform sampling of X is **non-uniform** sampling, with probability proportional to magnitude of entry

cf. Achlioptas, Z. Karnin, and E. Liberty 2013

Disadvantage of weighted sampling is the extra pass through data

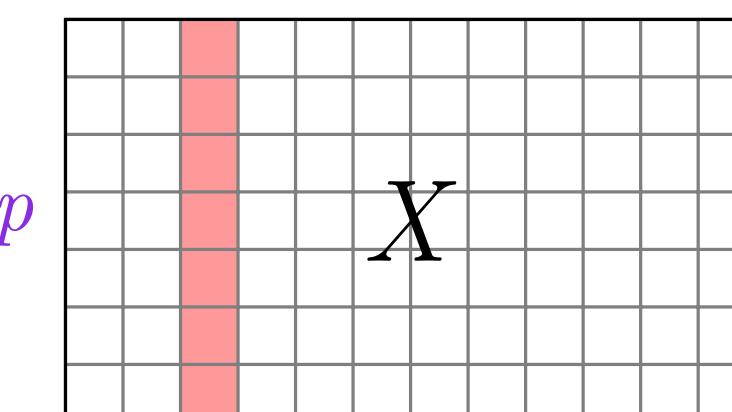
In our analysis, we have terms like (want this small)

$$\|X\|_{\text{max-entry}} = \max_{\substack{i=1, \dots, N \\ j=1, \dots, p}} |X_{ij}|$$

Assuming $\|x_i\|_2 = 1$,

- ▶ it is possible for $\|X\|_{\text{max-entry}} = 1$ (BAD)
- ▶ best case is $\|X\|_{\text{max-entry}} = 1/\sqrt{p}$

N



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Why precondition?

Easy to see that better than uniform sampling of X is **non-uniform** sampling, with probability proportional to magnitude of entry

cf. Achlioptas, Z. Karnin, and E. Liberty 2013

Disadvantage of weighted sampling is the extra pass through data

In our analysis, we have terms like (want this small)

$$\|X\|_{\text{max-entry}} = \max_{\substack{i=1, \dots, N \\ j=1, \dots, p}} |X_{ij}|$$

Assuming $\|x_i\|_2 = 1$,

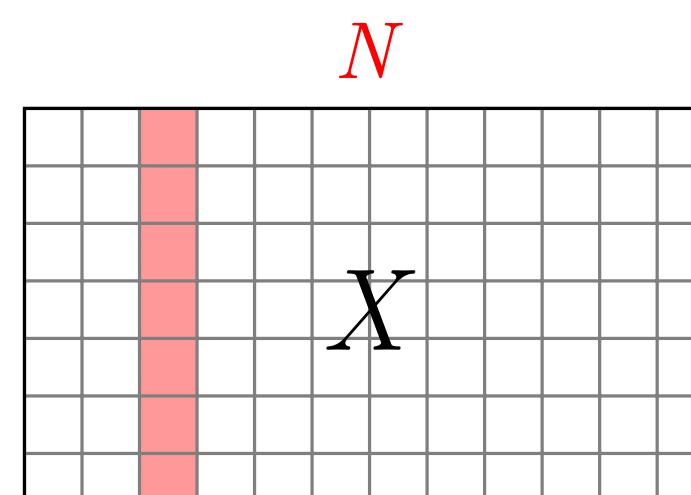
- ▶ it is possible for $\|X\|_{\text{max-entry}} = 1$ (BAD)
- ▶ best case is $\|X\|_{\text{max-entry}} = 1/\sqrt{p}$

Benefit of preconditioning

After applying ROS, exponentially small chance that

$$\|Y\|_{\text{max-entry}} > \sqrt{\log(Np)}/\sqrt{p}$$

N.B. Since $\mathcal{F}D$ is unitary, bounds in spectral/Frobenius norm are unchanged



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Theory: estimating the mean

Given a true mean $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$, and our estimate of it \hat{x} from sampled data,



Theorem (Pourkamali-Anaraki & B., *IEEE Trans. Info Theory* 2017)

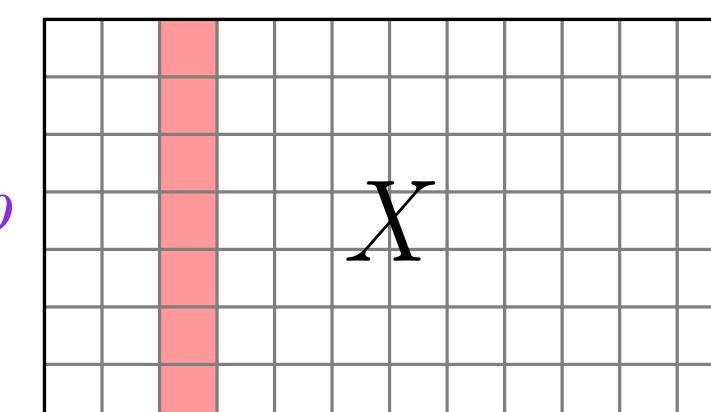
$\mathbb{E} \hat{x} = \bar{x}$ and $\|\hat{x} - \bar{x}\|_\infty \leq t$ with probability greater than

$$1 - 2p \exp\left(\frac{-N\gamma t^2/2}{\|X\|_{max-row}^2 + t/3\|X\|_{max-entry}}\right)$$

where $\gamma = p_{small}/p$

(simplifying to $p_{small} \ll p \ll N$)

N



p

X

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Theory: estimating the mean

Given a true mean $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$, and our estimate of it \hat{x} from sampled data,

Theorem (Pourkamali-Anaraki & B., *IEEE Trans. Info Theory* 2017)

$\mathbb{E} \hat{x} = \bar{x}$ and $\|\hat{x} - \bar{x}\|_\infty \leq t$ with probability greater than

$$1 - 2p \exp\left(\frac{-N\gamma t^2/2}{\|X\|_{\max\text{-row}}^2 + t/3\|X\|_{\max\text{-entry}}}\right)$$

where $\gamma = p_{small}/p$

(simplifying to $p_{small} \ll p \ll N$)

If X has normalized columns, then $\|X\|_{\max\text{-entry}} = 1$ and $\|X\|_{\max\text{-row}} = \sqrt{N}$ are possible, which is bad.

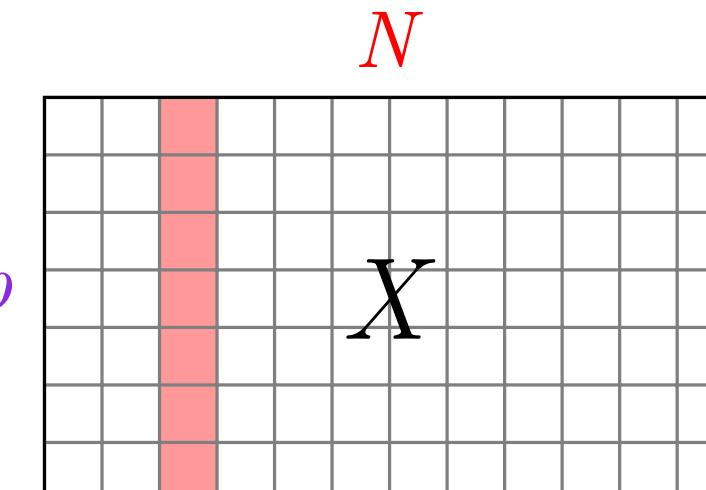
Lemma

If X is preconditioned, then

(for $Np = 10^{10}$, $\sqrt{\log(2Np) + 1000} = 32$)

$$\mathbb{P} \left\{ \|X\|_{\max\text{-entry}} \geq \frac{\sqrt{2}}{\sqrt{p}} \cdot \sqrt{\log(2Np) + 1000} \right\} \leq .001$$

$$\mathbb{P} \left\{ \|X\|_{\max\text{-row}} \geq \frac{\sqrt{2N}}{\sqrt{p}} \cdot \sqrt{\log(2Np) + 1000} \right\} \leq .001$$



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Theory: estimating the mean

Given a true mean μ

Theorem:

$$\mathbb{E} \hat{x} = \mu$$

where

If X has possible values

Lemmas:

If X is

$1, \dots, N$

Probability theory aside

(symmetric Bernoulli, aka Rademacher)
Let $y_i = \pm 1$ be iid Bernoulli, $\hat{x} = \frac{1}{N} \sum_{i=1}^N y_i$ (transform of Binomial)

- ▶ $\mathbb{E}(\hat{x}) = 0$
- ▶ $\text{Var}(\hat{x}) = 1/N$

Chebyshev, Markov inequalities

$$\mathbb{P}(|\hat{x}| \geq t) \leq \frac{1}{Nt^2}$$

Example: $N = 10^4$ and $t = 0.1$, $\mathbb{P}(|\hat{x}| \geq t) \leq 0.01$

Markov

$$X \geq 0, \mathbb{P}(X \geq t) \leq \frac{\mu}{t}$$

Chebyshev

$$\mathbb{P}(|X - \mu| \geq t\sigma^2) \leq \frac{1}{t^2}$$

Concentration ineq.

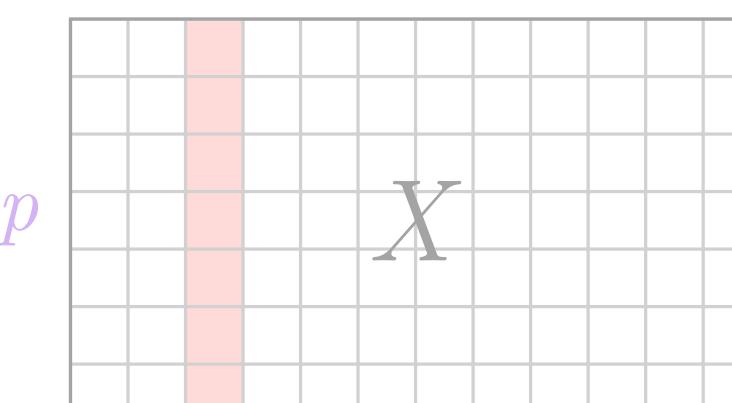
Bernstein, Hoeffding, Chernoff

$$\mathbb{P}(|\hat{x}| \geq t) \leq 2\exp\left(\frac{-Nt^2/2}{1 + t/3}\right)$$

Example: $N = 10^4$ and $t = 0.1$, $\mathbb{P}(|\hat{x}| \geq t) \leq 2 \cdot 10^{-21}$

intuition: central limit theorem says random sums of iid rv should look Gaussian, and for a Gaussian, we have exponential concentration

N



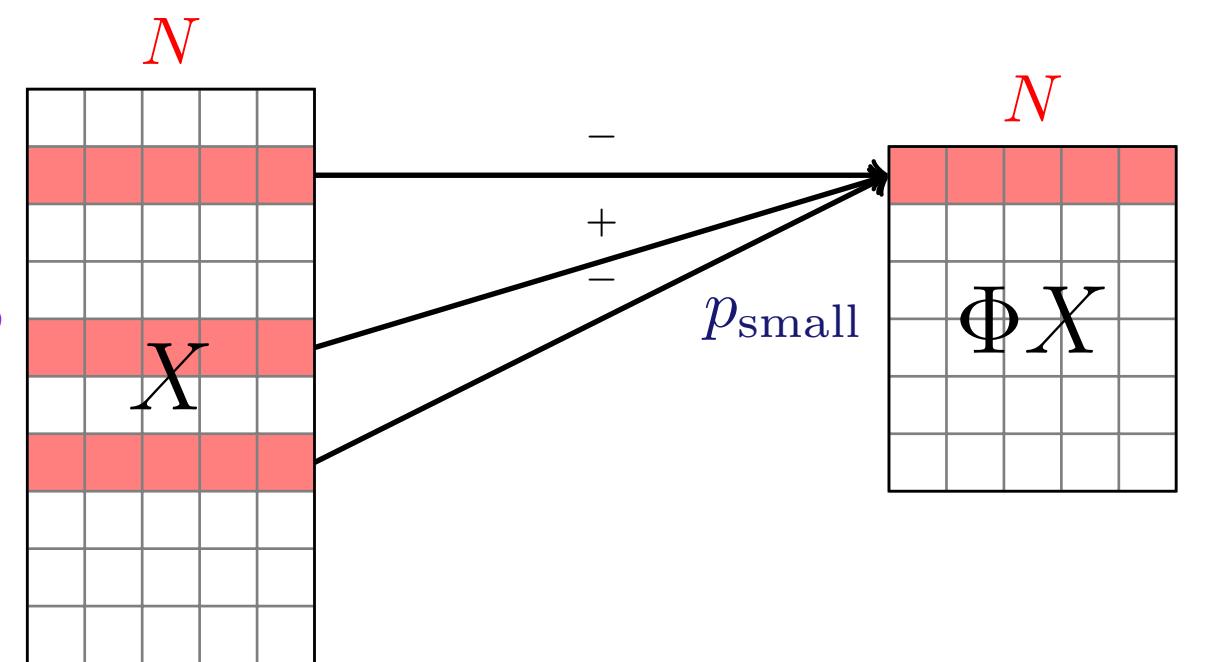
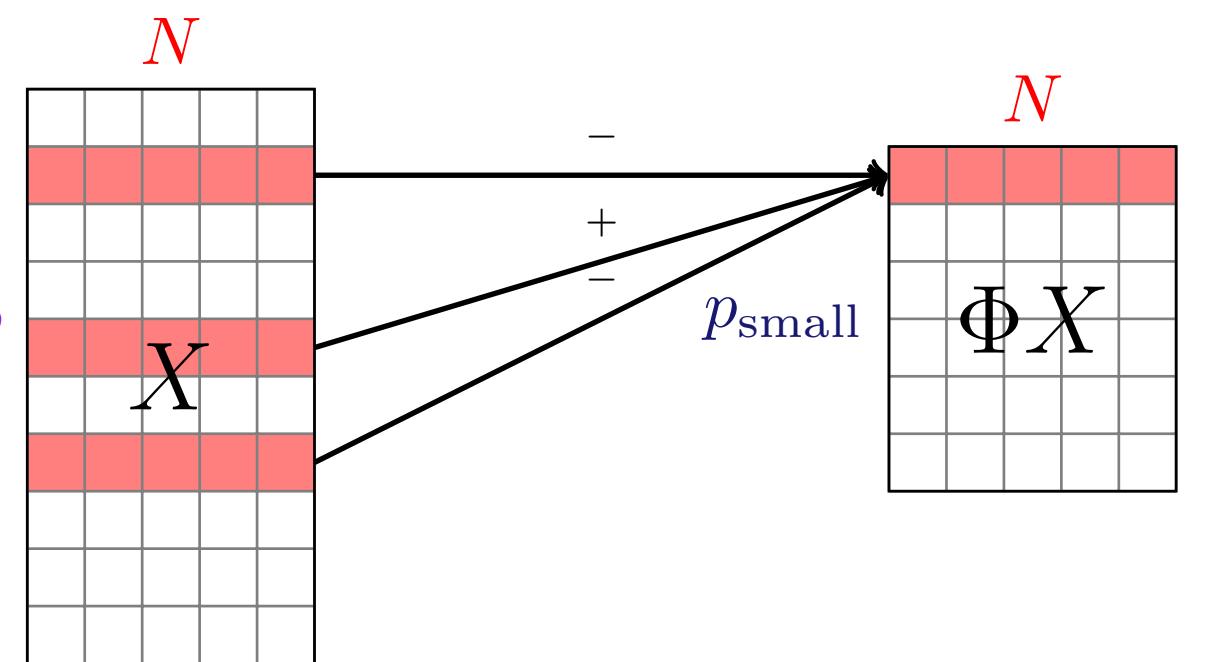
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

p

	± 1				± 1	
± 1						
		Φ		± 1		
				± 1		
	± 1	± 1				± 1

p_{small}



Introduced in Charikar et al. (2004), more analysis in, e.g., Clarkson and Woodruff (2017)

Based on **hash functions**

(as a side-effect, doesn't need full iid random variables for analysis to work)

See “appendix” of these slides for more details

Every column has exactly 1 nonzero entry, location chosen uniformly at random (and value is a Rademacher r.v.)

Has some guarantees, though not as good as a Gaussian (and not a JLT)

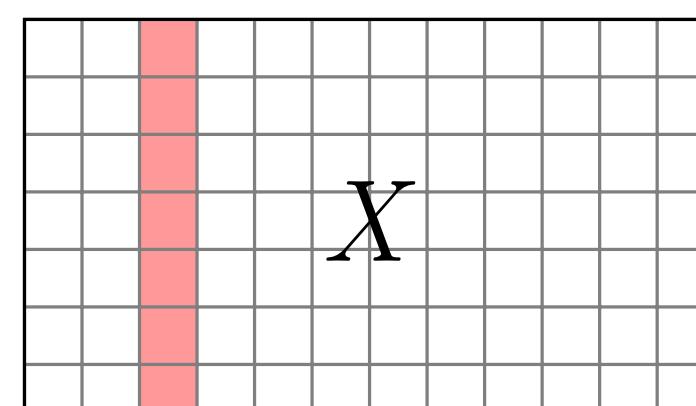
... but it's very **fast** to apply

cost: $\text{nnz}(X)$, so no more than pN (vs $p_{\text{small}}pN$ with a Gaussian)



new linear algebra application in: Malik & B. "Fast randomized matrix and tensor interpolative decomposition using CountSketch", Adv. Comp. Math (2020)

N



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 8: TensorSketch

Introduced in Pham, & Pagh (2013), more analysis
in, e.g., Diao, Zong, Sun, Woodruff (2018)

TensorSketch is just CountSketch when the input can be written as a tensor product
(for a special choice of the hash and sign functions)

$$\Phi_{\mathcal{T}} : \mathbb{R}^{\textcolor{violet}{p}} \rightarrow \mathbb{R}^{p_{\text{small}}}$$

$$\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)} \quad \text{where size is } \textcolor{violet}{p} = \textcolor{violet}{p}^{(1)} \cdot \textcolor{violet}{p}^{(2)}$$

Not (yet!) related to tensors

Two tricks:

- **Computationally**, combine small sketches in a convenient way
- For **analysis**, we lose **independence**, but CountSketch analysis didn't require full independence!

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 8: TensorSketch

Introduced in Pham, & Pagh (2013), more analysis
in, e.g., Diao, Zong, Sun, Woodruff (2018)

TensorSketch is just CountSketch when the input can be written as a tensor product
(for a special choice of the hash and sign functions)

$$\Phi_{\mathcal{T}} : \mathbb{R}^{\textcolor{violet}{p}} \rightarrow \mathbb{R}^{p_{\text{small}}}$$

$$\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)} \quad \text{where size is } \textcolor{violet}{p} = \textcolor{violet}{p}^{(1)} \cdot \textcolor{violet}{p}^{(2)}$$

Kronecker/tensor product of vectors

$$\mathbf{a} = \begin{bmatrix} \textcolor{red}{a}_1 \\ \textcolor{violet}{a}_2 \\ \textcolor{green}{a}_3 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} \textcolor{red}{b}_1 \\ \textcolor{violet}{b}_2 \\ \textcolor{green}{b}_3 \end{bmatrix} \quad \mathbf{a} \otimes \mathbf{b} = \begin{bmatrix} \textcolor{red}{a}_1 \mathbf{b} \\ \textcolor{violet}{a}_2 \mathbf{b} \\ \textcolor{green}{a}_3 \mathbf{b} \end{bmatrix} = \begin{bmatrix} \textcolor{red}{a}_1 \textcolor{red}{b}_1 \\ \textcolor{violet}{a}_1 \textcolor{violet}{b}_2 \\ \textcolor{green}{a}_1 \textcolor{green}{b}_3 \\ \textcolor{red}{a}_2 \textcolor{red}{b}_1 \\ \textcolor{violet}{a}_2 \textcolor{violet}{b}_2 \\ \textcolor{green}{a}_2 \textcolor{green}{b}_3 \\ \textcolor{red}{a}_3 \textcolor{red}{b}_1 \\ \textcolor{violet}{a}_3 \textcolor{violet}{b}_2 \\ \textcolor{green}{a}_3 \textcolor{green}{b}_3 \end{bmatrix}$$

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 8: TensorSketch

Introduced in Pham, & Pagh (2013), more analysis
in, e.g., Diao, Zong, Sun, Woodruff (2018)

TensorSketch is just CountSketch when the input can be written as a tensor product
(for a special choice of the hash and sign functions)

$$\Phi_{\mathcal{T}} : \mathbb{R}^{\textcolor{violet}{p}} \rightarrow \mathbb{R}^{p_{\text{small}}}$$

$$\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)} \quad \text{where size is } \textcolor{violet}{p} = \textcolor{violet}{p}^{(1)} \cdot \textcolor{violet}{p}^{(2)}$$

Kronecker/tensor product of vectors

$$\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \\ a_3 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

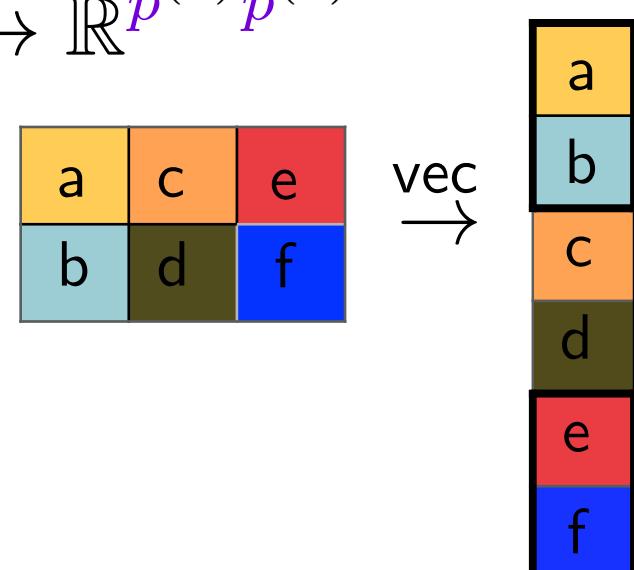
$$\mathbf{a} \otimes \mathbf{b} = \begin{bmatrix} \mathbf{a}_1 \mathbf{b} \\ \mathbf{a}_2 \mathbf{b} \\ \mathbf{a}_3 \mathbf{b} \end{bmatrix} = \begin{bmatrix} a_1 b_1 \\ a_1 b_2 \\ a_1 b_3 \\ a_2 b_1 \\ a_2 b_2 \\ a_2 b_3 \\ a_3 b_1 \\ a_3 b_2 \\ a_3 b_3 \end{bmatrix}$$

another equivalent definition:

$$\mathbf{a} \otimes \mathbf{b} = \text{vec}_{\text{col}} (\mathbf{b} \mathbf{a}^T)$$

$$\mathbf{b} \mathbf{a}^T = \begin{bmatrix} b_1 a_1 & b_1 a_2 & b_1 a_3 \\ b_2 a_1 & b_2 a_2 & b_2 a_3 \\ b_3 a_1 & b_3 a_2 & b_3 a_3 \end{bmatrix} = \begin{bmatrix} a_1 b_1 & a_2 b_1 & a_3 b_1 \\ a_1 b_2 & a_2 b_2 & a_3 b_2 \\ a_1 b_3 & a_2 b_3 & a_3 b_3 \end{bmatrix}$$

$$\text{vec}: \mathbb{R}^{\textcolor{violet}{p}^{(1)} \times \textcolor{violet}{p}^{(2)}} \rightarrow \mathbb{R}^{\textcolor{violet}{p}^{(1)} \textcolor{violet}{p}^{(2)}}$$



$$\text{mat} = \text{vec}^{-1}$$

more generally,

$$(\mathbf{A} \otimes \mathbf{B}) \text{vec}_{\text{col}} (\mathbf{X}) = \text{vec}_{\text{col}} (\mathbf{B} \mathbf{X} \mathbf{A}^T)$$

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch: applying to matrices

From vector case $\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)}$ to matrix case $\mathbf{A} = \mathbf{A}^{(1)} \otimes \mathbf{A}^{(2)} \otimes \dots \otimes \mathbf{A}^{(N)}$

Recall the **Kronecker** product:

$$\begin{aligned} \mathbf{A} \otimes \mathbf{B} &= \begin{bmatrix} a_{11}\mathbf{B} & a_{12}\mathbf{B} & \cdots & a_{1J}\mathbf{B} \\ a_{21}\mathbf{B} & a_{22}\mathbf{B} & \cdots & a_{2J}\mathbf{B} \\ \vdots & \vdots & \ddots & \vdots \\ a_{I1}\mathbf{B} & a_{I2}\mathbf{B} & \cdots & a_{IJ}\mathbf{B} \end{bmatrix} \in \mathbb{R}^{IK \times JL} \\ &= [\mathbf{a}_1 \otimes \mathbf{b}_1 \quad \mathbf{a}_1 \otimes \mathbf{b}_2 \quad \mathbf{a}_1 \otimes \mathbf{b}_3 \quad \cdots \quad \mathbf{a}_J \otimes \mathbf{b}_{L-1} \quad \mathbf{a}_J \otimes \mathbf{b}_L] \end{aligned}$$

Khatri-Rao product $\mathbf{A} \odot \mathbf{B} = [\mathbf{a}_1 \otimes \mathbf{b}_1 \quad \mathbf{a}_2 \otimes \mathbf{b}_2 \cdots \mathbf{a}_L \otimes \mathbf{b}_L] \in \mathbb{R}^{IK \times L} \quad (J = L)$

Observe

pointwise multiplication (aka Hadamard product)

$$\mathbf{a} \circ \mathbf{b} = (\mathbf{a}^\top \odot \mathbf{b}^\top)^\top$$

and with some work, $\Phi_{\mathcal{T}}(\mathbf{A}) = \text{FFT}^{-1} \left(\left(\bigodot_{n=1}^N \left(\text{FFT} \left(\mathbf{S}^{(n)} \mathbf{A}^{(n)} \right) \right)^\top \right)^\top \right)$
 if $\mathbf{A} = \mathbf{A}^{(1)} \otimes \mathbf{A}^{(2)} \otimes \dots \otimes \mathbf{A}^{(N)}$

small countSketch

FFTs are doing circular convolution,
 i.e., multiplying polynomials modulo
 a fixed modulus. See appendix

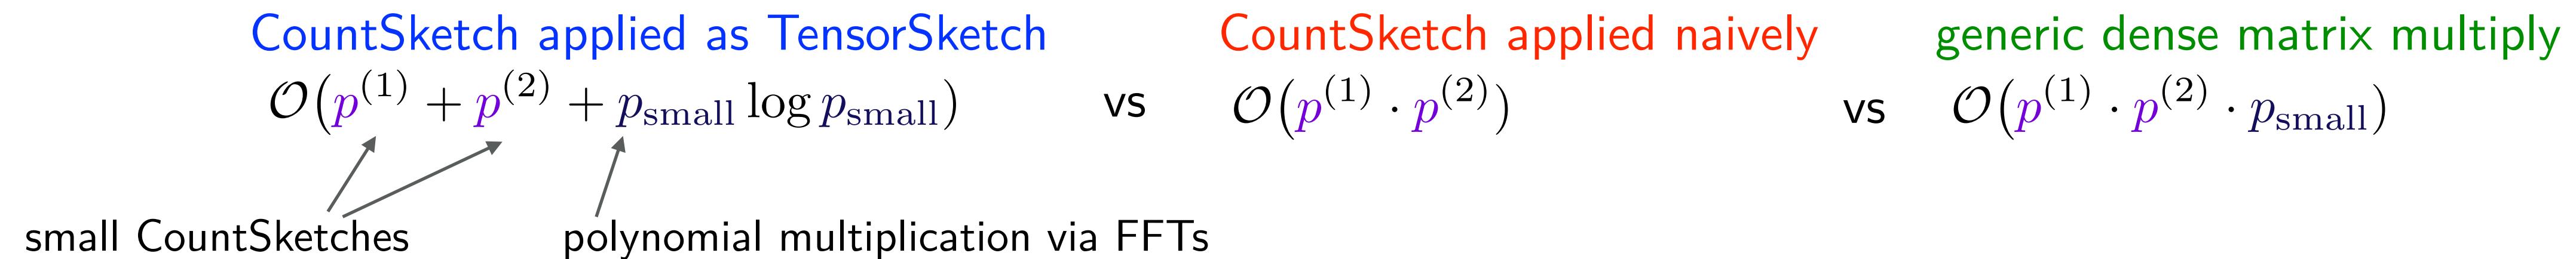
1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch: complexity analysis

$$\Phi_{\mathcal{T}} : \mathbb{R}^{\textcolor{violet}{p}} \rightarrow \mathbb{R}^{p_{\text{small}}}$$

$$p = \textcolor{violet}{p}^{(1)} \cdot \textcolor{violet}{p}^{(2)}$$

Complexity (per mat-vec):



 Savings grow as we have more tensor products

If $\textcolor{violet}{p} = \textcolor{violet}{p}^{(1)} \textcolor{violet}{p}^{(2)} \dots \textcolor{violet}{p}^{(q)}$

$$\mathcal{O}(\textcolor{violet}{p}^{(1)} + \textcolor{violet}{p}^{(2)} + \dots + \textcolor{violet}{p}^{(q)} + p_{\text{small}} \log p_{\text{small}}) \quad \text{vs} \quad \mathcal{O}(\textcolor{violet}{p}^{(1)} \textcolor{violet}{p}^{(2)} \dots \textcolor{violet}{p}^{(q)})$$

(Even larger savings if input factor matrices are **sparse**)

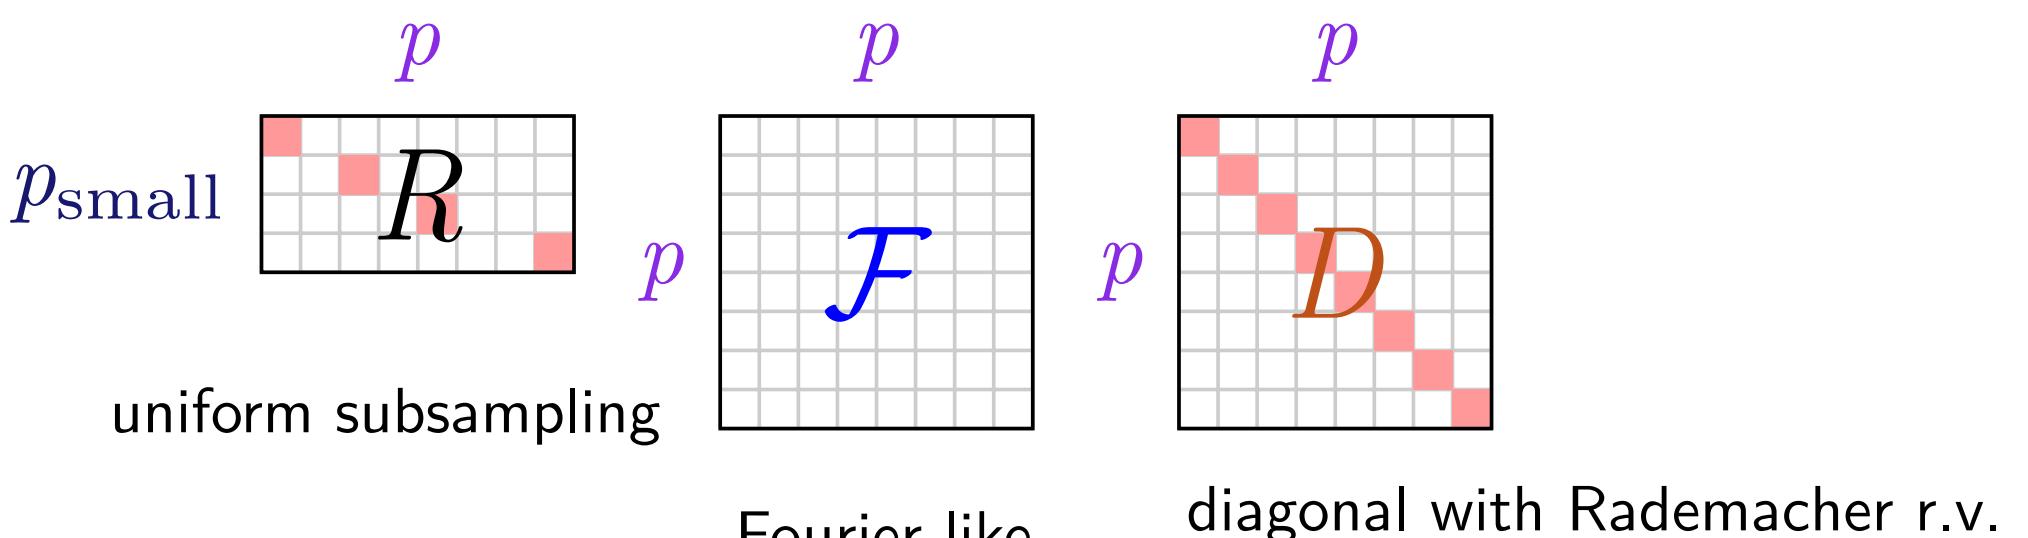
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 9: Kronecker Fast Johnson-Lindenstrauss sketch

As for TensorSketch, suppose each column of data looks like $\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)}$

$$\text{sizes: } p = p^{(1)} \cdot p^{(2)}$$

Recall the Fast Johnson-Lindenstrauss (FJLT): $\Phi = R\mathcal{F}\mathcal{D}$



The **Kronecker** FJLT has a similar structure:

$$\Phi_{\text{KFJLT}} = R \left(\mathcal{F}^{(1)} \mathcal{D}^{(1)} \otimes \mathcal{F}^{(2)} \mathcal{D}^{(2)} \right)$$

which is efficient to apply since

$$R \left(\mathcal{F}^{(1)} \mathcal{D}^{(1)} \otimes \mathcal{F}^{(2)} \mathcal{D}^{(2)} \right) \cdot \left(\mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)} \right) = R \left(\left(\mathcal{F}^{(1)} \mathcal{D}^{(1)} \mathbf{v}^{(1)} \right) \otimes \left(\mathcal{F}^{(2)} \mathcal{D}^{(2)} \mathbf{v}^{(2)} \right) \right)$$

and uniform subsampling can be done implicitly without forming this



- Battaglino, Ballard, Kolda, “A practical randomized CP tensor decomposition”, *SIAM J. Matrix Anal. Appl.* (2018)
- Jin, Kolda, Ward. “Faster Johnson–Lindenstrauss transforms via Kronecker products,” *Information and Inference: A Journal of the IMA* (2021)
- Malik & B. “Guarantees for the Kronecker fast Johnson–Lindenstrauss transform using a coherence and sampling argument,” *Lin. Alg. & its Applications* (2020)

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 9: Kronecker Fast Johnson-Lindenstrauss sketch

$$\Phi_{\text{KFJLT}} = R \left(\mathcal{F}^{(1)} \mathcal{D}^{(1)} \otimes \mathcal{F}^{(2)} \mathcal{D}^{(2)} \otimes \dots \otimes \mathcal{F}^{(q)} \mathcal{D}^{(q)} \right) \quad p = p^{(1)} p^{(2)} \dots p^{(q)}$$

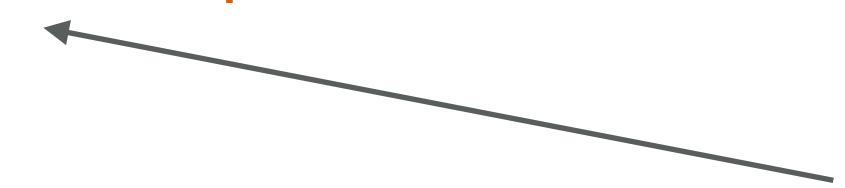
Theorem (Thm 4.2 in Malik & B. 2020)

If all datapoints have a Kronecker structure, $p_{\text{small}} > \frac{16}{2} \frac{4^q}{\varepsilon^2} \log \left(\frac{4N^2(q+1)}{\delta} \right) \log \left(\frac{4p^{(1)} N^2(q+1)}{\delta} \right) \dots \log \left(\frac{4p^{(q)} N^2(q+1)}{\delta} \right)$
and $\Phi = \Phi_{\text{KFJLT}}$

then with probability at least $1 - \delta$

$$\forall x_i, x_j \in \{x_1, \dots, x_N\} \subset \mathbb{R}^p \quad 1 - \varepsilon \leq \frac{\|\Phi x_i - \Phi x_j\|_2^2}{\|x_i - x_j\|_2^2} \leq 1 + \varepsilon$$

proof idea: FJLT keeps leverage scores almost uniform. With Kronecker product structure, this is (almost) true also, due to properties of **leverage scores of Kronecker products**



in fact, can extend. See: Malik, Xu, Cheng, B., Doostan, Narayan.

"Fast Algorithms for Monotone Lower Subsets of Kronecker Least Squares Problems",
<https://arxiv.org/abs/2209.05662>



o Malik & B. "Guarantees for the Kronecker fast Johnson–Lindenstrauss transform using a coherence and sampling argument," *Lin. Alg. & its Applications* (2020)

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 9: Kronecker Fast Johnson-Lindenstrauss sketch

$$\Phi_{\text{KFJLT}} = R \left(\mathcal{F}^{(1)} D^{(1)} \otimes \mathcal{F}^{(2)} D^{(2)} \otimes \dots \otimes \mathcal{F}^{(q)} D^{(q)} \right) \quad p = p^{(1)} p^{(2)} \dots p^{(q)}$$

Theorem (Thm 4.2 in Malik & B. 2020)

If all datapoints have a Kronecker structure, $p_{\text{small}} > \frac{16}{2} \frac{4^q}{\varepsilon^2} \log \left(\frac{4N^2(q+1)}{\delta} \right) \log \left(\frac{4p^{(1)} N^2(q+1)}{\delta} \right) \dots \log \left(\frac{4p^{(q)} N^2(q+1)}{\delta} \right)$
and $\Phi = \Phi_{\text{KFJLT}}$

then with probability at least $1 - \delta$

$$\forall x_i, x_j \in \{x_1, \dots, x_N\} \subset \mathbb{R}^p \quad 1 - \varepsilon \leq \frac{\|\Phi x_i - \Phi x_j\|_2^2}{\|x_i - x_j\|_2^2} \leq 1 + \varepsilon$$

Theorem (Johnson-Lindenstrauss, 1984 (and Indyk-Motwani, 1998))

Choose $\Phi = \Phi_{\text{randn}}$ with $p_{\text{small}} \propto \varepsilon^{-2} \log N$ iid rows each $\mathcal{N}(0, p/p_{\text{small}})$, then for all $x_i, x_j \in \{x_1, \dots, x_N\} \subset \mathbb{R}^p$,

for reference:

$$1 - \varepsilon \leq \frac{\|\Phi x_i - \Phi x_j\|_2}{\|x_i - x_j\|_2} \leq 1 + \varepsilon$$

with constant probability.

no dependence on p

OUTLINE

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

One application of sketching: least-squares

Least-squares

Solve $x_{\text{LS}} = \operatorname{argmin} \|Ax - b\|$, A is $p \times N$ with $p \gg N$

Approach 1: randomization to quickly find preconditioner

BLENDENPIK, LSRN

“sketch-to-precondition”

- ▶ If $A = QR$ is a QR-decomposition, then AR^{-1} is well-conditioned
- ▶ Idea: do QR-decomp on reduced-dimension matrix ΦA
- ▶ On large, very ill-conditioned matrices (and tall), about $4\times$ faster than LAPACK

One application of sketching: least-squares

Least-squares

Solve $x_{\text{LS}} = \operatorname{argmin} \|Ax - b\|$, A is $p \times N$ with $p \gg N$

Approach 1: randomization to quickly find preconditioner

BLENDENPIK, LSRN

“sketch-to-precondition”

- ▶ If $A = QR$ is a QR-decomposition, then AR^{-1} is well-conditioned
- ▶ Idea: do QR-decomp on reduced-dimension matrix ΦA
- ▶ On large, very ill-conditioned matrices (and tall), about $4\times$ faster than LAPACK

Approach 2: directly solve sketched problem

Sarlos, Woodruff, Mahoney, etc.

“sketch-to-solve”

- ▶ Directly solve $\min_x \|\Phi(Ax - b)\|$
- ▶ Theoretical bounds on objective error if Φ has $p_{\text{small}} = N \text{polylog}(N)$ rows
 - ▶ i.e., p_{small} is independent of p

For a sub-sampling sketch, this lets you avoid even computing some entries of b , which can be very useful when each entry of b is, e.g., the output of a simulation

One application of sketching: least-squares

Least-squares

Solve $x_{\text{LS}} = \operatorname{argmin} \|Ax - b\|$, A is $p \times N$ with $p \gg N$

Approach 1: randomization to quickly find preconditioner

BLENDENPIK, LSRN

- ▶ If $A = QR$ is a QR-decomposition, then AR^{-1} is well-conditioned
- ▶ Idea: do QR-decomp on reduced-dimension matrix ΦA
- ▶ On large, very ill-conditioned matrices (and tall), about $4\times$ faster than LAPACK

Approach 2: directly solve sketched problem

Sarlos, Woodruff, Mahoney, etc.

- ▶ Directly solve $\min_x \|\Phi(Ax - b)\|$
- ▶ Theoretical bounds on objective error if Φ has $p_{\text{small}} = N \text{polylog}(N)$ rows
 - ▶ i.e., p_{small} is independent of p

Other applications: SVDs, matrix multiplies, Nyström/CUR/Interpolative Decompositions, QR without pivoting

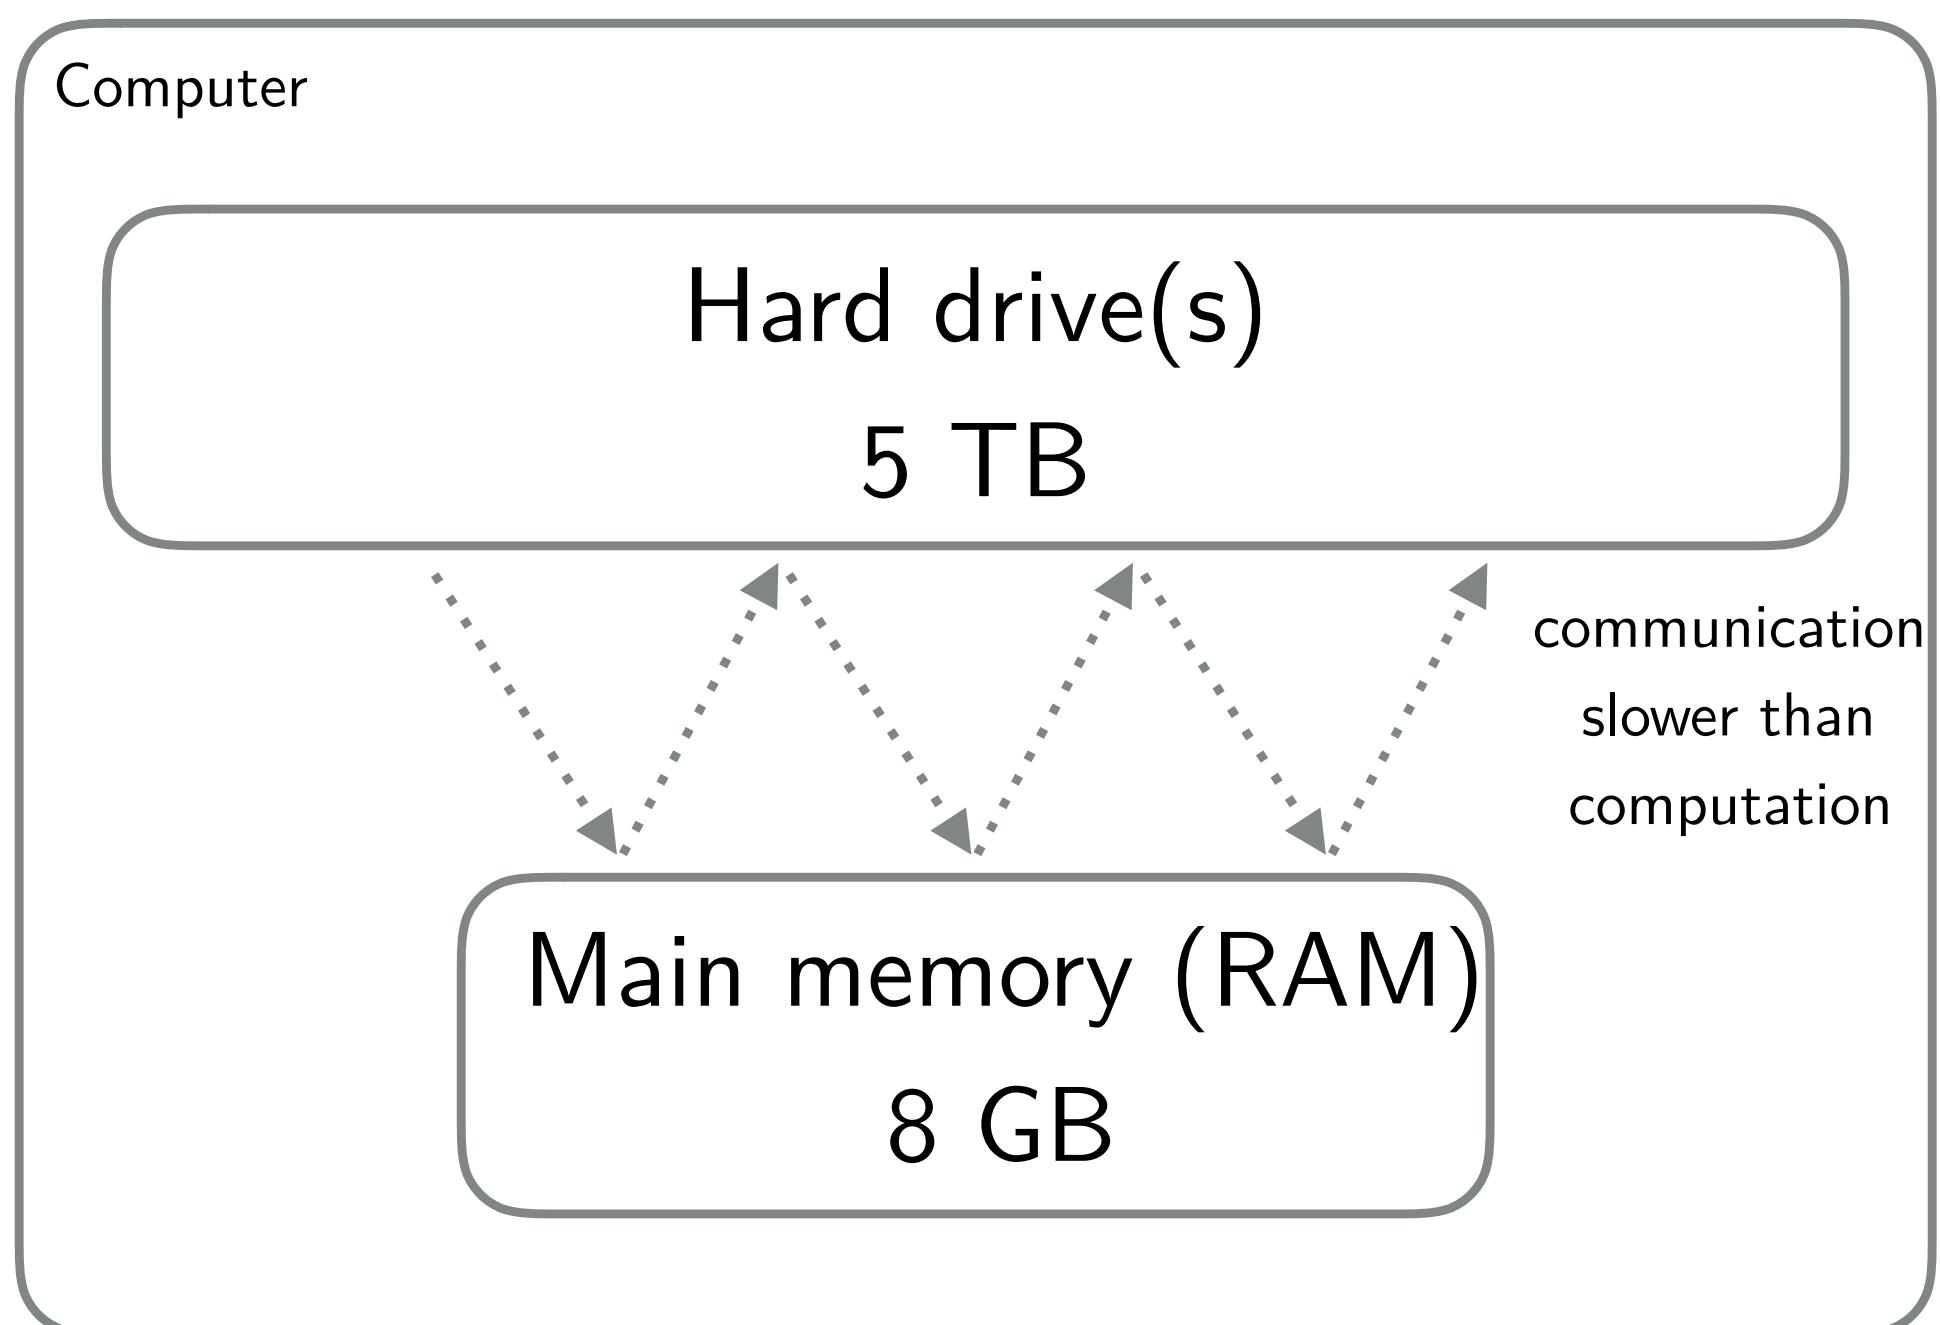
OUTLINE

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Interlude: one-pass methods

Single computer



Often want to minimize number of **passes** through a dataset if it's large

- *communication cost* can be significant
- in the extreme case, new data is constantly **streaming** in and we could never store it all

Think of a **pass** as a “for” loop through the dataset

Streaming is a special case of a “**one-pass**” method

Example: google indexing the web is streaming

1. Google finds a website and downloads the site
2. Applies algorithm to index/categorize/rank it
3. Discards the website (keeps only meta-data, ranking, etc.)
4. Never has to store entire WWW *all at once*

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Interlude: one-pass methods

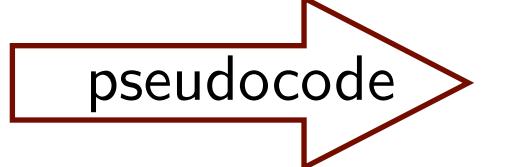
Example: computing the sample variance $\hat{\sigma}^2$

2-pass method

pass #1

$$\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$$

pseudocode



pass #2

$$\hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \hat{\mu})^2$$

pseudocode



$\left\{ \begin{array}{l} \hat{\mu} = 0 \\ \text{for } i = 1, \dots, n \\ \quad \hat{\mu} \leftarrow \hat{\mu} + x_i \\ \hat{\mu} \leftarrow \frac{1}{n} \hat{\mu} \\ \hat{\sigma}^2 = 0 \\ \text{for } i = 1, \dots, n \\ \quad \hat{\sigma}^2 \leftarrow \hat{\sigma}^2 + (x_i - \hat{\mu})^2 \\ \hat{\sigma}^2 \leftarrow \frac{1}{n-1} \hat{\sigma}^2 \end{array} \right.$

1-pass method

motivation: $\text{Var}[X] = \mathbb{E}[X^2] - \mathbb{E}[X]^2$

pass #1

$\left\{ \begin{array}{l} \hat{\mu} = 0, s = 0 \\ \text{for } i = 1, \dots, n \\ \quad \hat{\mu} \leftarrow \hat{\mu} + x_i \\ \quad s \leftarrow s + x_i^2 \\ \hat{\sigma}^2 \leftarrow \frac{1}{n-1} (s - \hat{\mu}^2) \end{array} \right.$

Interlude: one-pass methods

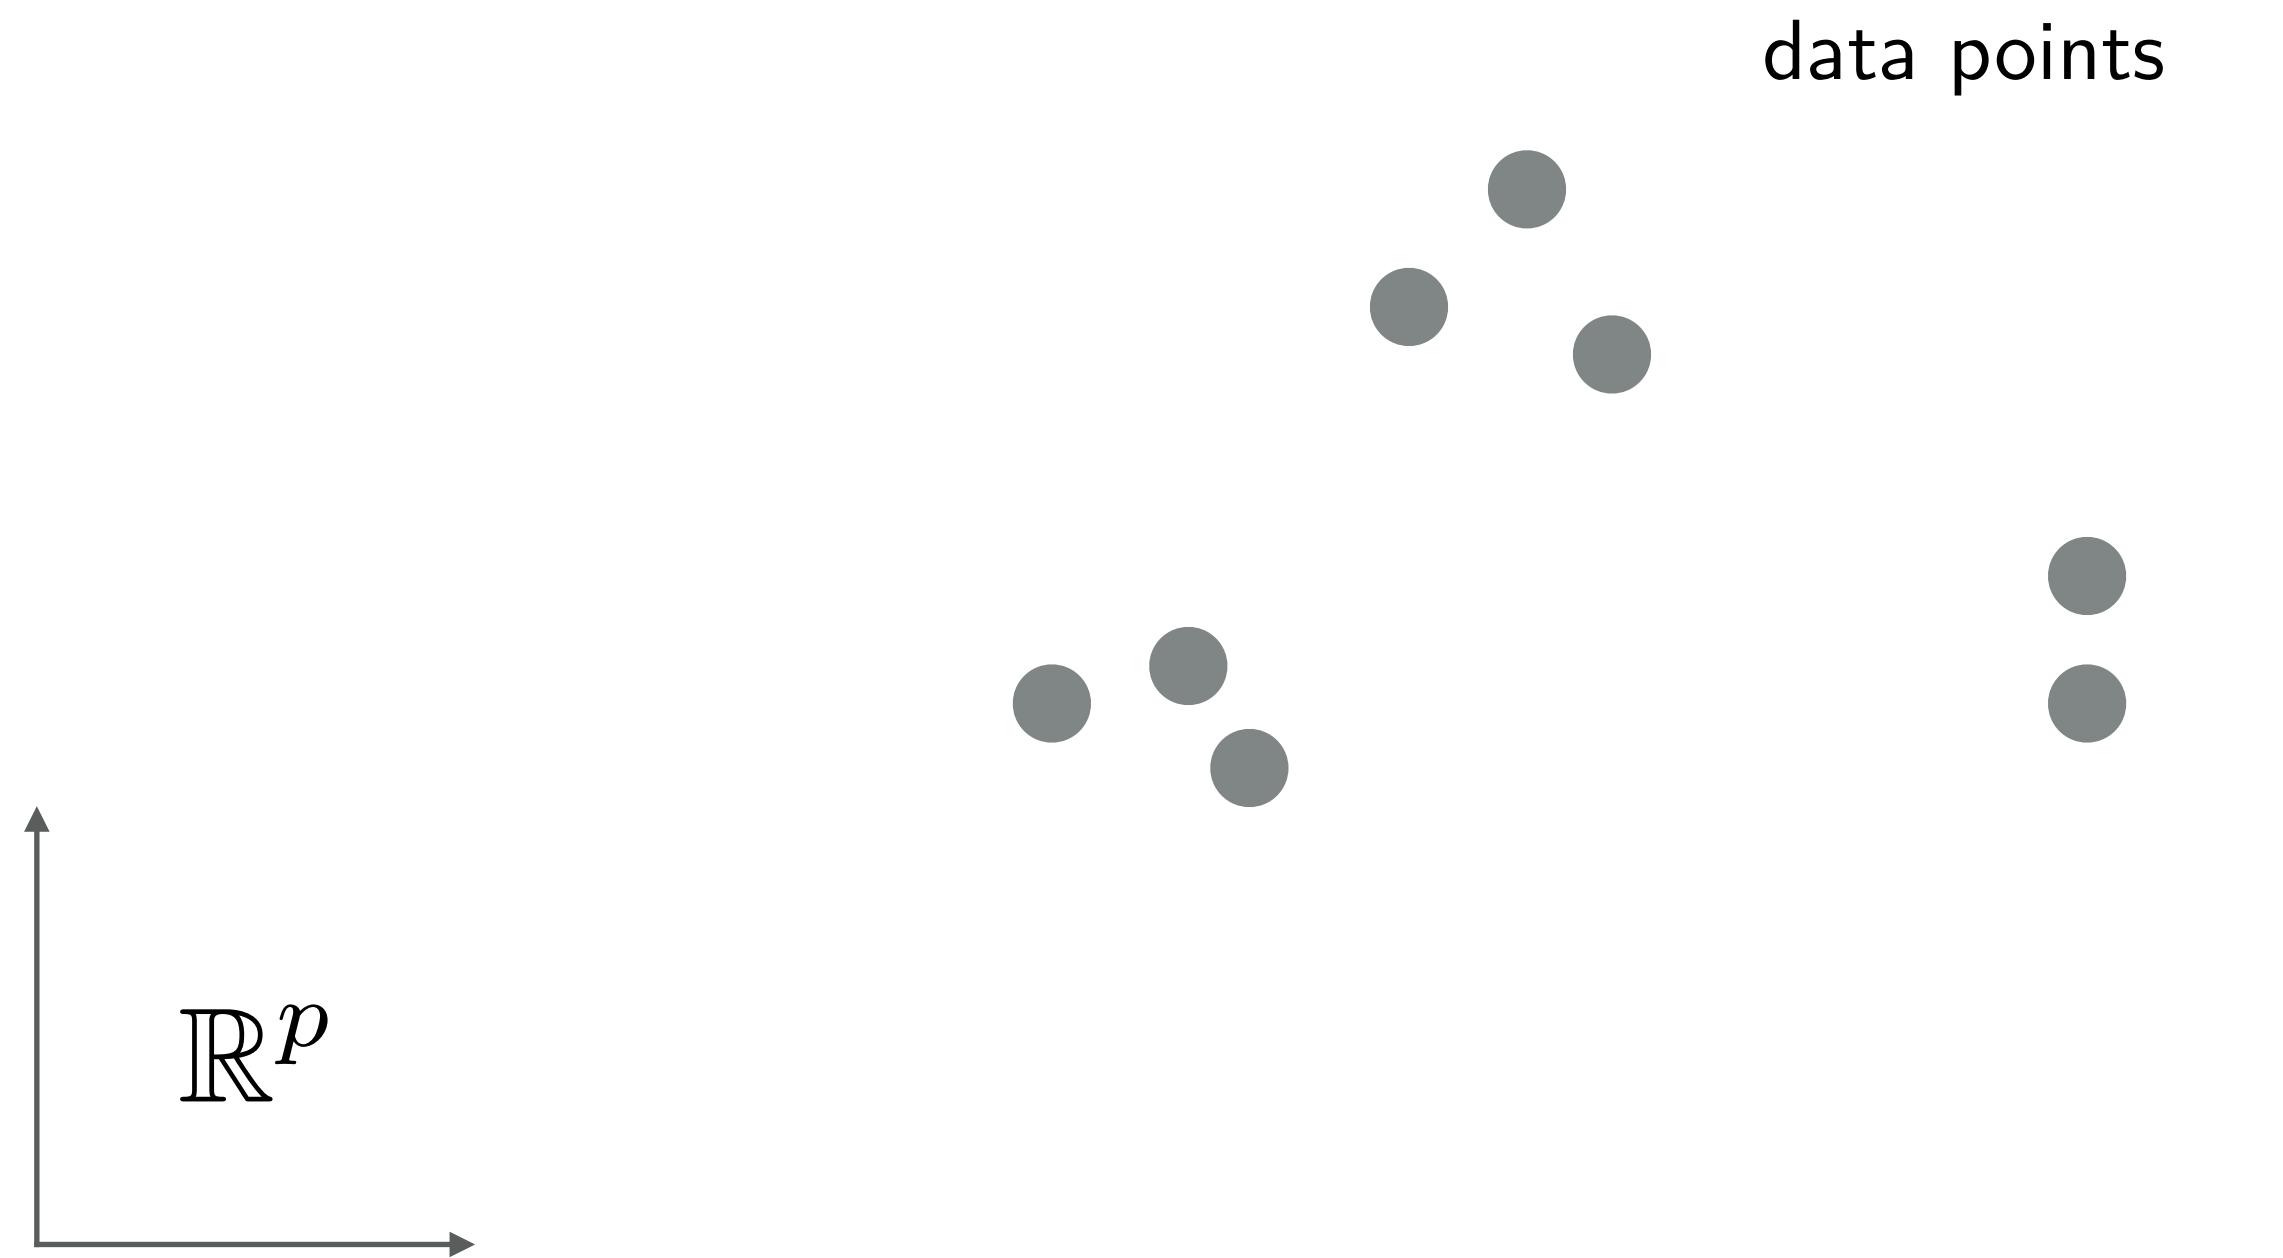
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

A major application of **sketching** is to create **one-pass** algorithms:

we spend one-pass to apply the sketch, and the sketched data is then small enough that we don't “count it” in our budget anymore

Machine Learning 101: K-means clustering

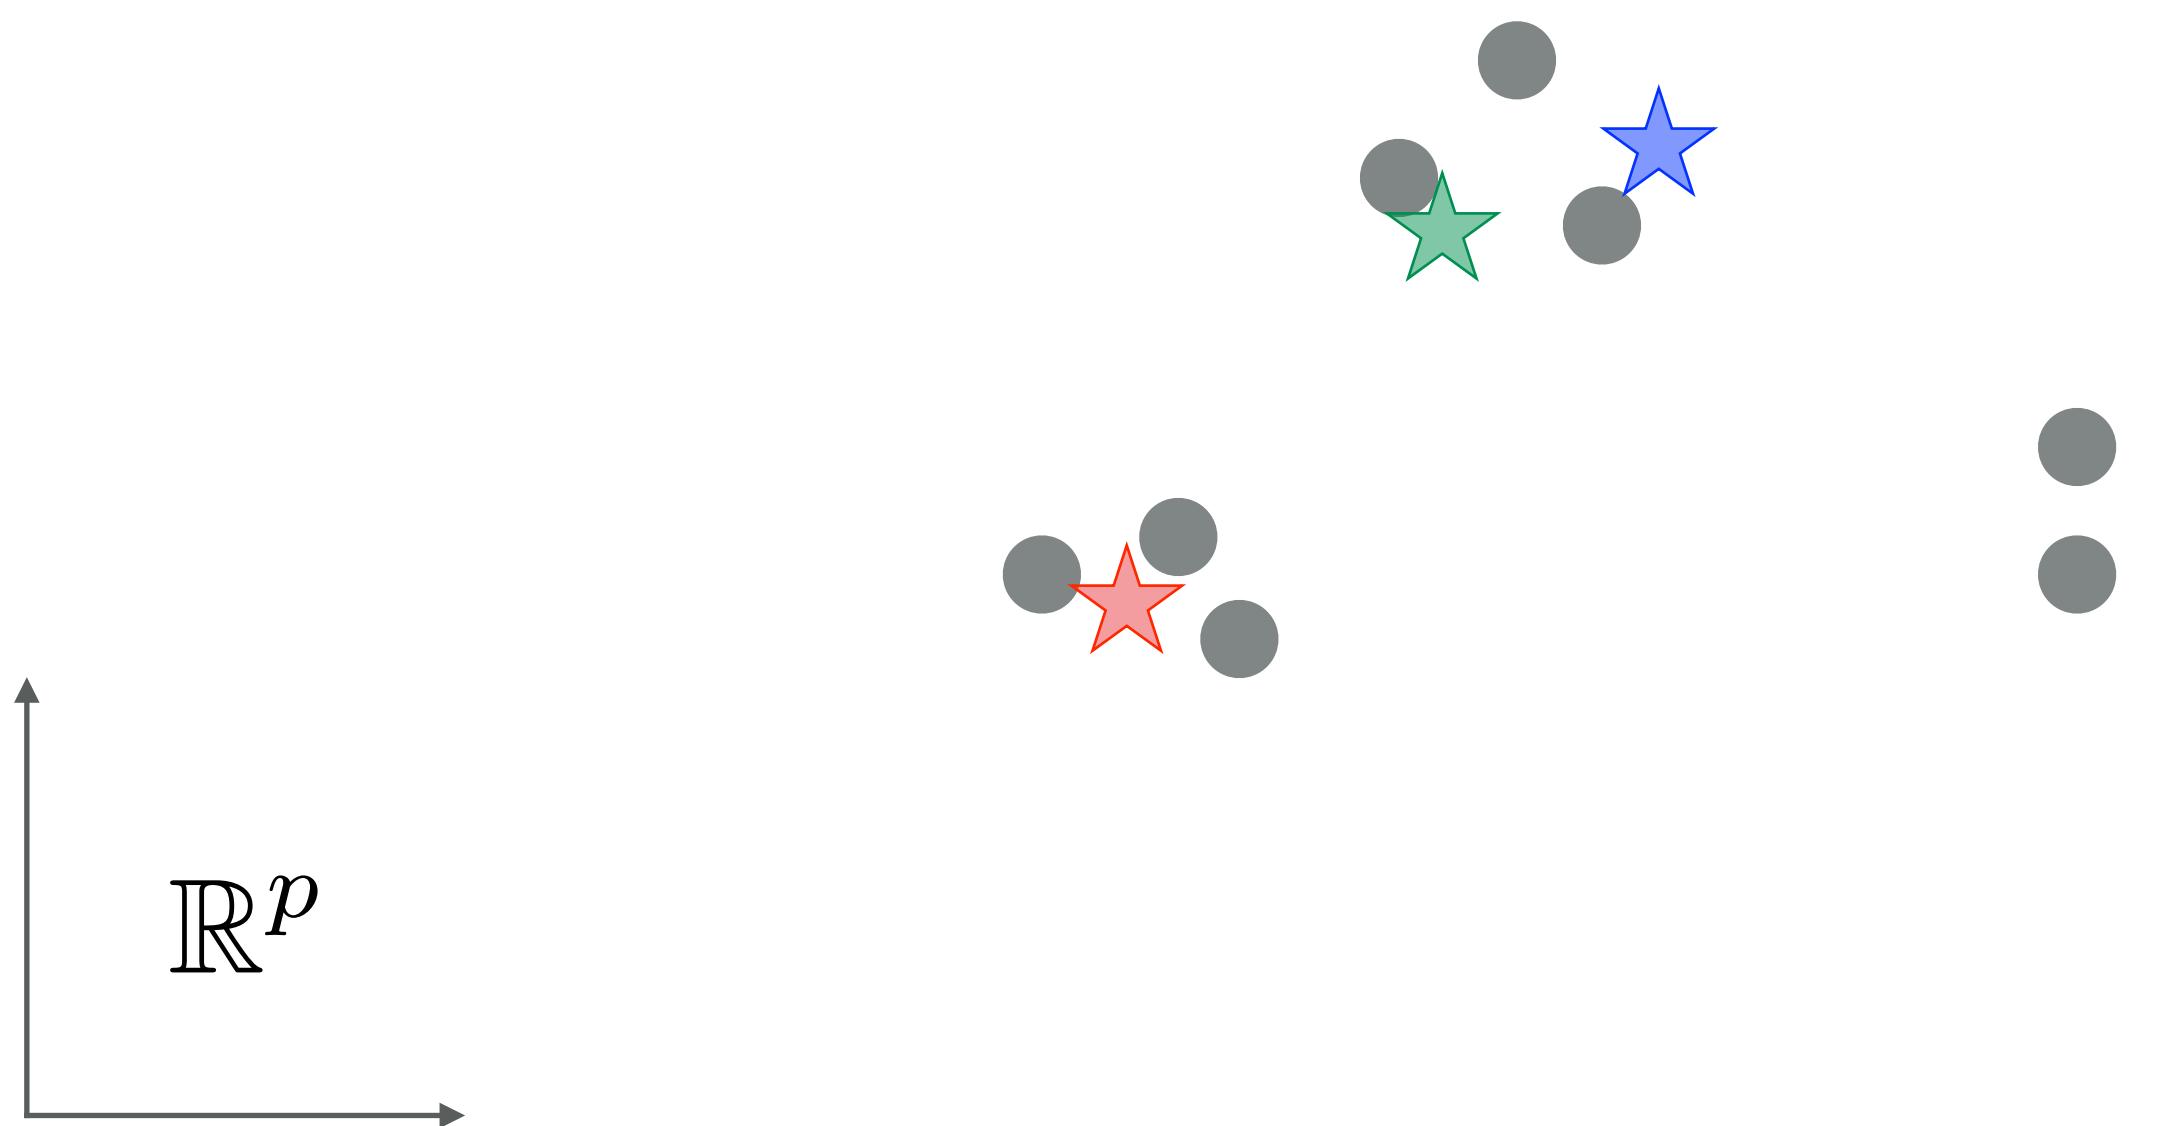
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Machine Learning 101: K-means clustering

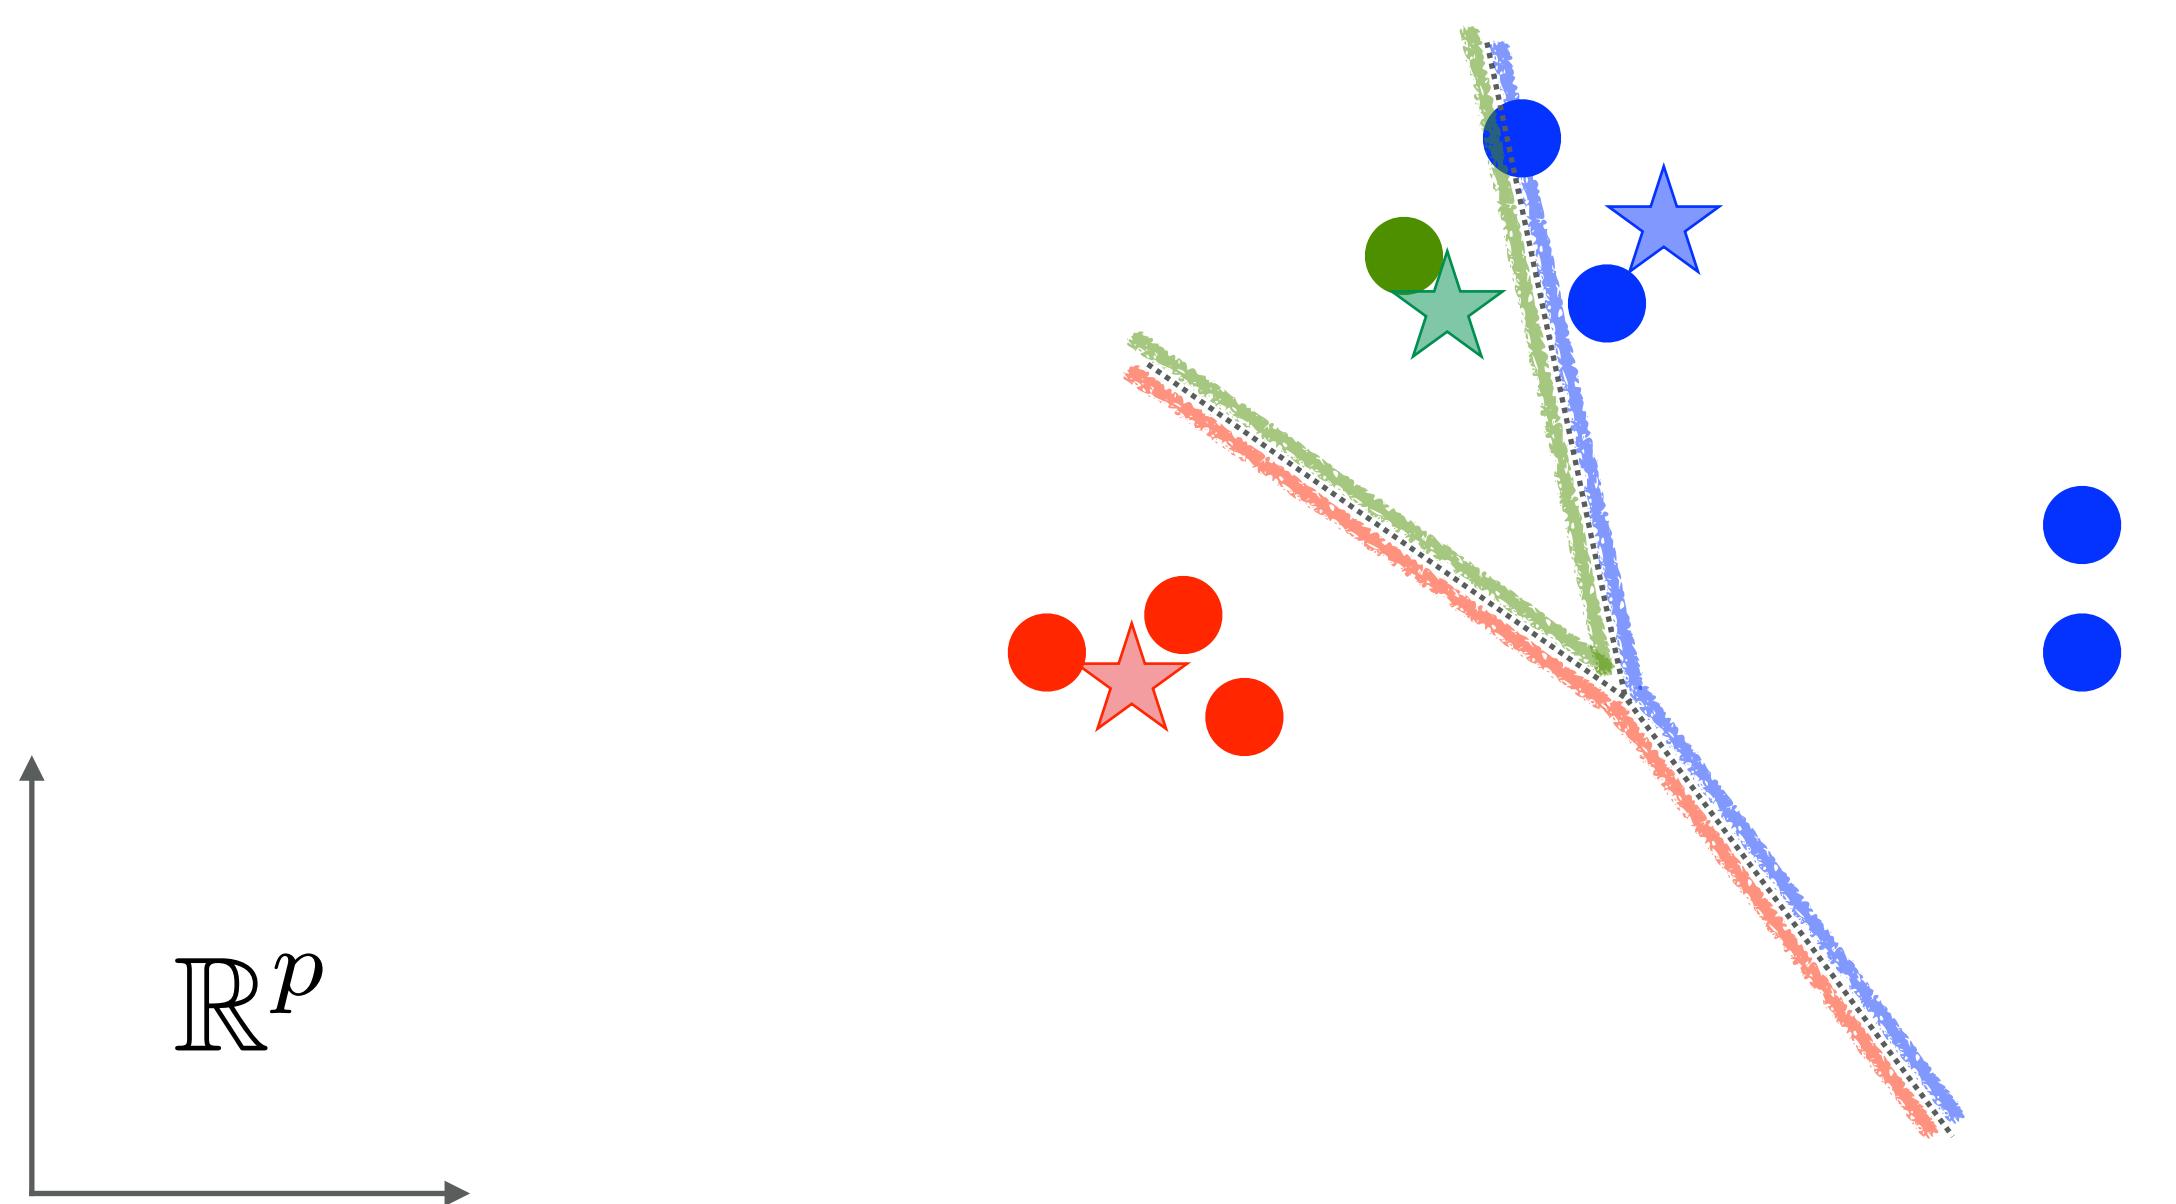
Step: initialize with guesses for cluster **centers**



Machine Learning 101: K-means clustering

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

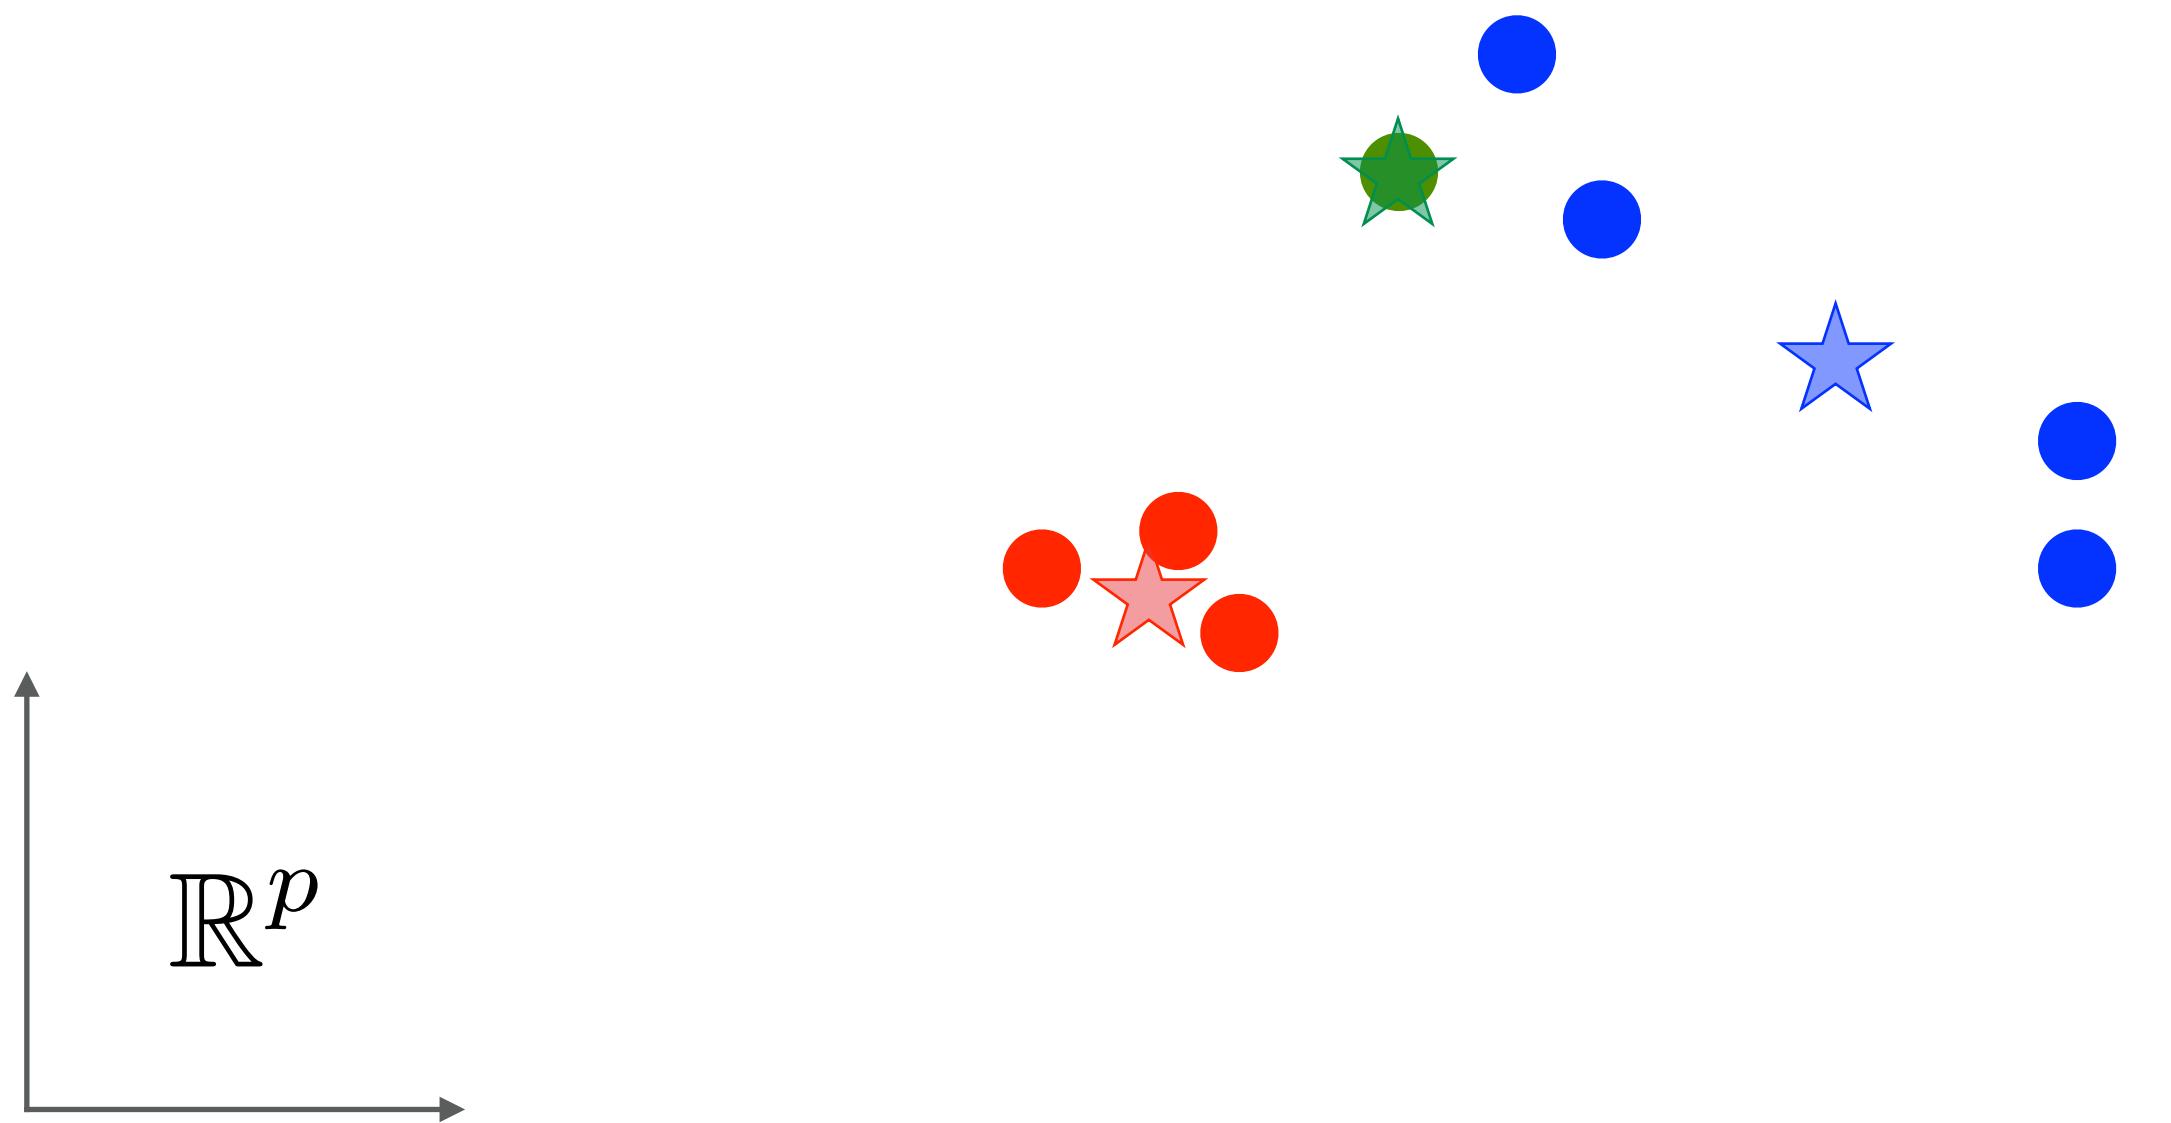
Step: update **assignments**



Machine Learning 101: K-means clustering

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

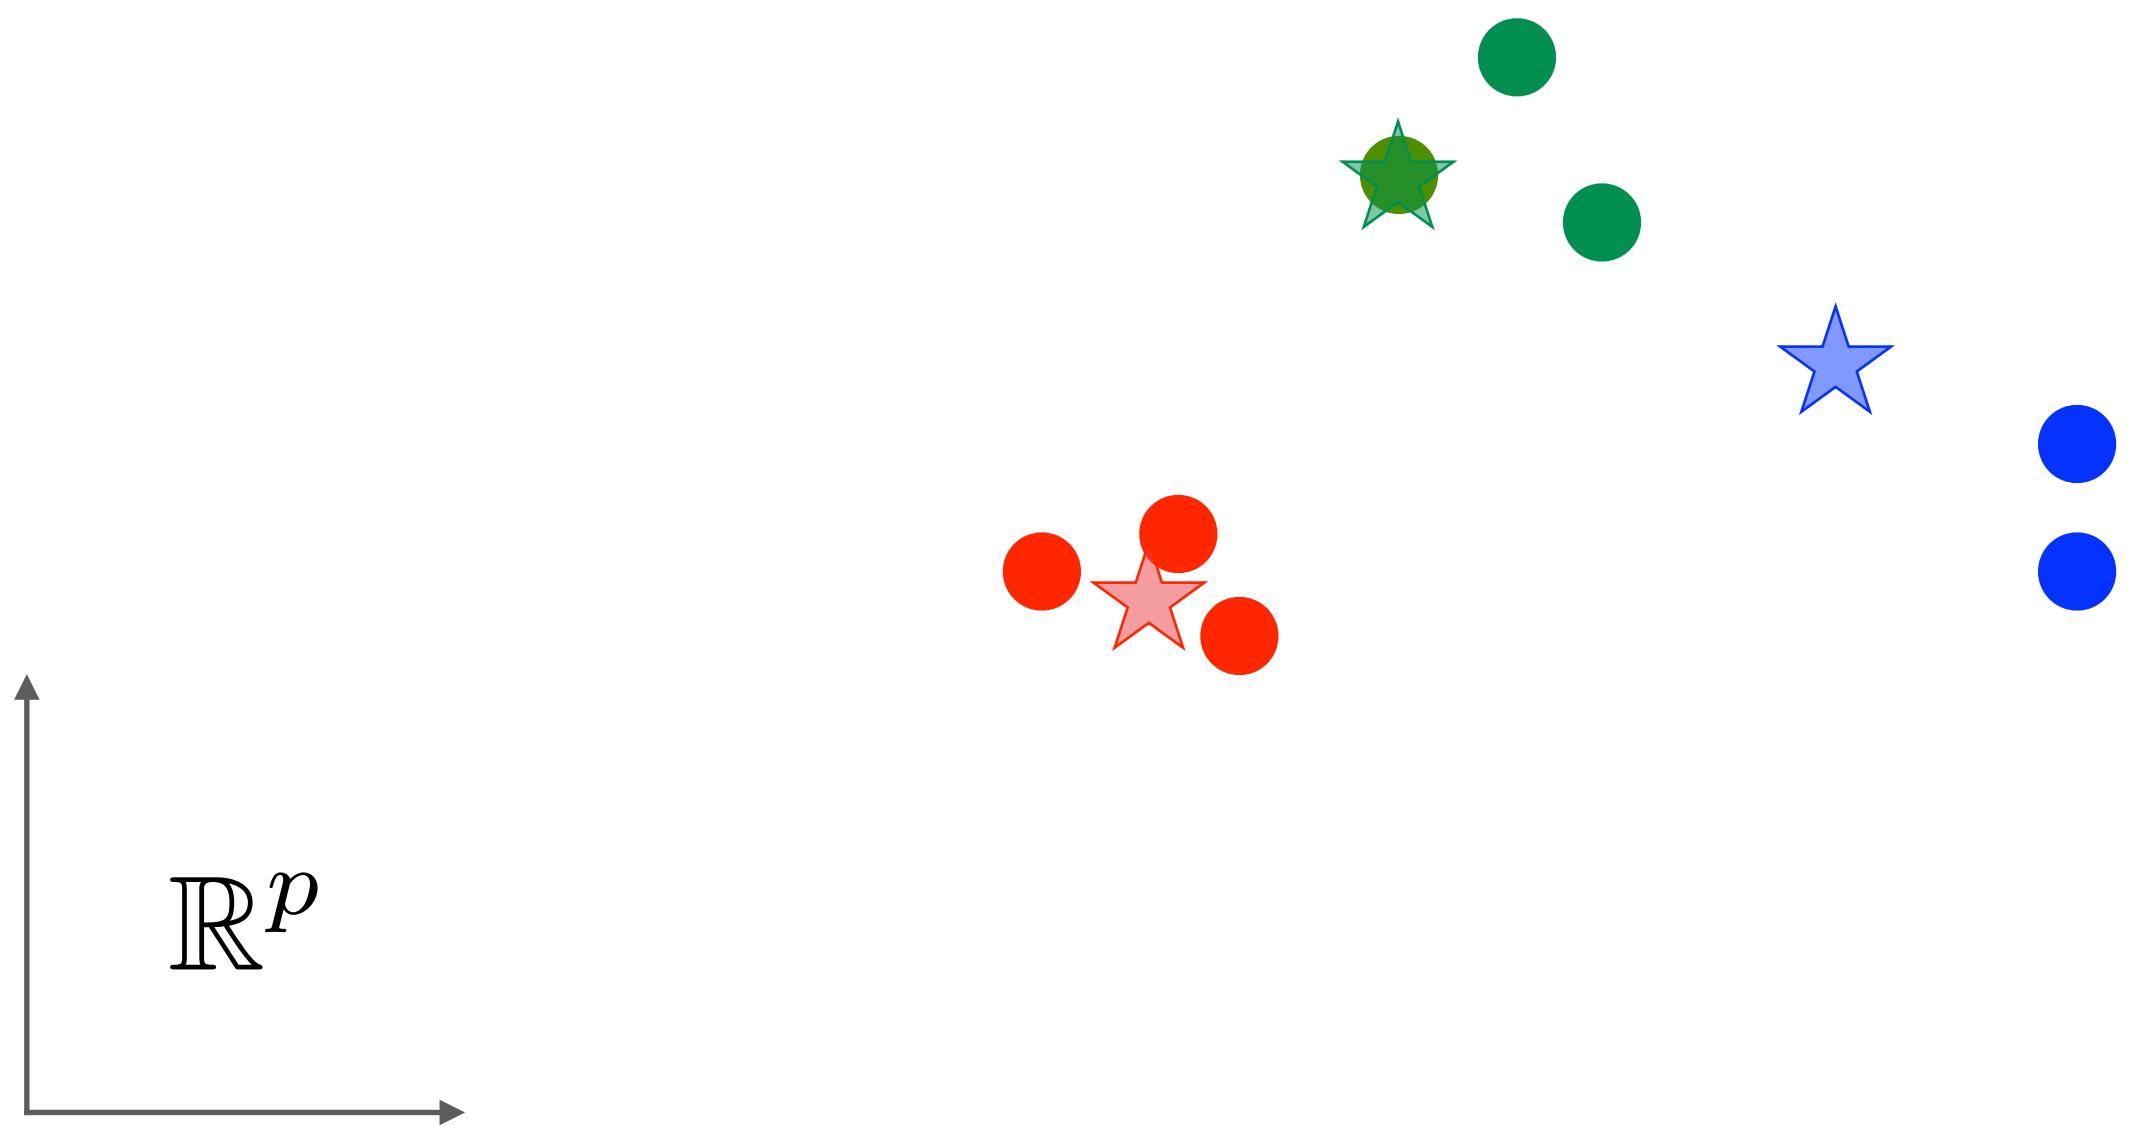
Step: recompute **cluster means**



Machine Learning 101: K-means clustering

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

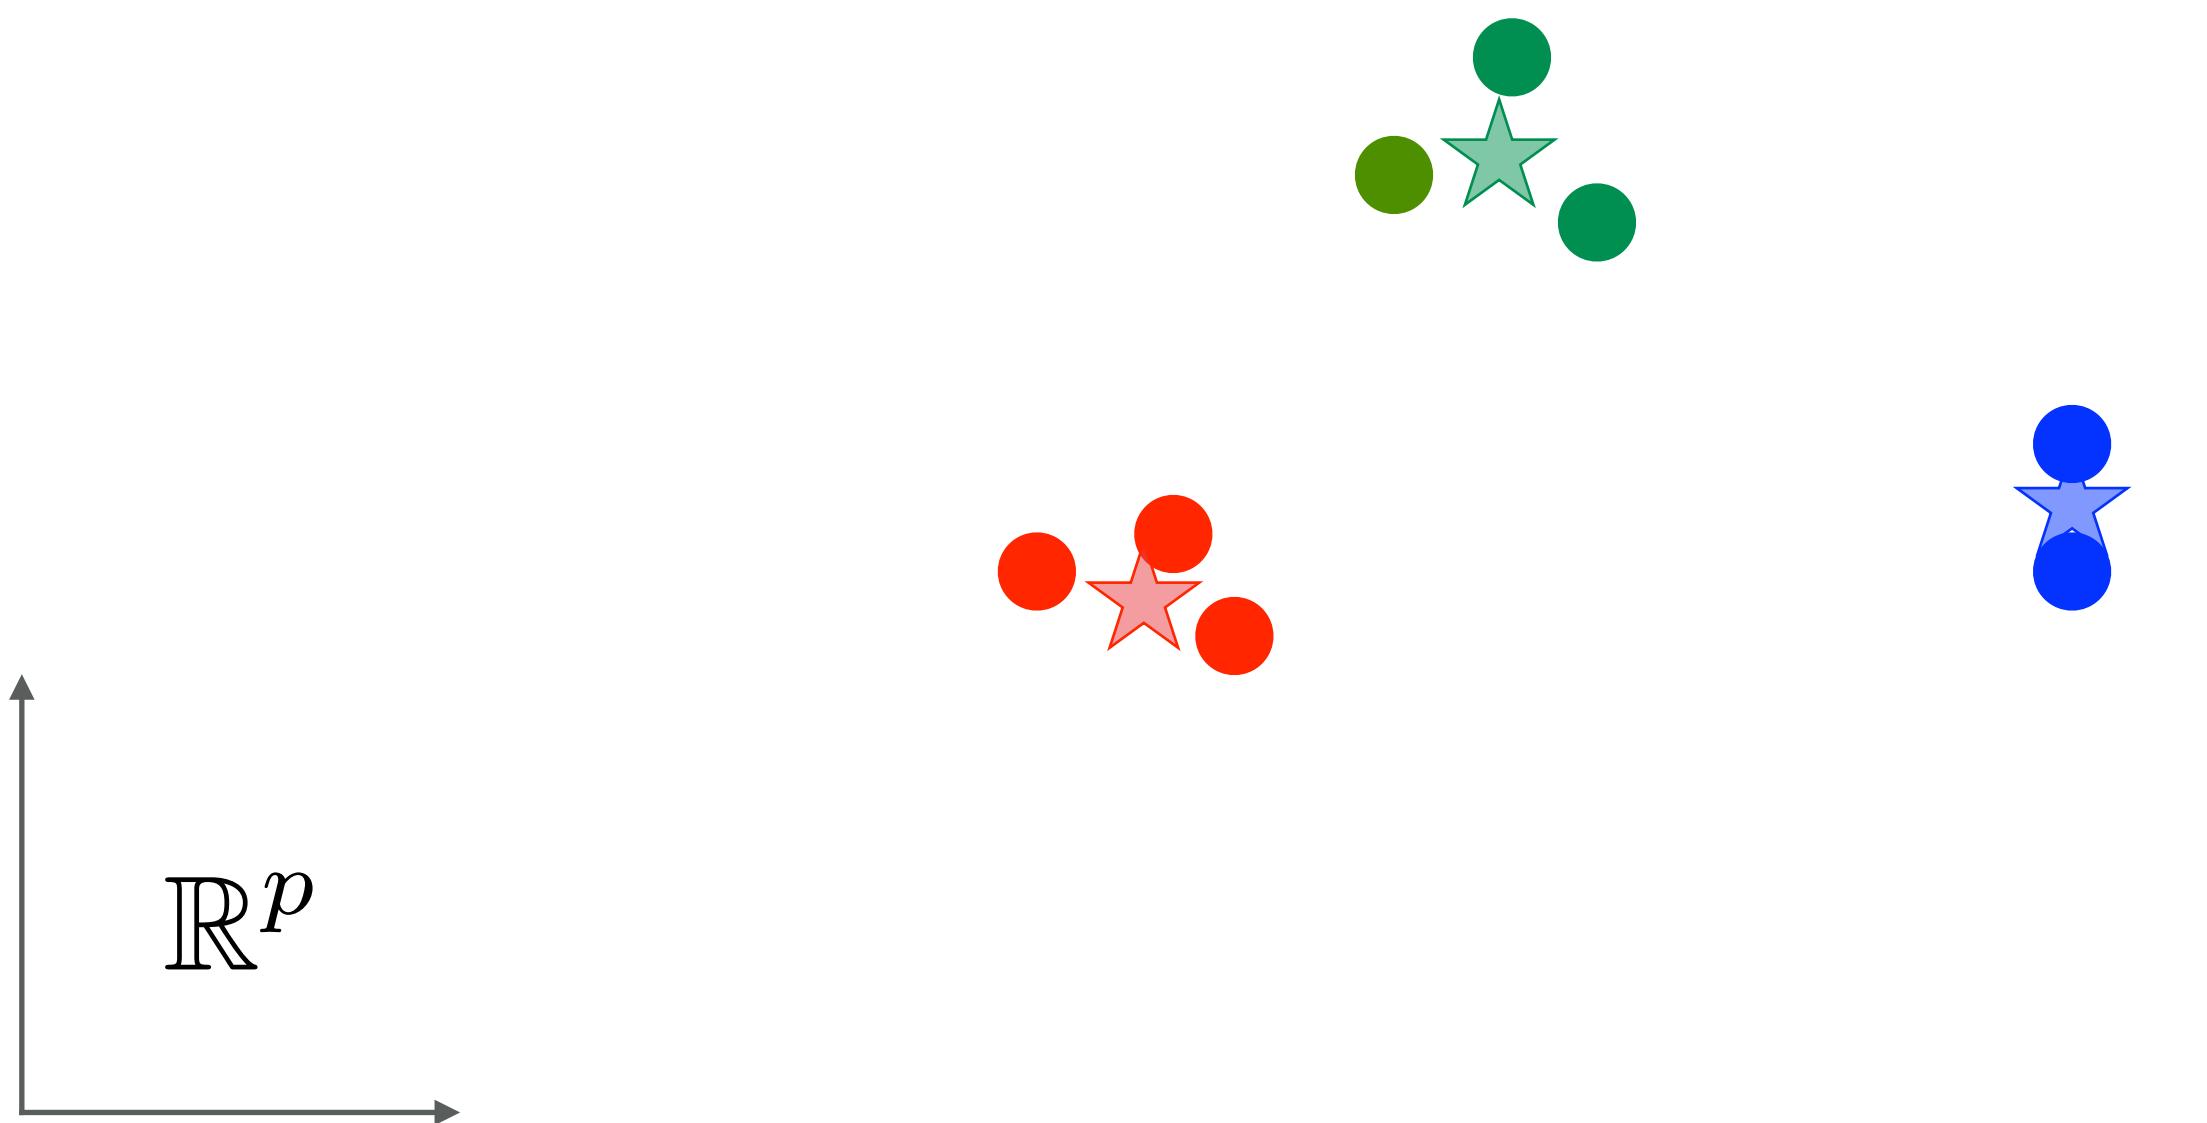
Step: update **assignments**



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Machine Learning 101: K-means clustering

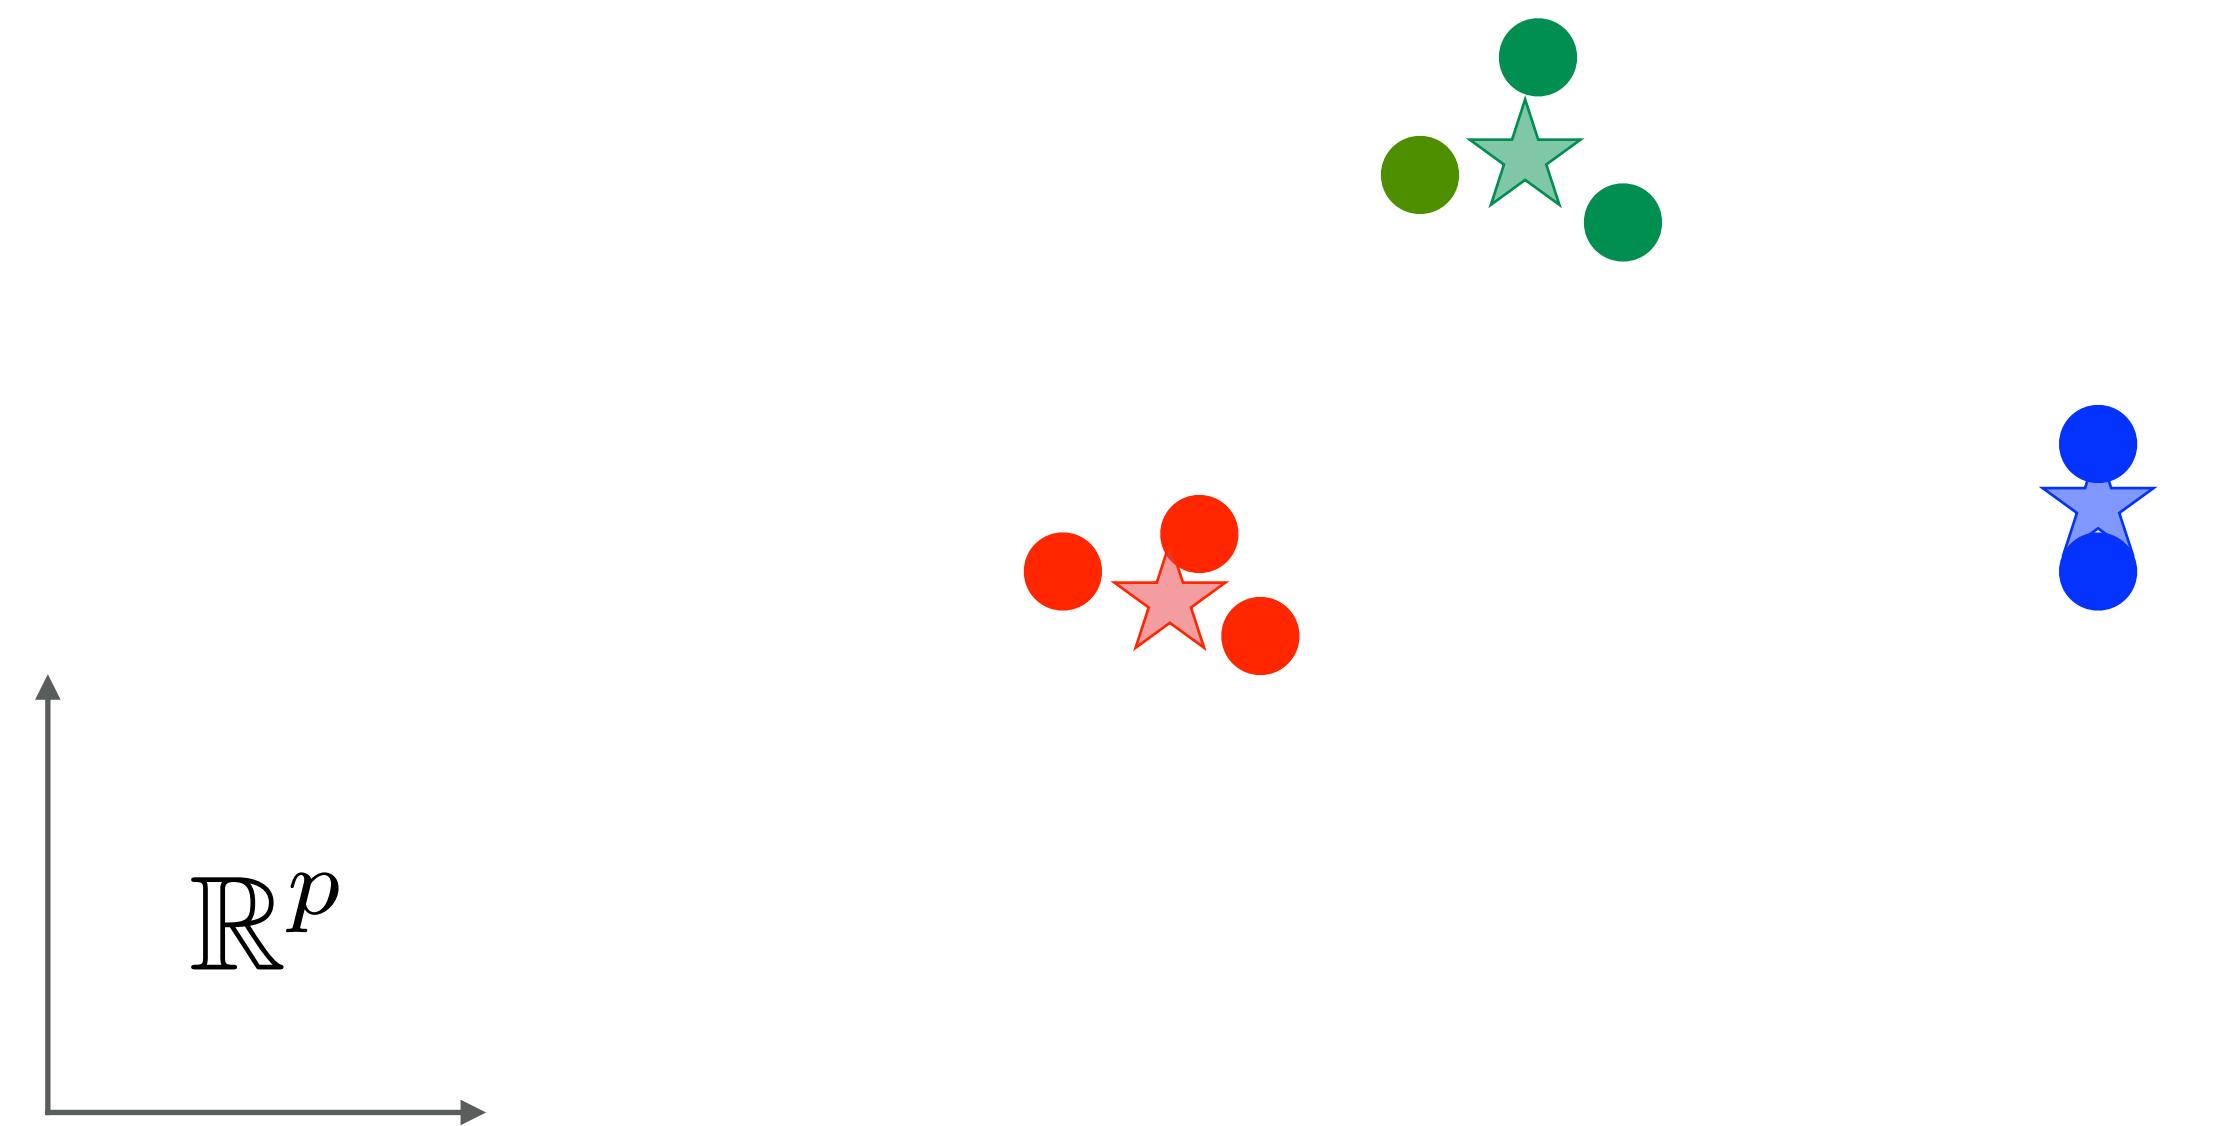
Step: recompute **cluster means**



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Machine Learning 101: K-means clustering

Step: recompute **cluster means**



* but it need not converge to the optimal solution!

Theorem:

K-means procedure (Lloyd's algorithm) converges*

Proof:

There are only a finite number of assignments, and we can never cycle since each iteration is an improvement

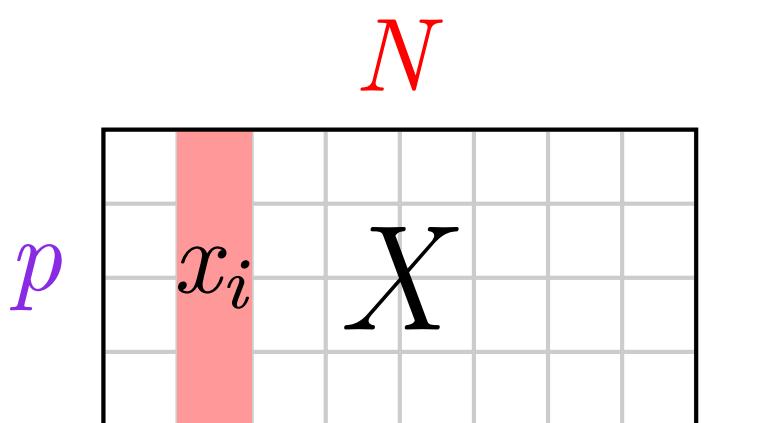
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

K-means clustering: complexity

Hard K-means / Lloyd's algorithm:

$$\{x_i\}_{i=1,\dots,\textcolor{red}{N}} \subset \mathbb{R}^{\textcolor{violet}{p}}$$

1. Update cluster centers $\mu_k \in \mathbb{R}^{\textcolor{violet}{p}}$ for $k = 1, \dots, K$
2. Update assignments $c_i \in \{1, \dots, K\}$ for $i = 1, \dots, \textcolor{red}{N}$



Let $X = (x_i)_{i=1,\dots,\textcolor{red}{N}}$

To update assignment of x_i :

- ▶ Compute $\|x_i - \mu_k\|$ (cost: $\textcolor{violet}{p}$ flops).
- ▶ For all k , and all i , this means $\mathcal{O}(KN\textcolor{violet}{p})$ cost.

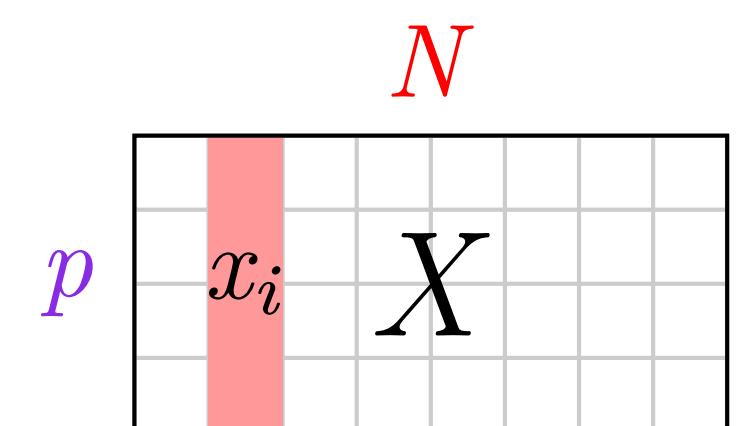
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

K-means clustering: complexity

Hard K-means / Lloyd's algorithm:

$$\{x_i\}_{i=1,\dots,\textcolor{red}{N}} \subset \mathbb{R}^{\textcolor{violet}{p}}$$

1. Update cluster centers $\mu_k \in \mathbb{R}^{\textcolor{violet}{p}}$ for $k = 1, \dots, K$
2. Update assignments $c_i \in \{1, \dots, K\}$ for $i = 1, \dots, \textcolor{red}{N}$



Let $X = (x_i)_{i=1,\dots,\textcolor{red}{N}}$

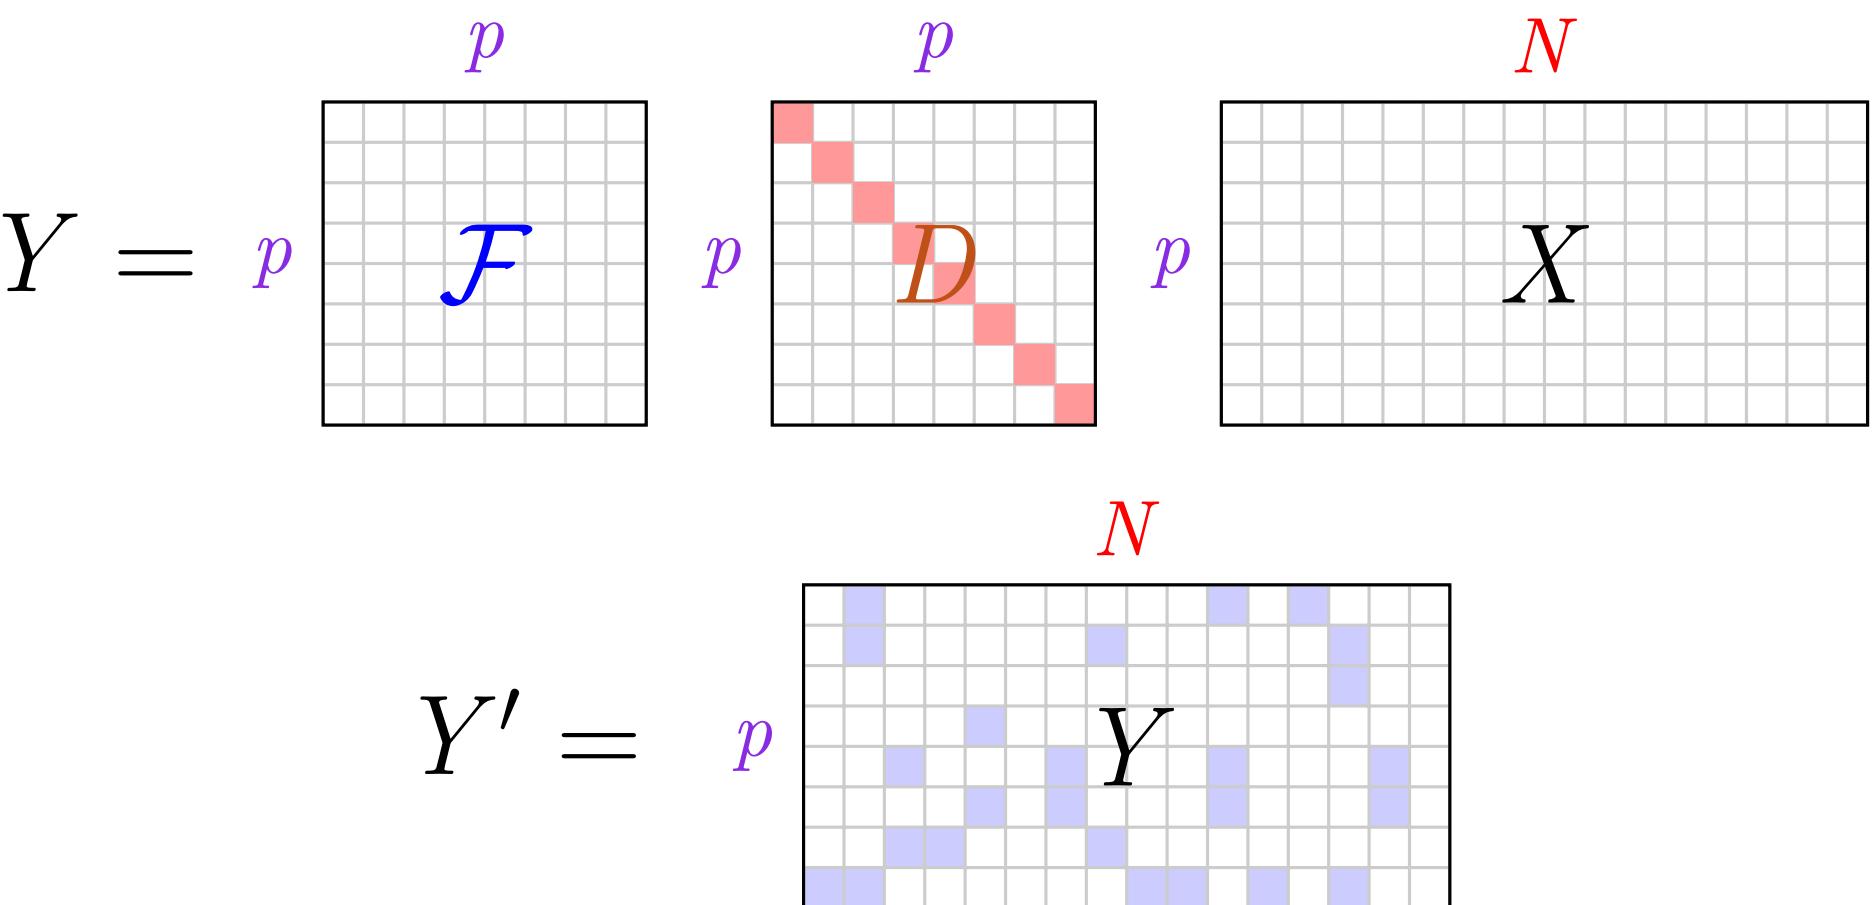
To update assignment of x_i :

- ▶ Compute $\|x_i - \mu_k\|$ (cost: $\textcolor{violet}{p}$ flops).
- ▶ For all k , and all i , this means $\mathcal{O}(KN\textcolor{violet}{p})$ cost.

Because we're estimating the mean, we can use the theorems discussed earlier!

Our sampling is suitable for **streaming** data

Our idea: apply “preconditioned” entry-wise subsampling



Given a true mean $\bar{x} = \frac{1}{\textcolor{red}{N}} \sum_{i=1}^{\textcolor{red}{N}} x_i$, and our estimate of it \hat{x} from sampled data,

Theorem (Pourkamali-Anaraki & B., *IEEE Trans. Info Theory* 2017)

$\mathbb{E} \hat{x} = \bar{x}$ and $\|\hat{x} - \bar{x}\|_\infty \leq \textcolor{violet}{t}$ with probability greater than

$$1 - 2\textcolor{violet}{p} \exp\left(\frac{-\textcolor{red}{N}\gamma\textcolor{violet}{t}^2/2}{\|X\|_{\max\text{-row}}^2 + \textcolor{violet}{t}/3\|X\|_{\max\text{-entry}}}\right)$$

where $\gamma = p_{\text{small}}/\textcolor{violet}{p}$

(simplifying to $p_{\text{small}} \ll \textcolor{violet}{p} \ll \textcolor{red}{N}$)

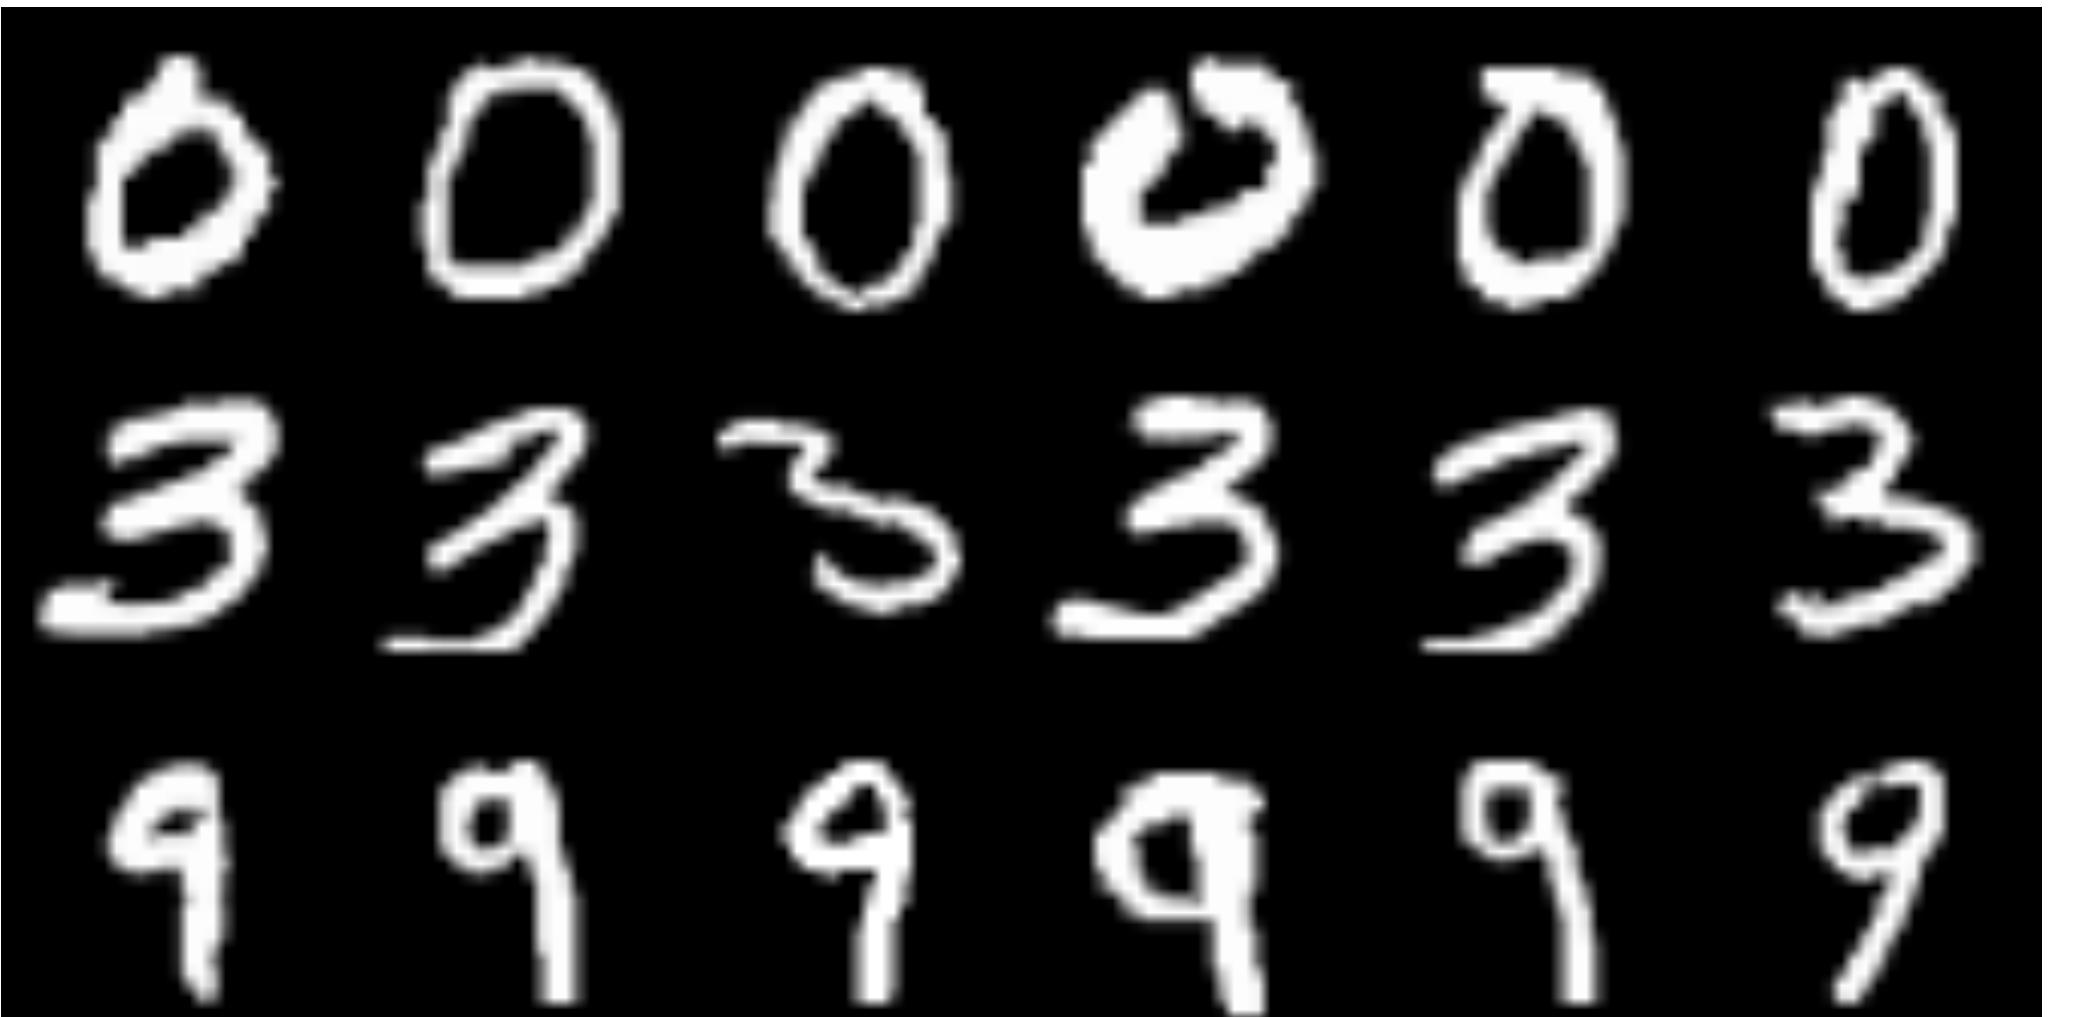


1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

MNIST handwritten digits

Goal: recover cluster centers, using K-means clustering algorithm

Data: $N = 21,002$ examples of 28×28 pixel images ($p = 784$)



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Experiments: MNIST handwritten digits



(a) true cluster centers



(b) K-means, many passes

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Experiments: MNIST handwritten digits



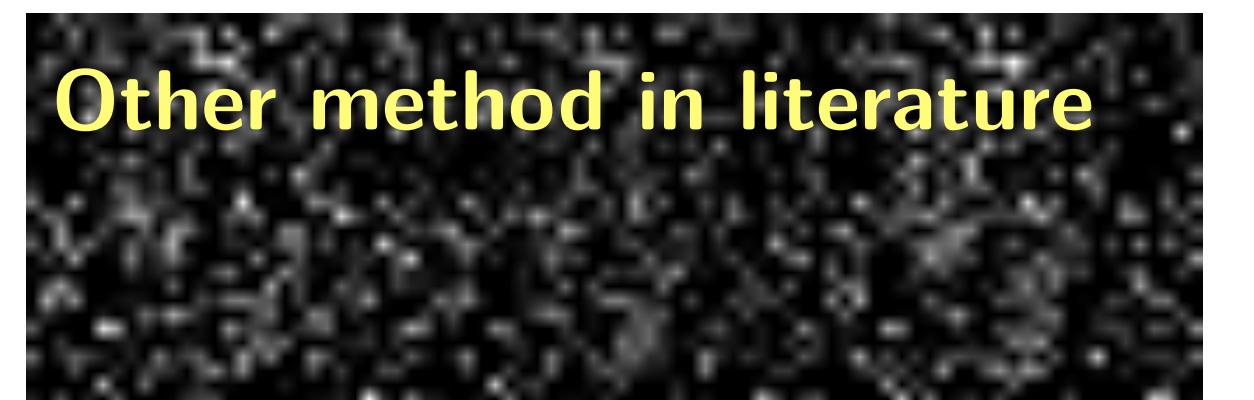
(a) true cluster centers



(b) K-means, many passes



(c) sparsified K-means, 1 pass,
no preconditioning



(g) feature extraction, 1 pass



(h) feature selection, 3 passes

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Experiments: MNIST handwritten digits



(a) true cluster centers



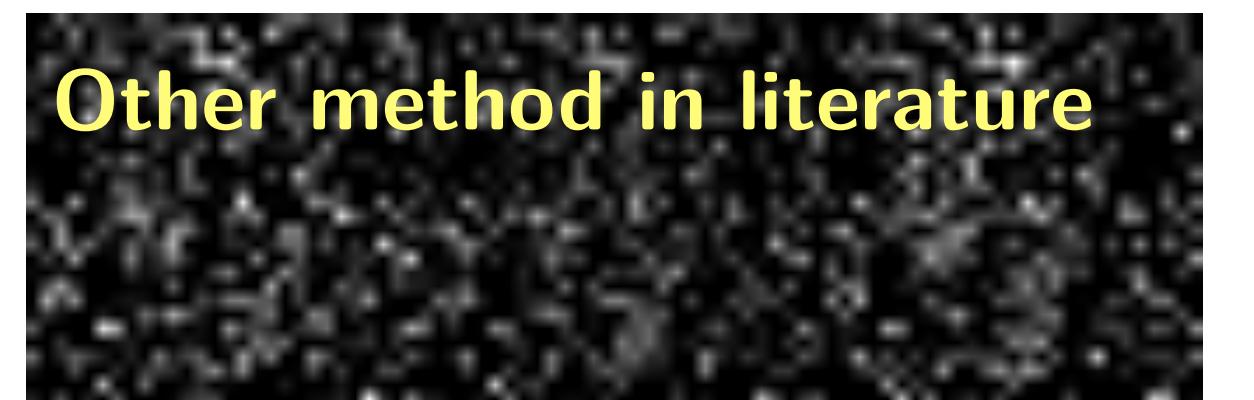
(b) K-means, many passes



(c) sparsified K-means, 1 pass,
no preconditioning



(d) sparsified K-means, 1 pass,
preconditioned



(g) feature extraction, 1 pass

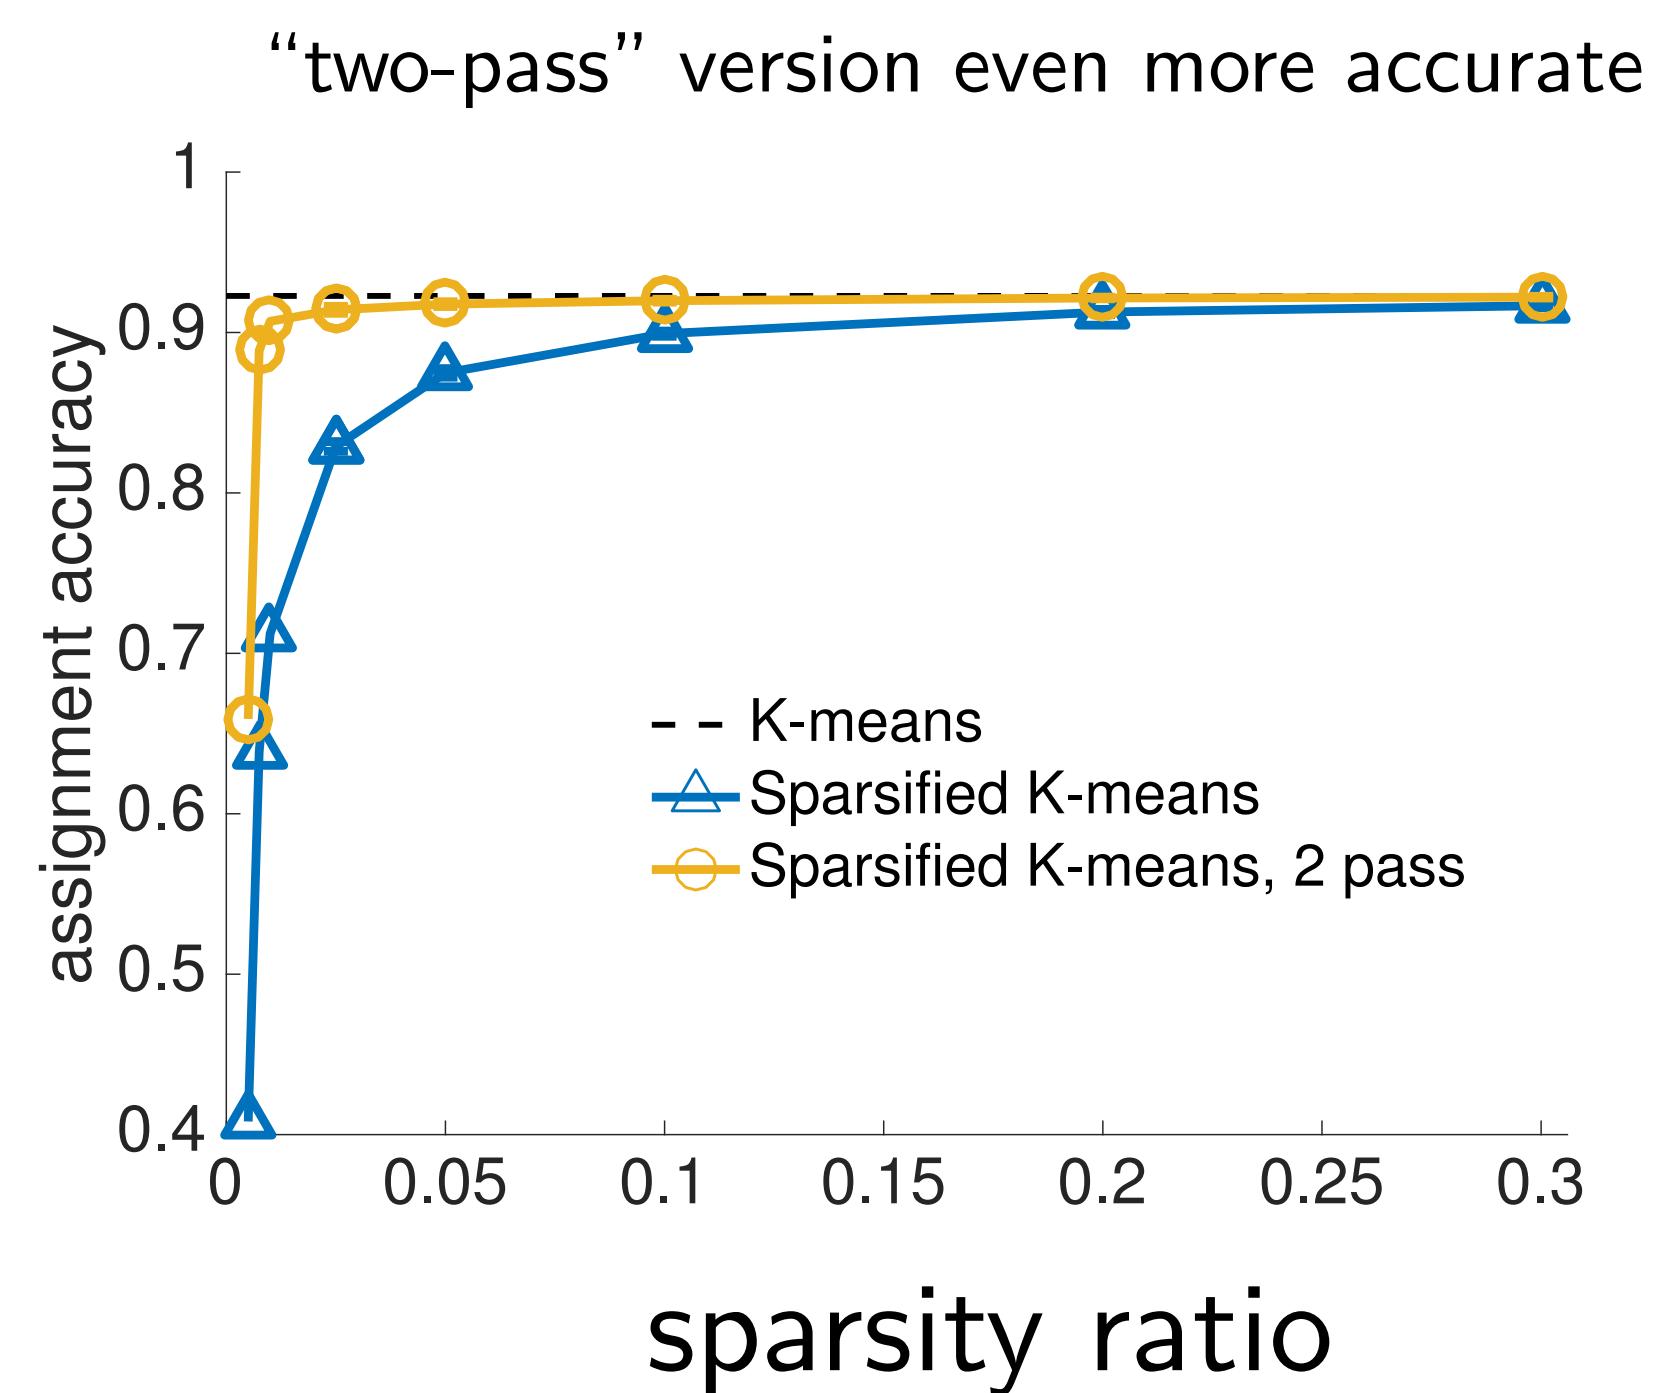
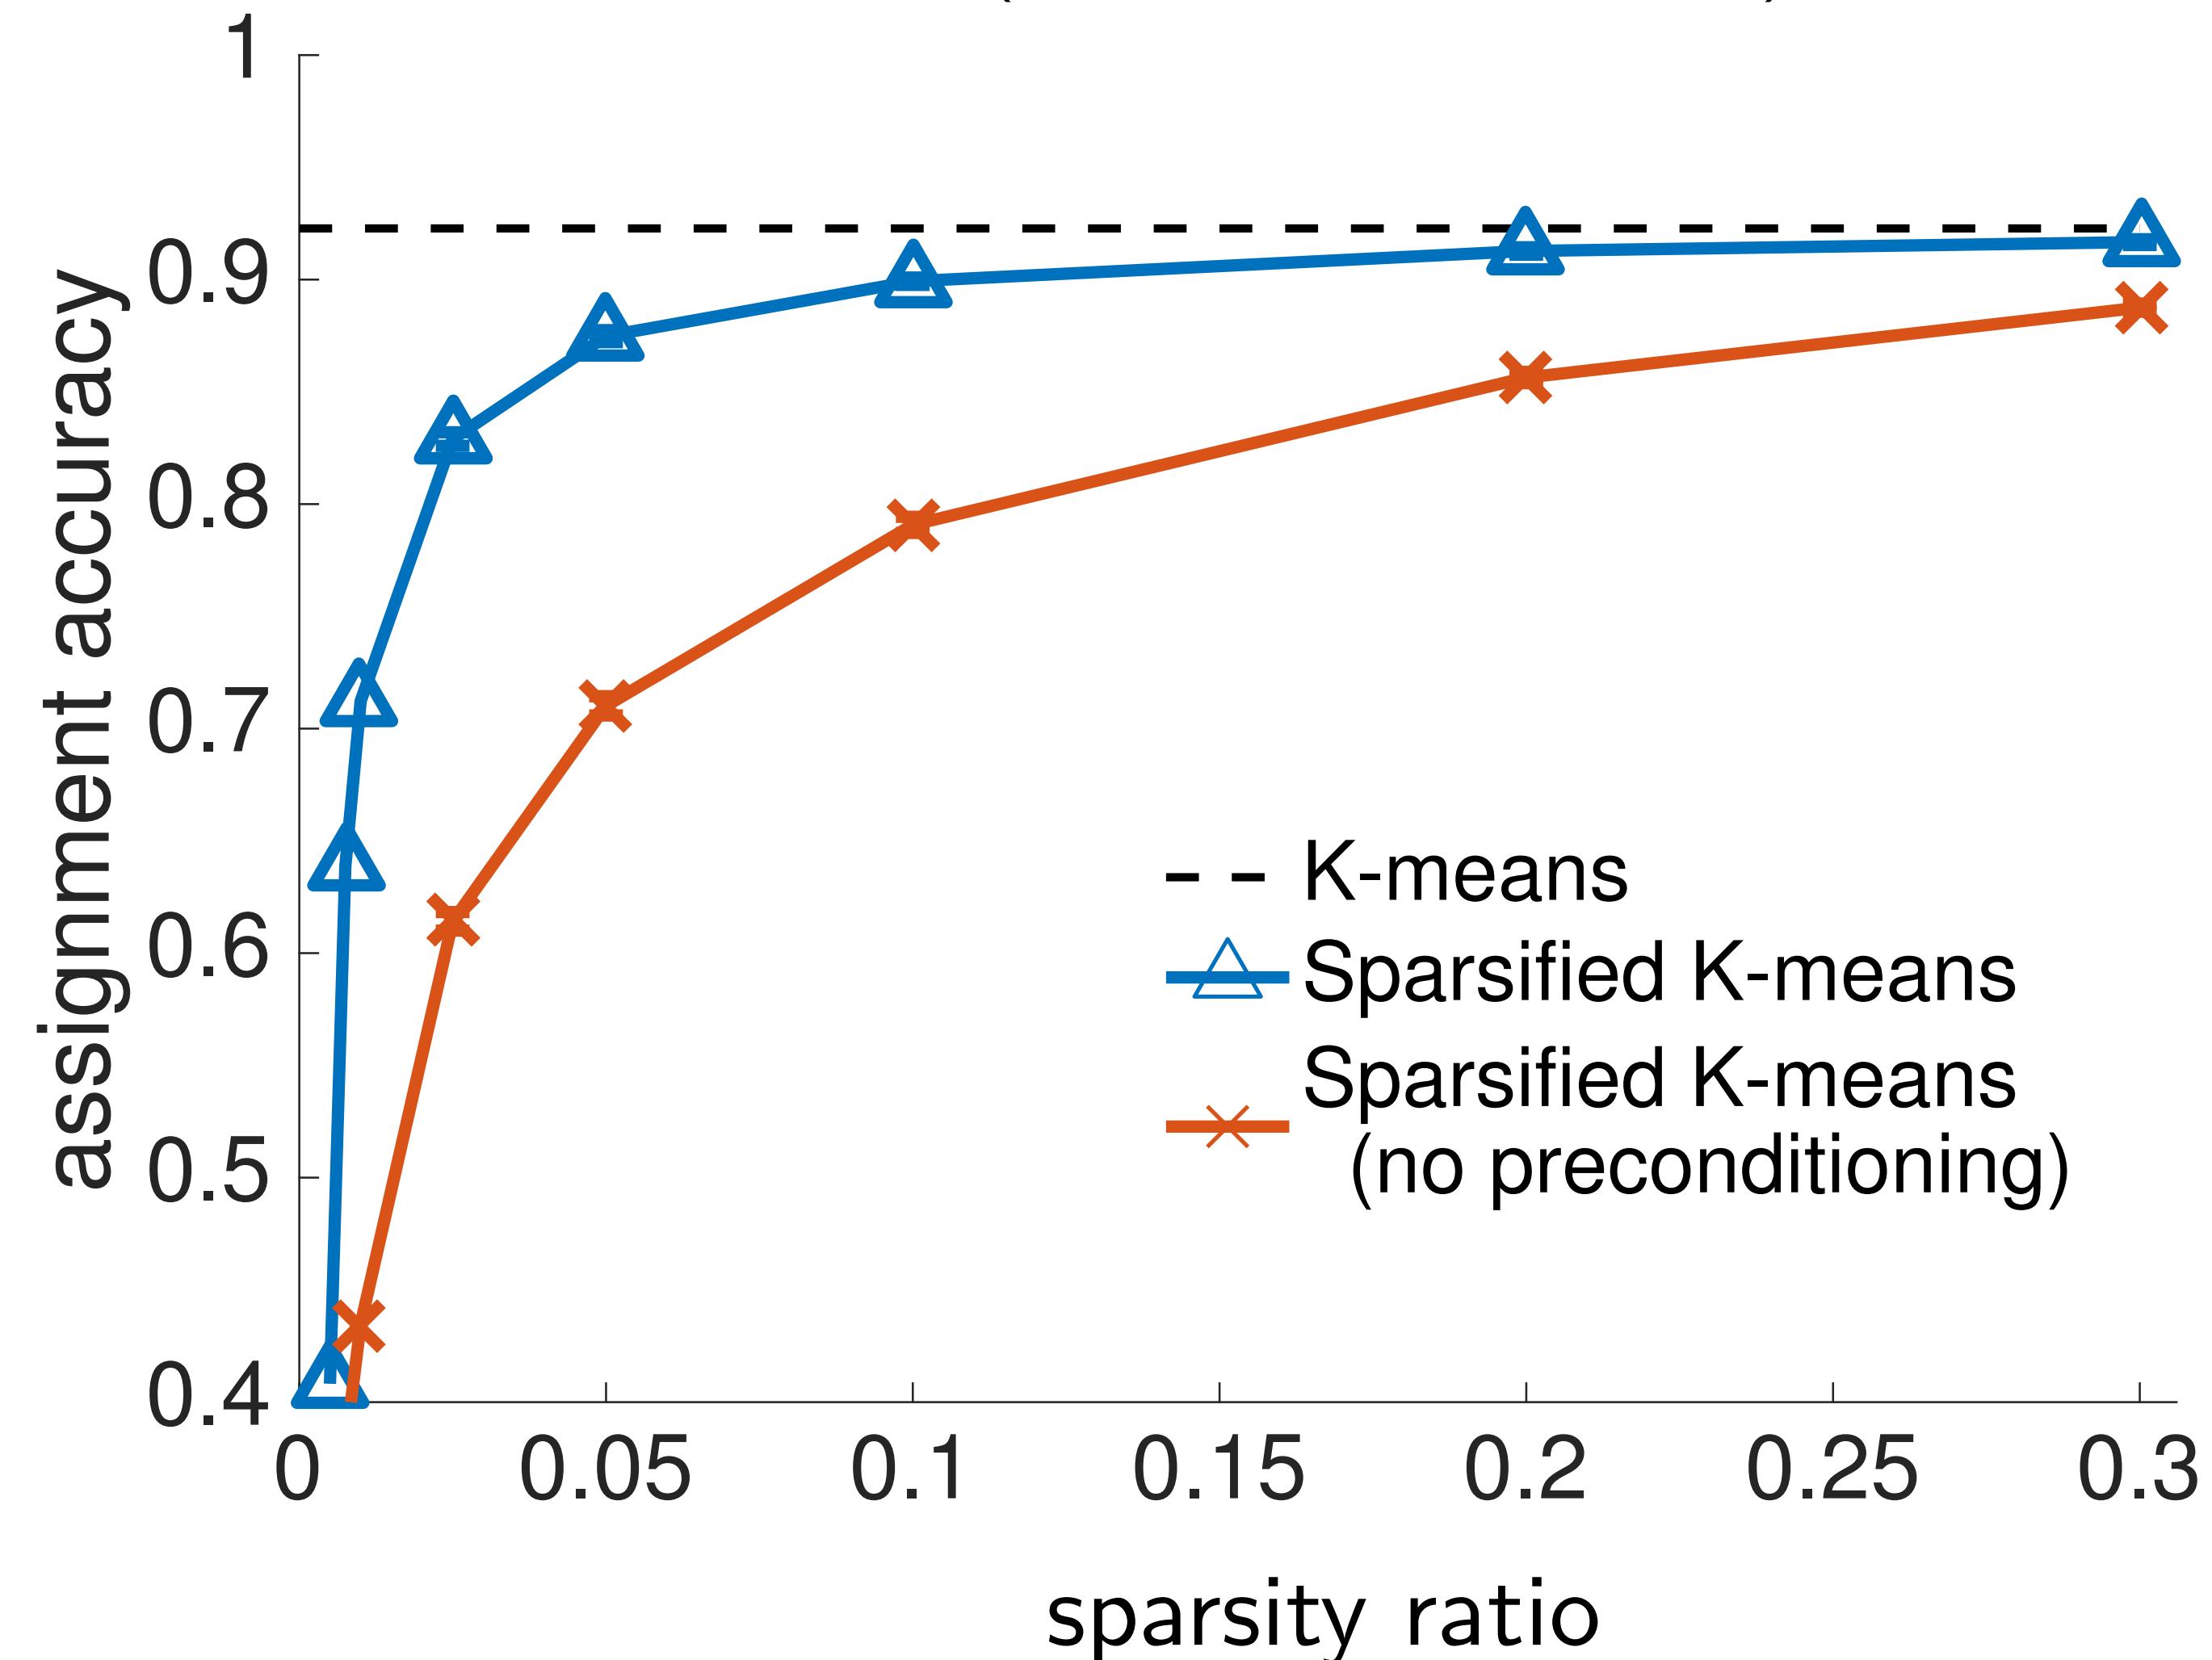


(h) feature selection, 3 passes

Accuracy on MNIST

Now look at **classification accuracy** (averaged over 50 trials)

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. **K-means clustering**
 - c. Tensor factorizations
 - d. Gradient-free optimization

Speed on “Infinite MNIST”

“Infinite MNIST”, $p = 784$, but now $N = 9,631,605$ instead of 21,002

56 GB data, split into 58 chunks to keep RAM usage around 1 GB

Algorithm	Time to find assignments		Time to update all centers		Combined time	
	Absolute	Speedup	Absolute	Speedup	Absolute	Speedup
K-means	130.0s	1×	150.8s	1×	280.8s	1×
Sparsified	1.3s	100×	5.7s	26.4×	7.0s	40.1×

at 20x undersampling; accuracy is 89% (vs 92% for K-means)

OUTLINE

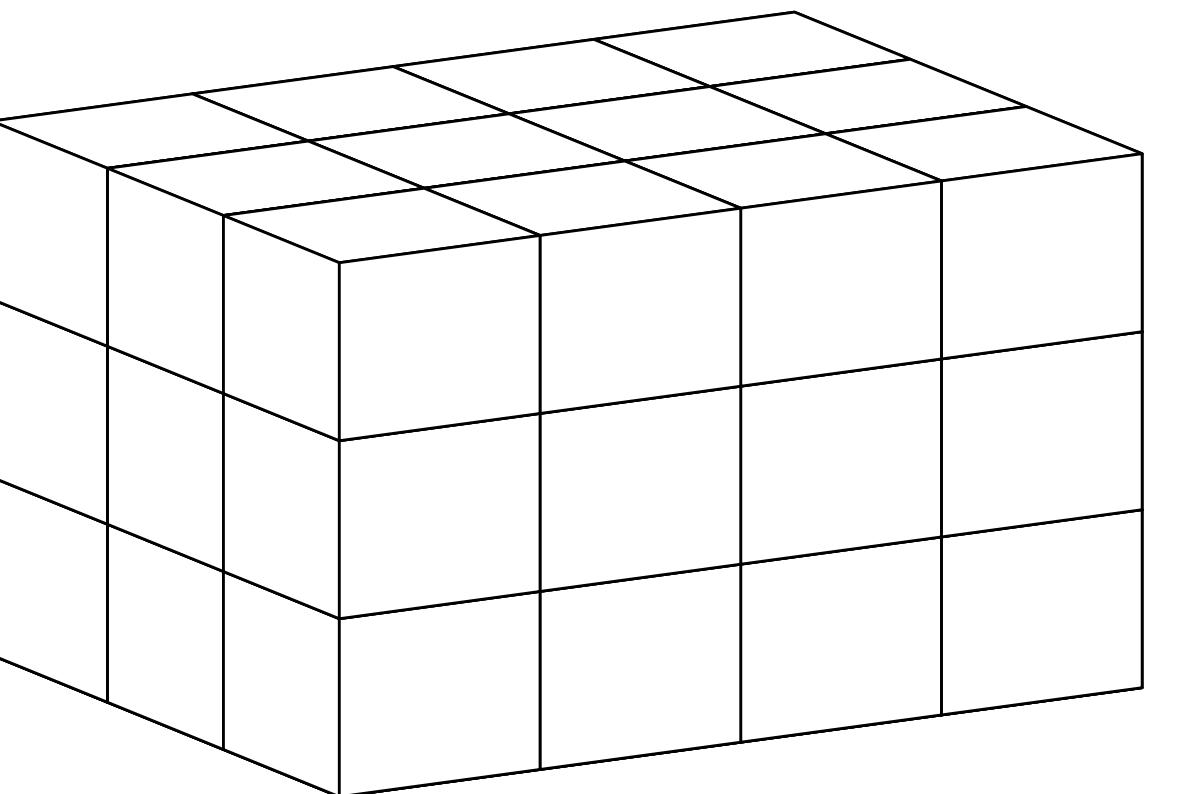
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Tensor Background: what is a tensor?



Change in notation temporarily



- ▶ A *tensor* $\mathcal{X} \in \mathbb{R}^{I_1 \times I_2 \times \dots \times I_N}$ is an array of dimension N , also called an N -way *tensor* or *order N tensor*.
- ▶ A matrix $\mathbf{X} \in \mathbb{R}^{I_1 \times I_2}$ is a 2-way tensor.
- ▶ A vector $\mathbf{x} \in \mathbb{R}^{I_1}$ is a 1-way tensor.
- ▶ A scalar $x \in \mathbb{R}$ is a 0-way tensor.

Examples of tensors

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
 2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

For example, [Kolda & Sun, '08] consider these datasets:

- ▶ The **Enron dataset** has size $1K \times 1K \times 1.1K \times 200$ and consists of elements of the form (user,user,keyword,day). Element (i, j, k, l) is 1 if user i sent an email to user j with keyword k on day l , 0 otherwise.
 - ▶ The **DBLP dataset** has size $5K \times 1K \times 1K$ and consists of elements of the form (author,conference,keyword). Element (i, j, k) is 1 if author i published a paper at conference j containing the keyword k in the title, 0 otherwise.

These are examples of **sparse** tensors

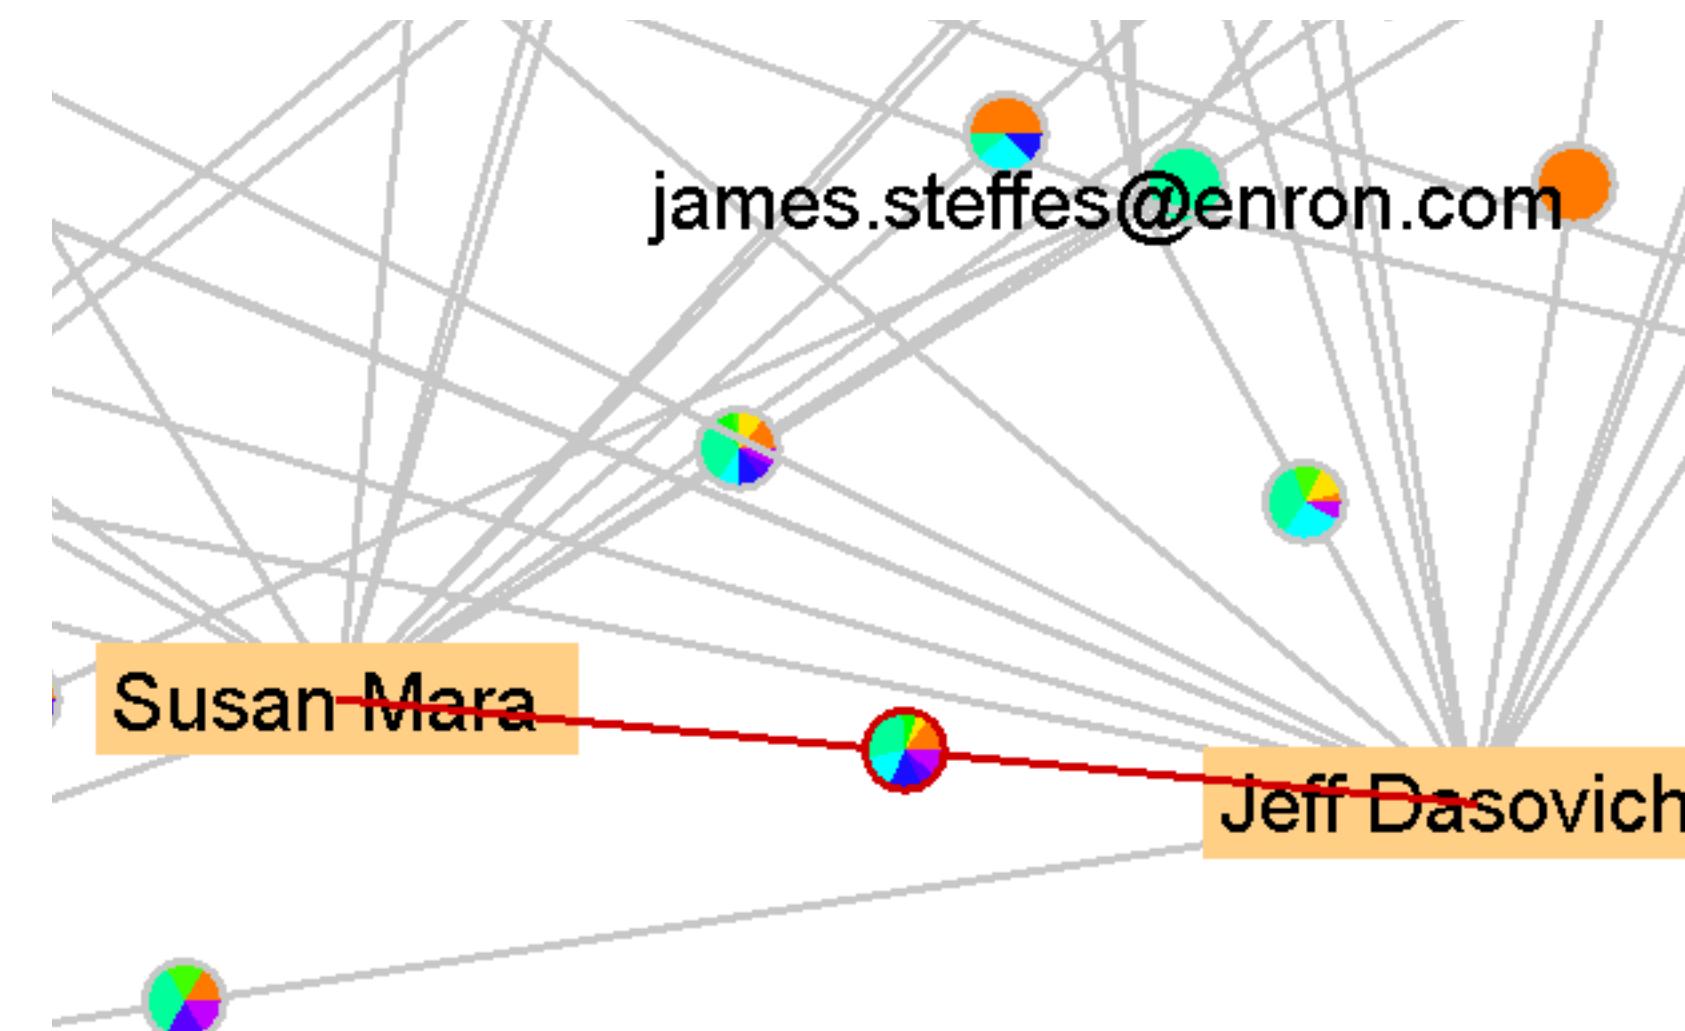


Image: <https://homes.cs.washington.edu/~jheer/projects/enron/>

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Examples of tensors

For example, [Kolda & Sun, '08] consider these datasets:

- ▶ **The Enron dataset** has size $1K \times 1K \times 1.1K \times 1K$ and consists of elements of the form (user, user, keyword, day). Element (i, j, k, l) is 1 if user i sent an email to user j containing the keyword k on day l , 0 otherwise.
- ▶ **The DBLP dataset** has size $5K \times 1K \times 1K$ and consists of elements of the form (author, conference, keyword). Element (i, j, k) is 1 if author i published a paper at conference j containing the keyword k in the title, 0 otherwise.

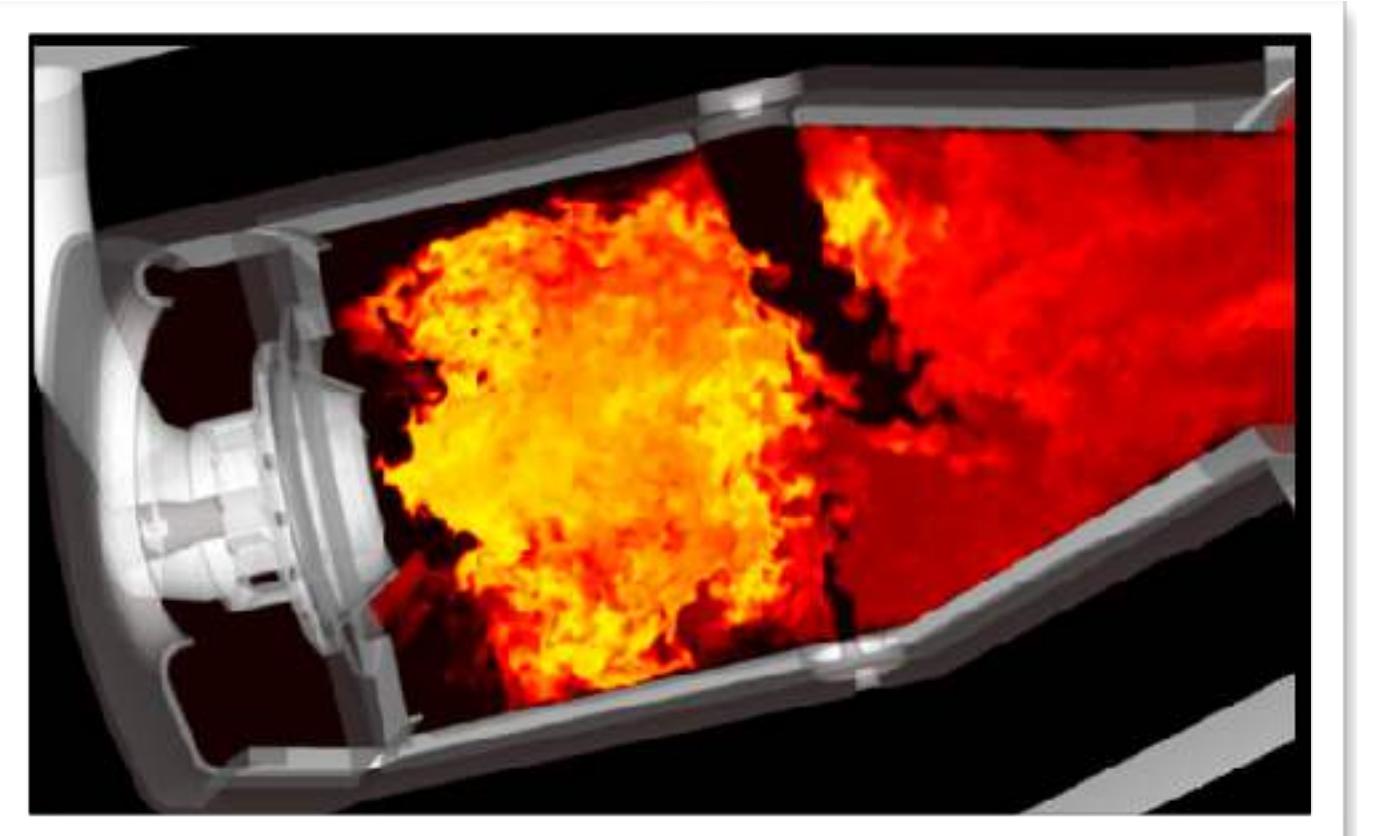
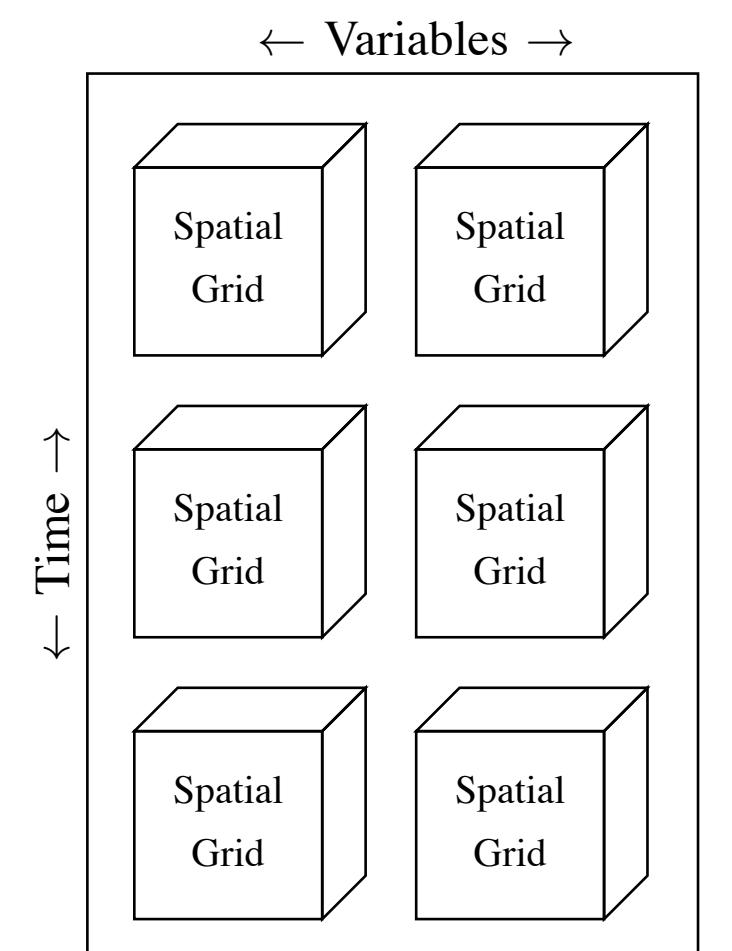


Image: Center for Turbulence Research (Stanford)

[Austin et al., '16] consider tensor data produced by high-fidelity combustion simulations with the following properties:

- ▶ 3-dimensional spatial grid with 512 points per dimension
- ▶ 64 variables are tracked per grid point
- ▶ 128 time steps $512 \times 512 \times 512 \times 64 \times 128$ (8 TB)



This is a **dense tensor**

... we'll want methods that work with both sparse and dense tensors

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Examples of tensors

For example, [Kolda & Sun, '08] consider these datasets:

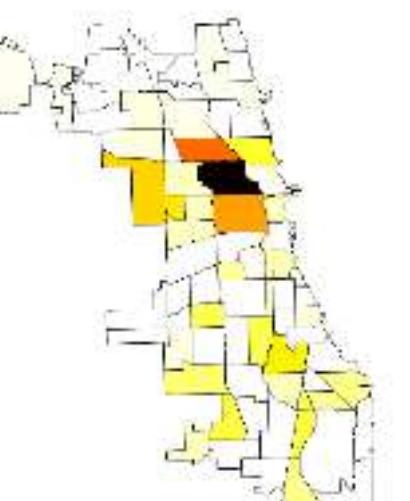
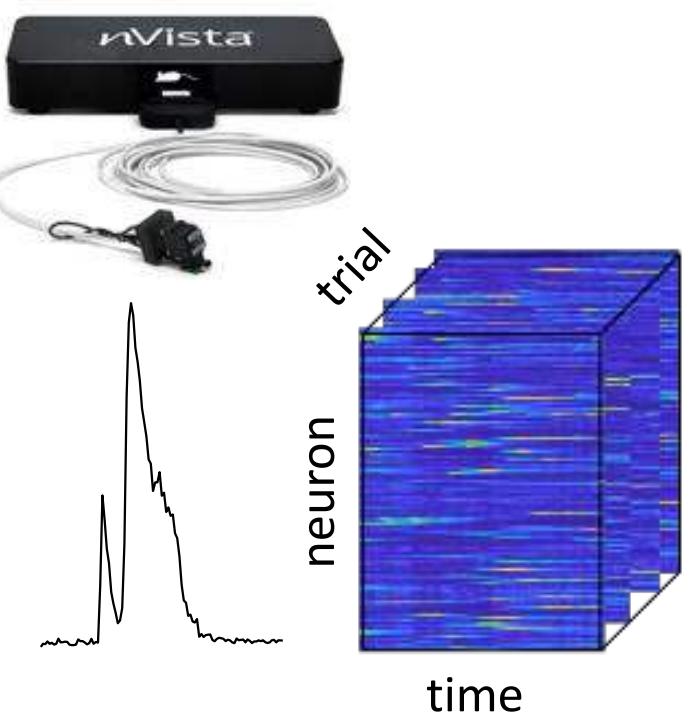
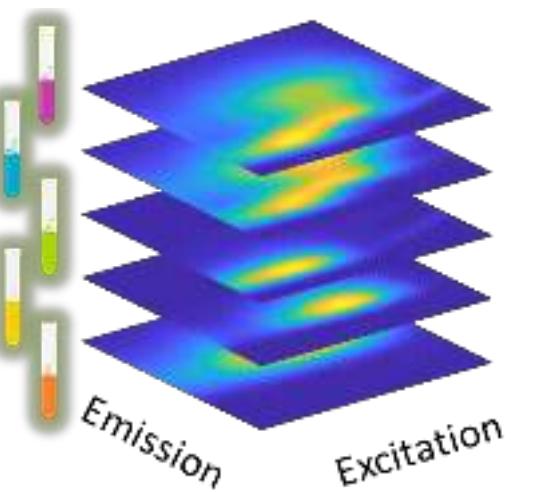
- ▶ **The Enron dataset** has size $1K \times 1K \times 1.1K \times 200$ and consists of elements of the form (user,user,keyword,day). Element (i, j, k, l) is 1 if user i sent an email to user j with keyword k on day l , 0 otherwise.
- ▶ **The DBLP dataset** has size $5K \times 1K \times 1K$ and consists of elements of the form (author,conference,keyword). Element (i, j, k) is 1 if author i published a paper at conference j containing the keyword k in the title, 0 otherwise.

[Austin et al., '16] consider tensor data produced by high-fidelity combustion simulations with the following properties:

- ▶ 3-dimensional spatial grid with 512 points per dimension
- ▶ 64 variables are tracked per grid point
- ▶ 128 time steps $512 \times 512 \times 512 \times 64 \times 128$ (8 TB)

Others:

- ▶ Chemometrics (emission \times excitation \times samples, for fluorescence spectroscopy)
- ▶ Neuroscience (neuron \times time \times trial/stimulus, for calcium imaging)
- ▶ Criminology (day \times hour \times location \times crime, e.g., Chicago crime dataset)



Images: Tamara Kolda

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Examples of tensors

For example, [Kolda & Sun, '08] consider these datasets:

- ▶ **The Enron dataset**

consists of
Element (i, j)
keyword k

- ▶ **The DBLP dataset**

elements of
 (i, j, k) is
containing

[Austin et al.,
combustion sim]

- ▶ 3-dimensional
- ▶ 64 variables
- ▶ 128 time steps

One-pass methods

All these tensors have a **time or sample index**

As time/samples increase, we don't want to store old ones

i.e., want a **streaming/one-pass method**

THIS IS A BIG MOTIVATION FOR SKETCHING

$512 \times 512 \times 512 \times 04 \times 128$ (8 TB)

Others:

- ▶ Chemometrics (emission \times excitation \times samples, for fluorescence spectroscopy)
- ▶ Neuroscience (neuron \times time \times trial/stimulus, for calcium imaging)
- ▶ Criminology (day \times hour \times location \times crime, e.g., Chicago crime dataset)

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Tensor operations 1: matricization

Matricization (i.e., flattening) turns a tensor \mathcal{X} into a matrix $\mathbf{X}_{(n)}$

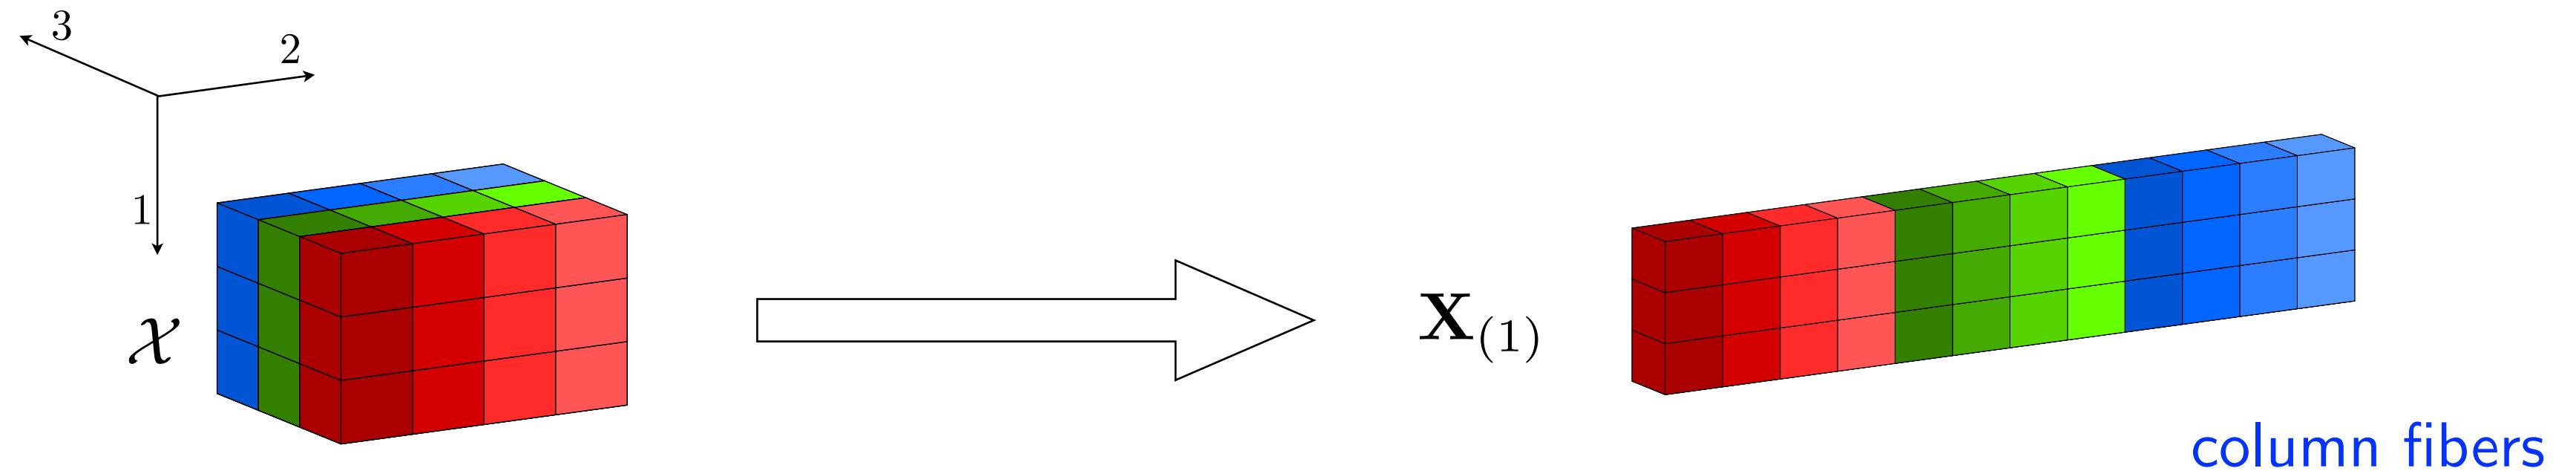
Same idea as turning a matrix (e.g., an image) into a vector

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Tensor operations 1: matricization

Matricization (i.e., flattening) turns a tensor \mathcal{X} into a matrix $\mathbf{X}_{(n)}$

Same idea as turning a matrix (e.g., an image) into a vector

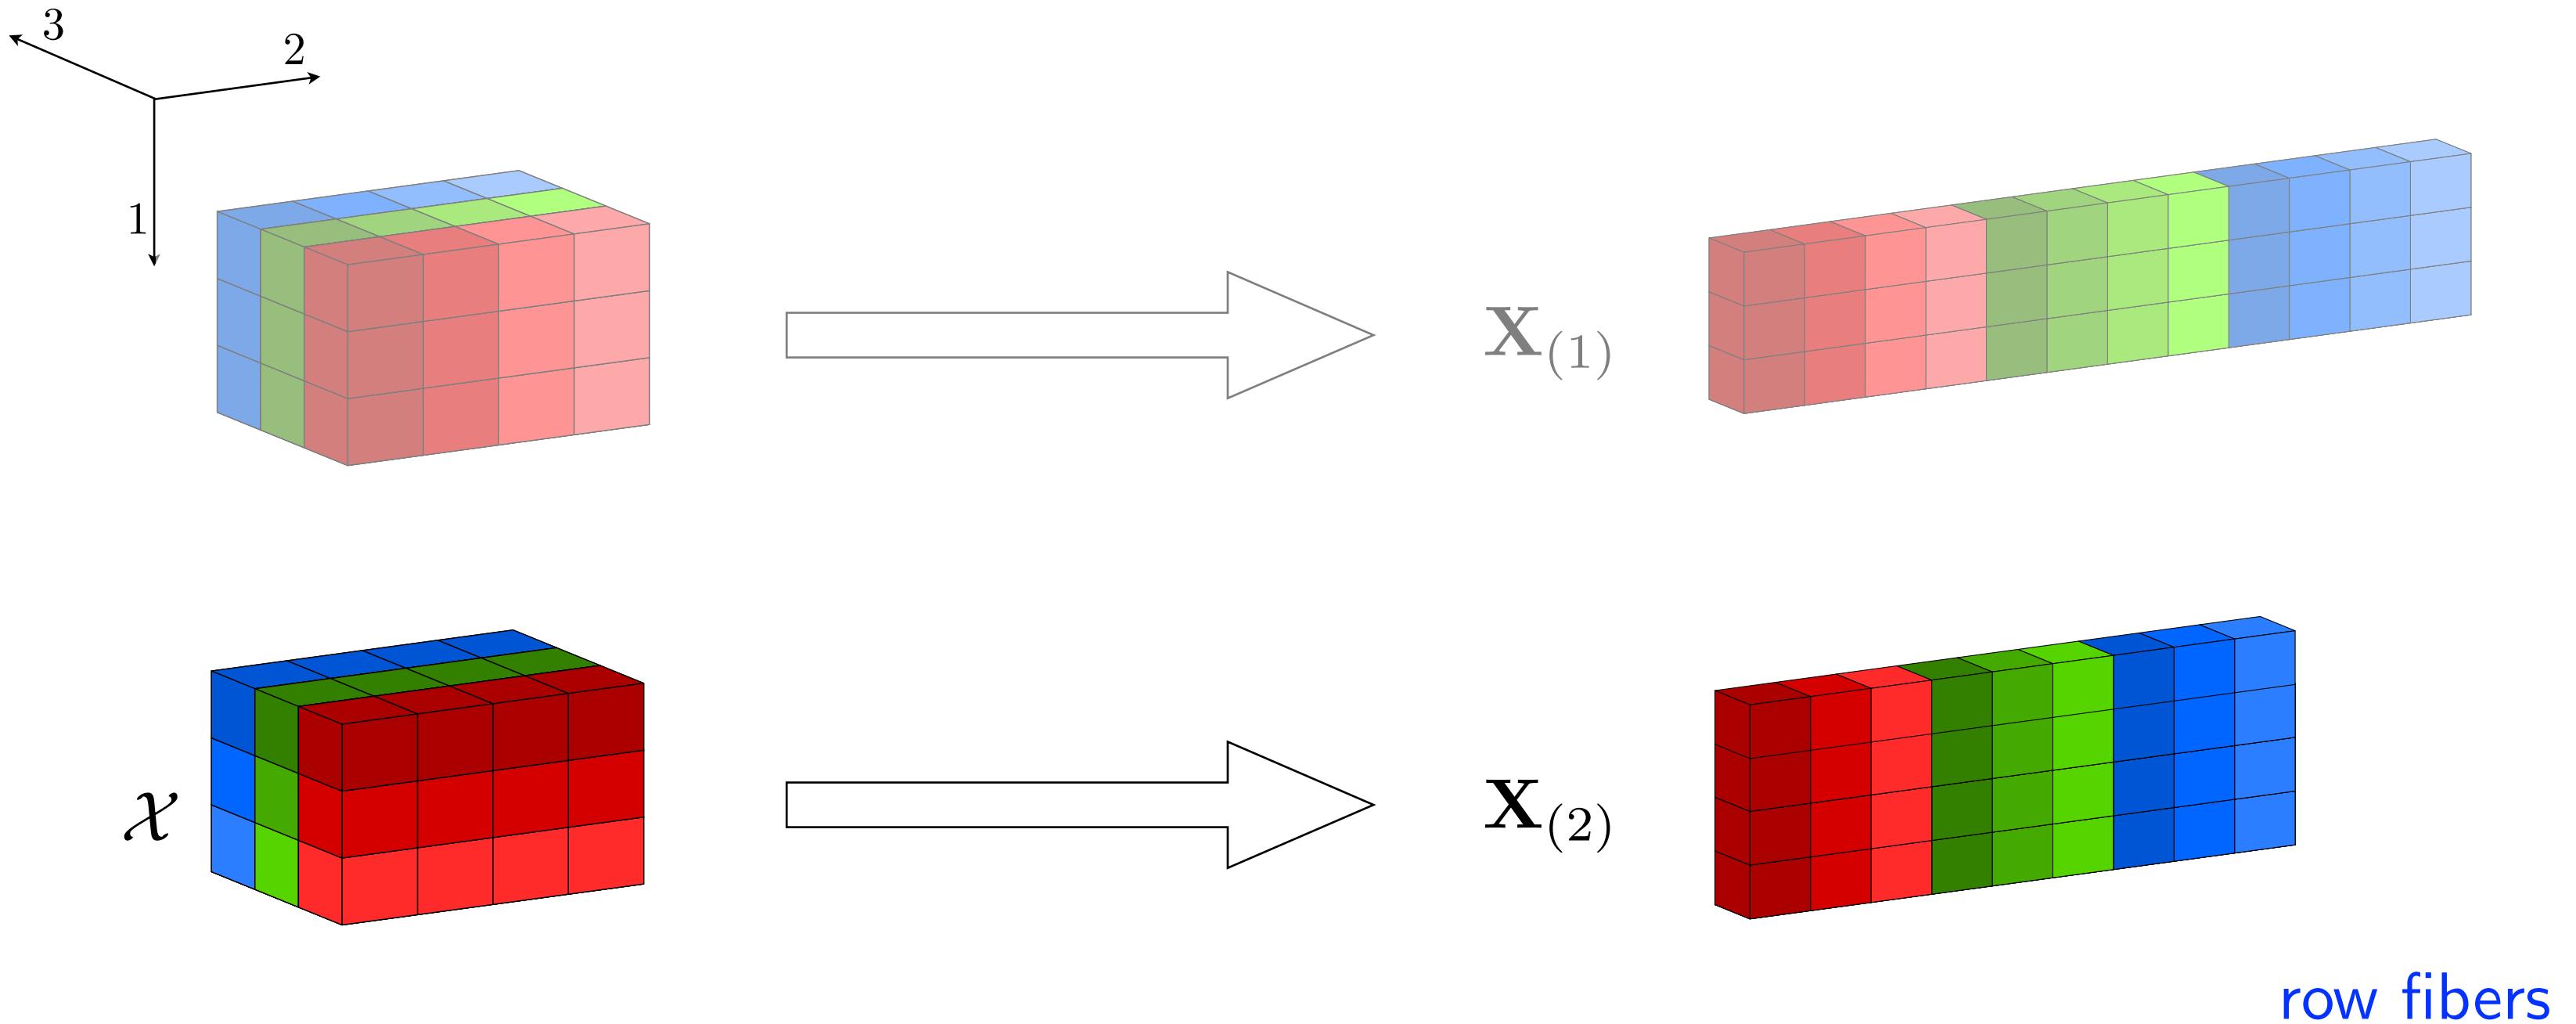


1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Tensor operations 1: matricization

Matricization (i.e., flattening) turns a tensor \mathcal{X} into a matrix $\mathbf{X}_{(n)}$

Same idea as turning a matrix (e.g., an image) into a vector

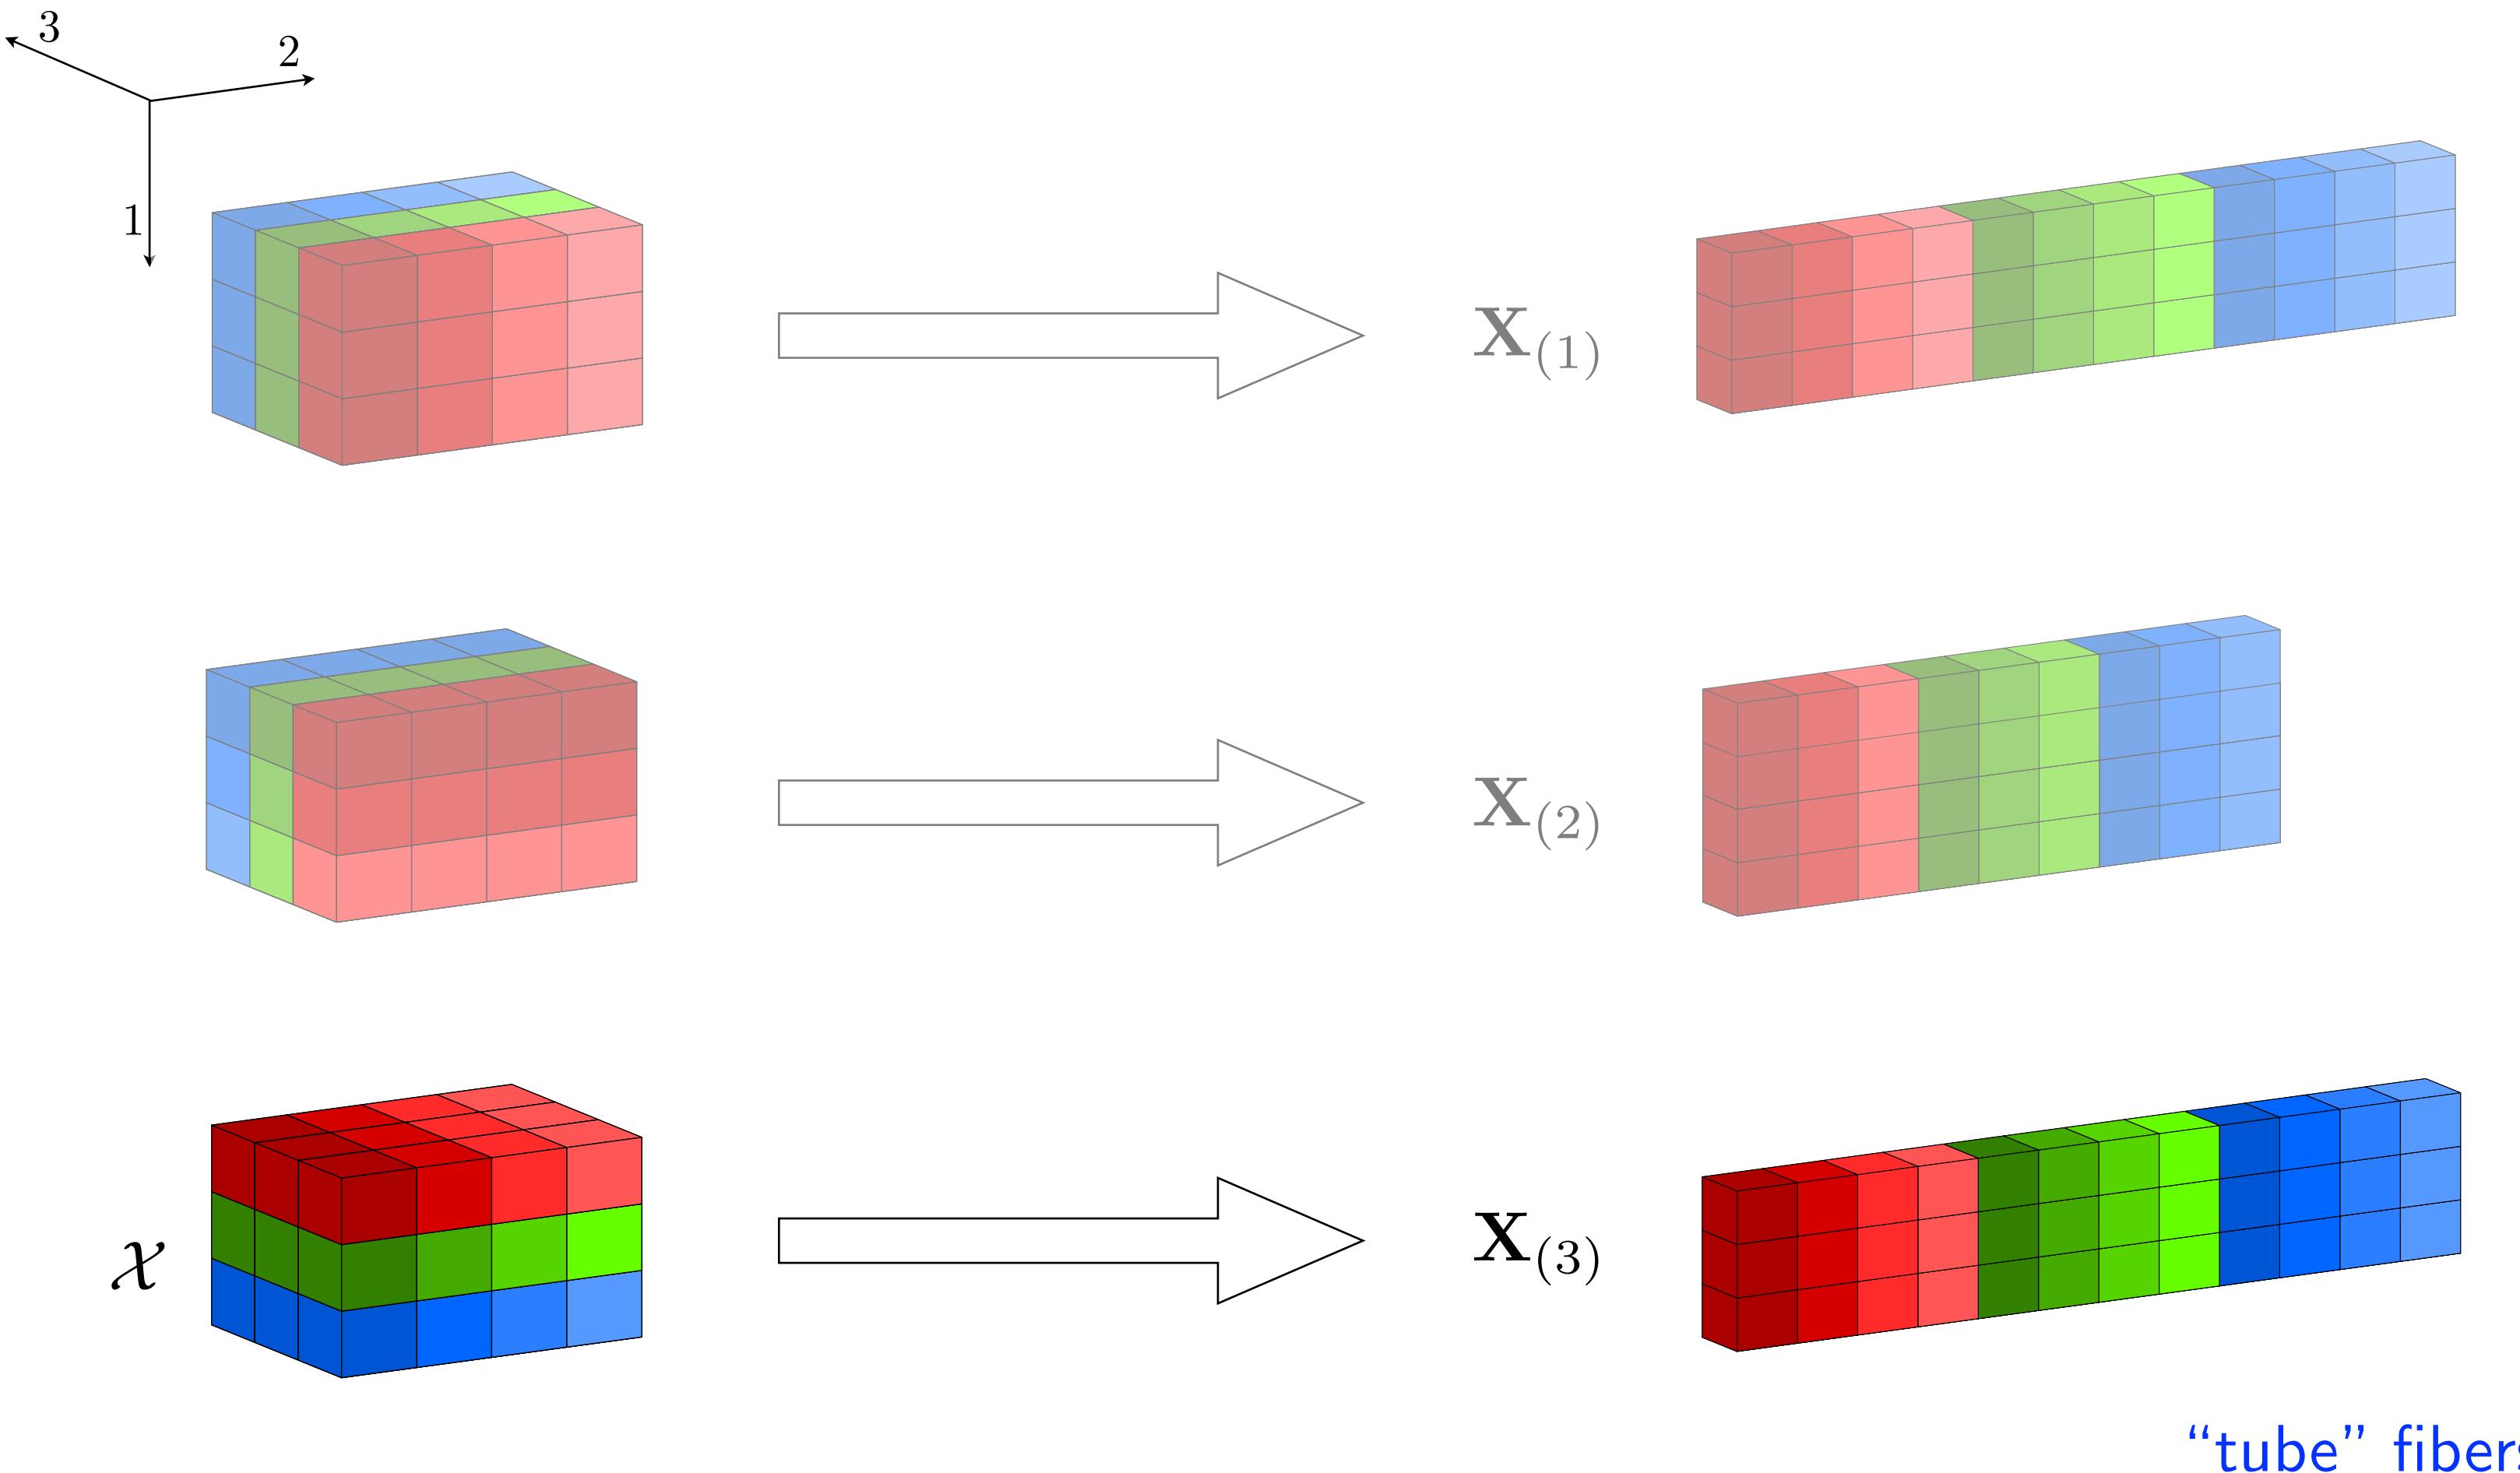


1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Tensor operations 1: matricization

Matricization (i.e., flattening) turns a tensor \mathcal{X} into a matrix $\mathbf{X}_{(n)}$

Same idea as turning a matrix (e.g., an image) into a vector



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Tensor operations 2: tensor-times-matrix

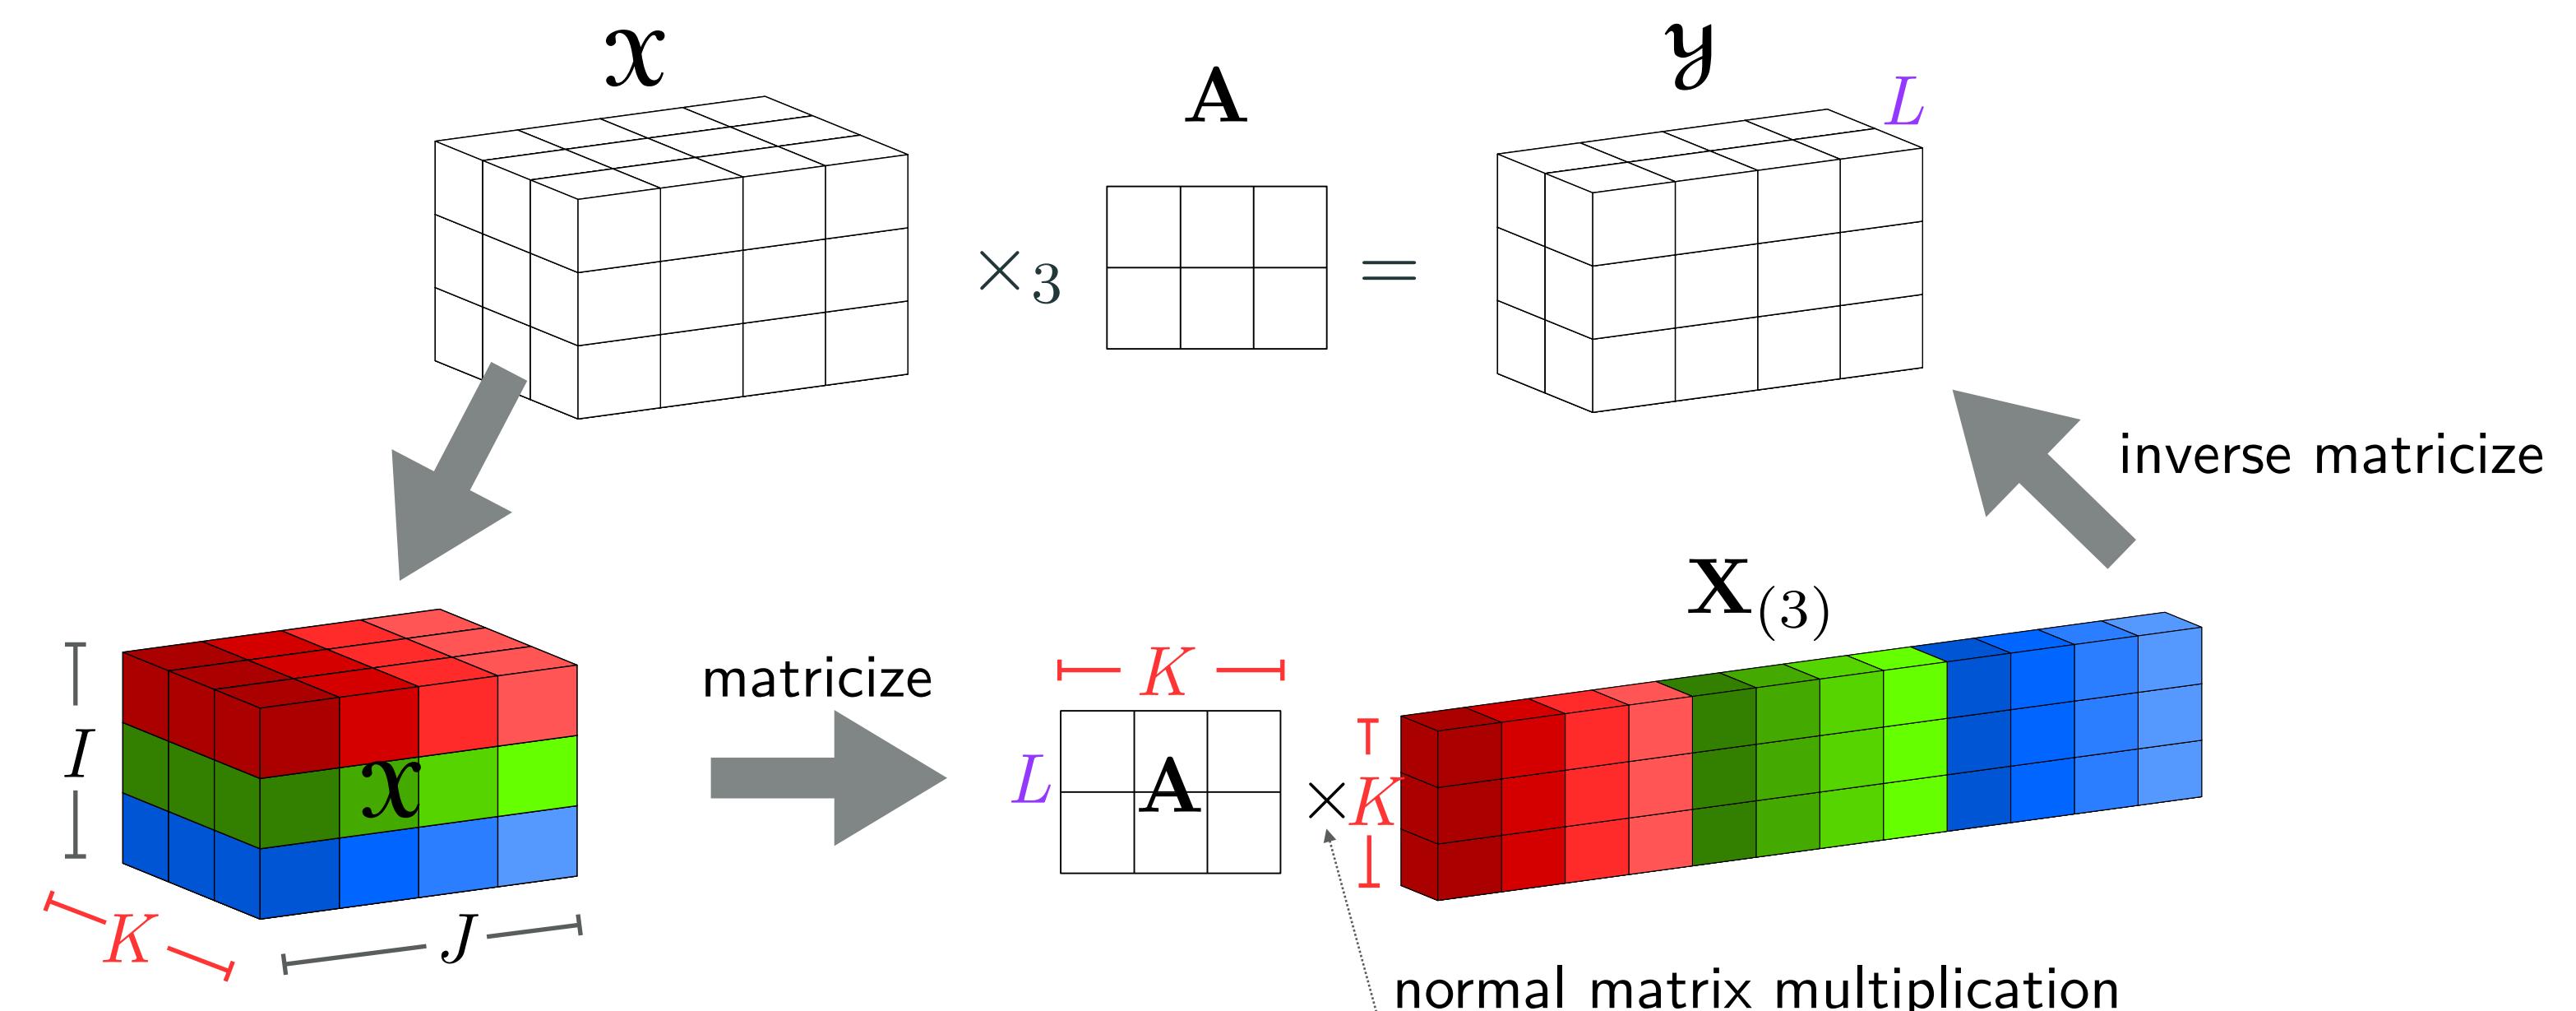
Example:

- Let $\mathcal{X} \in \mathbb{R}^{I \times J \times K}$ and $\mathbf{A} \in \mathbb{R}^{L \times K}$.
- We define $\mathcal{Y} \stackrel{\text{def}}{=} \mathcal{X} \times_3 \mathbf{A} \in \mathbb{R}^{I \times J \times L}$ elementwise by

$$y_{ijl} = \sum_{k=1}^K x_{ijk} a_{kl}.$$

- In matrix terms, $\mathbf{Y}_{(3)} = \mathbf{AX}_{(3)}$.

Conveniently, order doesn't matter*

$$\mathcal{X} \times_1 \mathbf{A}_1 \times_2 \mathbf{A}_2 = \mathcal{X} \times_2 \mathbf{A}_2 \times_1 \mathbf{A}_1$$


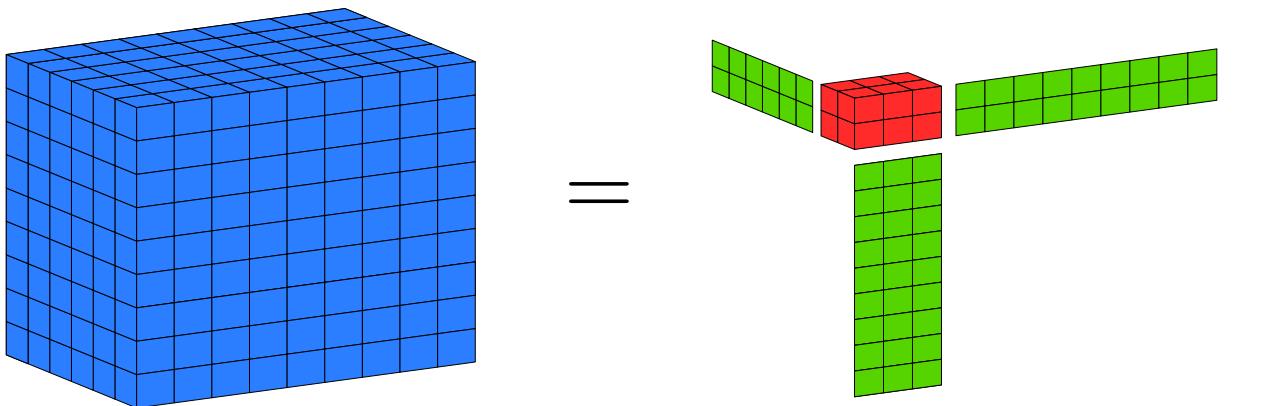
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

The Tucker decomposition

A Tucker decomposition of $\mathcal{X} \in \mathbb{R}^{I_1 \times I_2 \times \dots \times I_N}$ is of the form

$$\mathcal{X} = \mathcal{G} \times_1 \mathbf{A}^{(1)} \times_2 \mathbf{A}^{(2)} \dots \times_N \mathbf{A}^{(N)} =: [\![\mathcal{G}; \mathbf{A}^{(1)}, \mathbf{A}^{(2)}, \dots, \mathbf{A}^{(N)}]\!] \text{ (shorthand notation)}$$

where $\mathcal{G} \in \mathbb{R}^{R_1 \times R_2 \times \dots \times R_N}$ and each $\mathbf{A}^{(n)} \in \mathbb{R}^{I_n \times R_n}$.



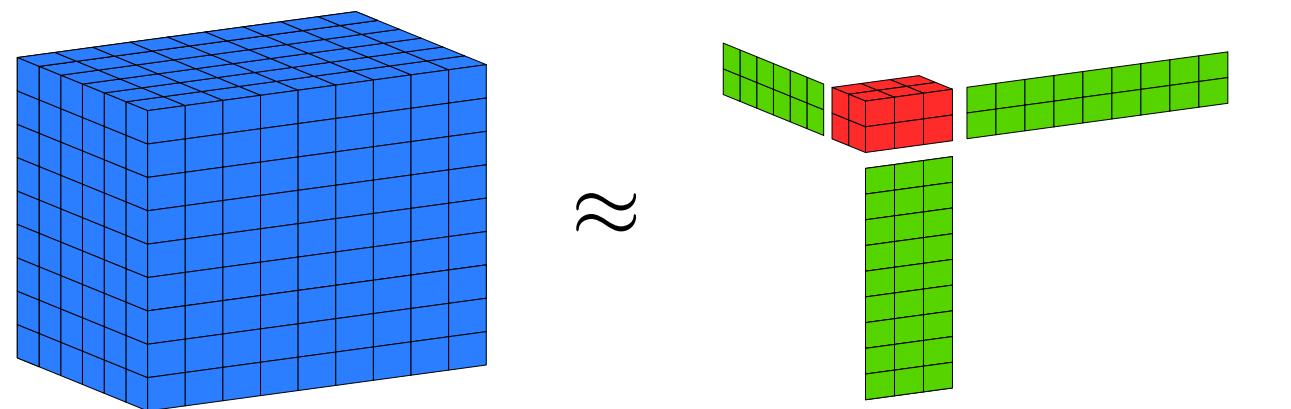
\mathcal{G} is the *core tensor* and $\mathbf{A}^{(n)}$, $n = 1, 2, \dots, N$, are *factor matrices*.

We say that \mathcal{X} is a rank- (R_1, R_2, \dots, R_N) tensor.

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Applications of the Tucker decomposition

- **Data compression.** Storing the **red** and **green** objects require less storage than the full **blue** tensor.
- **Data analysis.** Use rows of factor matrices as feature vectors—*unsupervised learning*.



Computing the Tucker decomposition

Given a data tensor $\mathcal{X} \in \mathbb{R}^{I_1 \times I_2 \times \dots \times I_N}$, we can find a decomposition $\mathcal{X} \approx [\mathcal{G}; \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)}]$ by solving the optimization problem

$$\min_{\mathcal{G}, \mathbf{A}^{(1)}, \dots, \mathbf{A}^{(N)}} \|\mathcal{X} - [\mathcal{G}; \mathbf{A}^{(1)}, \mathbf{A}^{(2)}, \dots, \mathbf{A}^{(N)}]\|^2, \quad (4)$$

where $\mathcal{G} \in \mathbb{R}^{R_1 \times R_2 \times \dots \times R_N}$ and $\mathbf{A}^{(n)} \in \mathbb{R}^{I_n \times R_n}$ for $n = 1, 2, \dots, N$.

- Nonlinear!
- A common approach to this issue is to use *alternating least squares (ALS)*: Minimize with respect to each $\mathbf{A}^{(n)}$ and \mathcal{G} one at a time. Repeat! i.e., Gauss-Siedel style
- With ALS, each subproblem is linear and easily solved.

ALS same as *higher-order orthogonal iteration* (HOOI)
and if you do just one iteration, it's like *higher-order SVD* (HOSVD)

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Computing the Tucker decomposition

The algorithm looks like this. Repeat the following until convergence:

Repeat until convergence

1. For $n = 1, \dots, N$, update (10)

$$\mathbf{A}^{(n)} = \arg \min_{\mathbf{A} \in \mathbb{R}^{I_n \times R_n}} \left\| \left(\bigotimes_{\substack{i=1 \\ i \neq n}}^N \mathbf{A}^{(i)} \right) \mathbf{G}_{(n)}^\top \mathbf{A}^\top - \mathbf{X}_{(n)}^\top \right\|_F^2. \quad (11)$$

2. Update $\mathcal{G} = \mathcal{X} \times_1 \mathbf{A}^{(1)\top} \times_2 \mathbf{A}^{(2)\top} \cdots \times_N \mathbf{A}^{(N)\top}$. (12)

$$= \arg \min_{\mathbf{Z} \in \mathbb{R}^{R_1 \times \cdots \times R_N}} \left\| \left(\bigotimes_{i=1}^N \mathbf{A}^{(i)} \right) \mathbf{Z}(:, :) - \mathbf{X}(:, :) \right\|_2^2.$$

(Line 12 follows if the factor matrices are orthogonalized)

Make extensive use of this identity:

$$\mathcal{Y} = \mathcal{X} \times_1 \mathbf{A}^{(1)} \times_2 \mathbf{A}^{(2)} \cdots \times_N \mathbf{A}^{(N)} \iff$$

$$\mathbf{Y}_{(n)} = \mathbf{A}^{(n)} \mathbf{X}_{(n)} \left(\mathbf{A}^{(N)} \otimes \cdots \otimes \mathbf{A}^{(n+1)} \otimes \mathbf{A}^{(n-1)} \otimes \cdots \otimes \mathbf{A}^{(1)} \right)^\top$$

Computing the Tucker decomposition

The algorithm looks like this. Repeat the following until convergence:

1. For $n = 1, \dots, N$, update (10)

$$\mathbf{A}^{(n)} = \arg \min_{\mathbf{A} \in \mathbb{R}^{I_n \times R_n}} \left\| \left(\bigotimes_{\substack{i=1 \\ i \neq n}}^N \mathbf{A}^{(i)} \right) \mathbf{G}_{(n)}^\top \mathbf{A}^\top - \mathbf{X}_{(n)}^\top \right\|_F^2. \quad (11)$$

2. Update $\mathcal{G} = \mathcal{X} \times_1 \mathbf{A}^{(1)\top} \times_2 \mathbf{A}^{(2)\top} \cdots \times_N \mathbf{A}^{(N)\top}$. (12)

Issues when dealing with large, sparse tensors:

- The Kronecker product matrix is huge.
- The series of tensor-times-matrix products is very costly, and can require lots of additional memory when \mathcal{X} is sparse.
- We address this problem by using *TensorSketch*.

Aside: we initially thought that **solving** the least-squares problem would be an issue, and something to fix with sketching, but this is not exactly the case

(partly because of many clever tricks due to orthogonality of \mathbf{A})

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Our modified algorithm: Tucker-TS



We simply sketch each of these least-squares problems with TensorSketch operators $\mathbf{S}^{(n)}$, $n = 1, \dots, N + 1$, which are defined at the start of the algorithm:

Repeat until convergence

1. For $n = 1, \dots, N$, update

$$\mathbf{A}^{(n)} = \arg \min_{\mathbf{A} \in \mathbb{R}^{I_n \times R_n}} \left\| \left(\mathbf{S}^{(n)} \bigotimes_{\substack{i=1 \\ i \neq n}}^N \mathbf{A}^{(i)} \right) \mathbf{G}_{(n)}^\top \mathbf{A}^\top - \mathbf{S}^{(n)} \mathbf{x}_{(n)}^\top \right\|_F^2. \quad (18)$$

(and orthogonalize)

2. Update

$$\mathbf{G} = \arg \min_{\mathbf{z} \in \mathbb{R}^{R_1 \times \dots \times R_N}} \left\| \left(\mathbf{S}^{(N+1)} \bigotimes_{i=1}^N \mathbf{A}^{(i)} \right) \mathbf{z}_{(:)} - \mathbf{S}^{(N+1)} \mathbf{x}_{(:)} \right\|_2^2. \quad (19)$$

Tricks:

- don't draw new sketches each iteration
- the TensorSketches have [CountSketches/hashes in common](#) $\mathbf{S}^{(n)}, \mathbf{S}^{(n')}$ (each is leave-one-out)
- G update solved via conjugate gradient (CG) algo, and can prove it is well-conditioned
- advanced tricks and variants in the paper

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

Numerical experiment: sparse 3D tensor

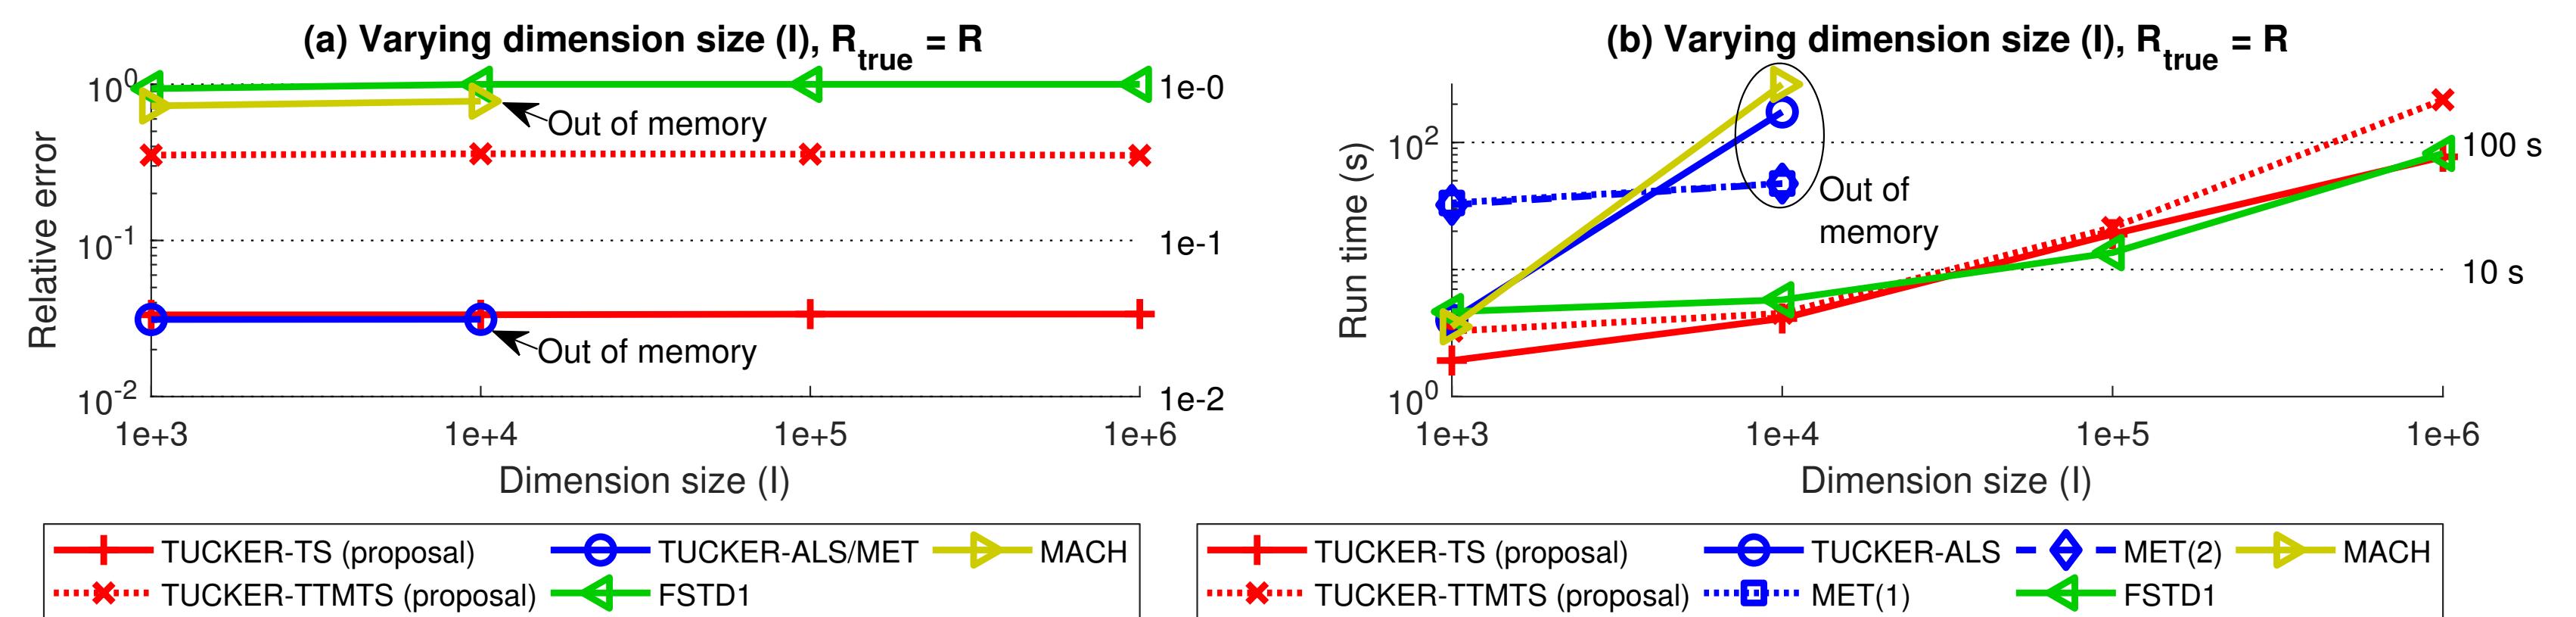


Figure 3: Relative error and run time for random sparse 3-way tensors with varying dimension size I and $\text{nnz}(\mathcal{Y}) \approx 1e+6$. Both the true and target ranks are $(10, 10, 10)$.

Tucker-TS approximates least-squares like this:

$$\begin{aligned} \operatorname{argmin}_x \frac{1}{2} \|Ax - b\|_2^2 &\approx \operatorname{argmin}_x \frac{1}{2} \|\Phi(Ax - b)\|_2^2 \\ &= (A^\top \Phi^\top \Phi A)^{-1} A^\top \Phi^\top \Phi b \end{aligned}$$

Tucker-TTMTS approximates least-squares like this:

$$\begin{aligned} \operatorname{argmin}_x \frac{1}{2} \|Ax - b\|_2^2 &= (A^\top A)^{-1} A^\top b \\ &\approx (A^\top A)^{-1} A^\top \Phi^\top \Phi b \\ &\quad (\text{and exploit orthogonality here}) \end{aligned}$$

Video

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. **Tensor factorizations**
 - d. Gradient-free optimization

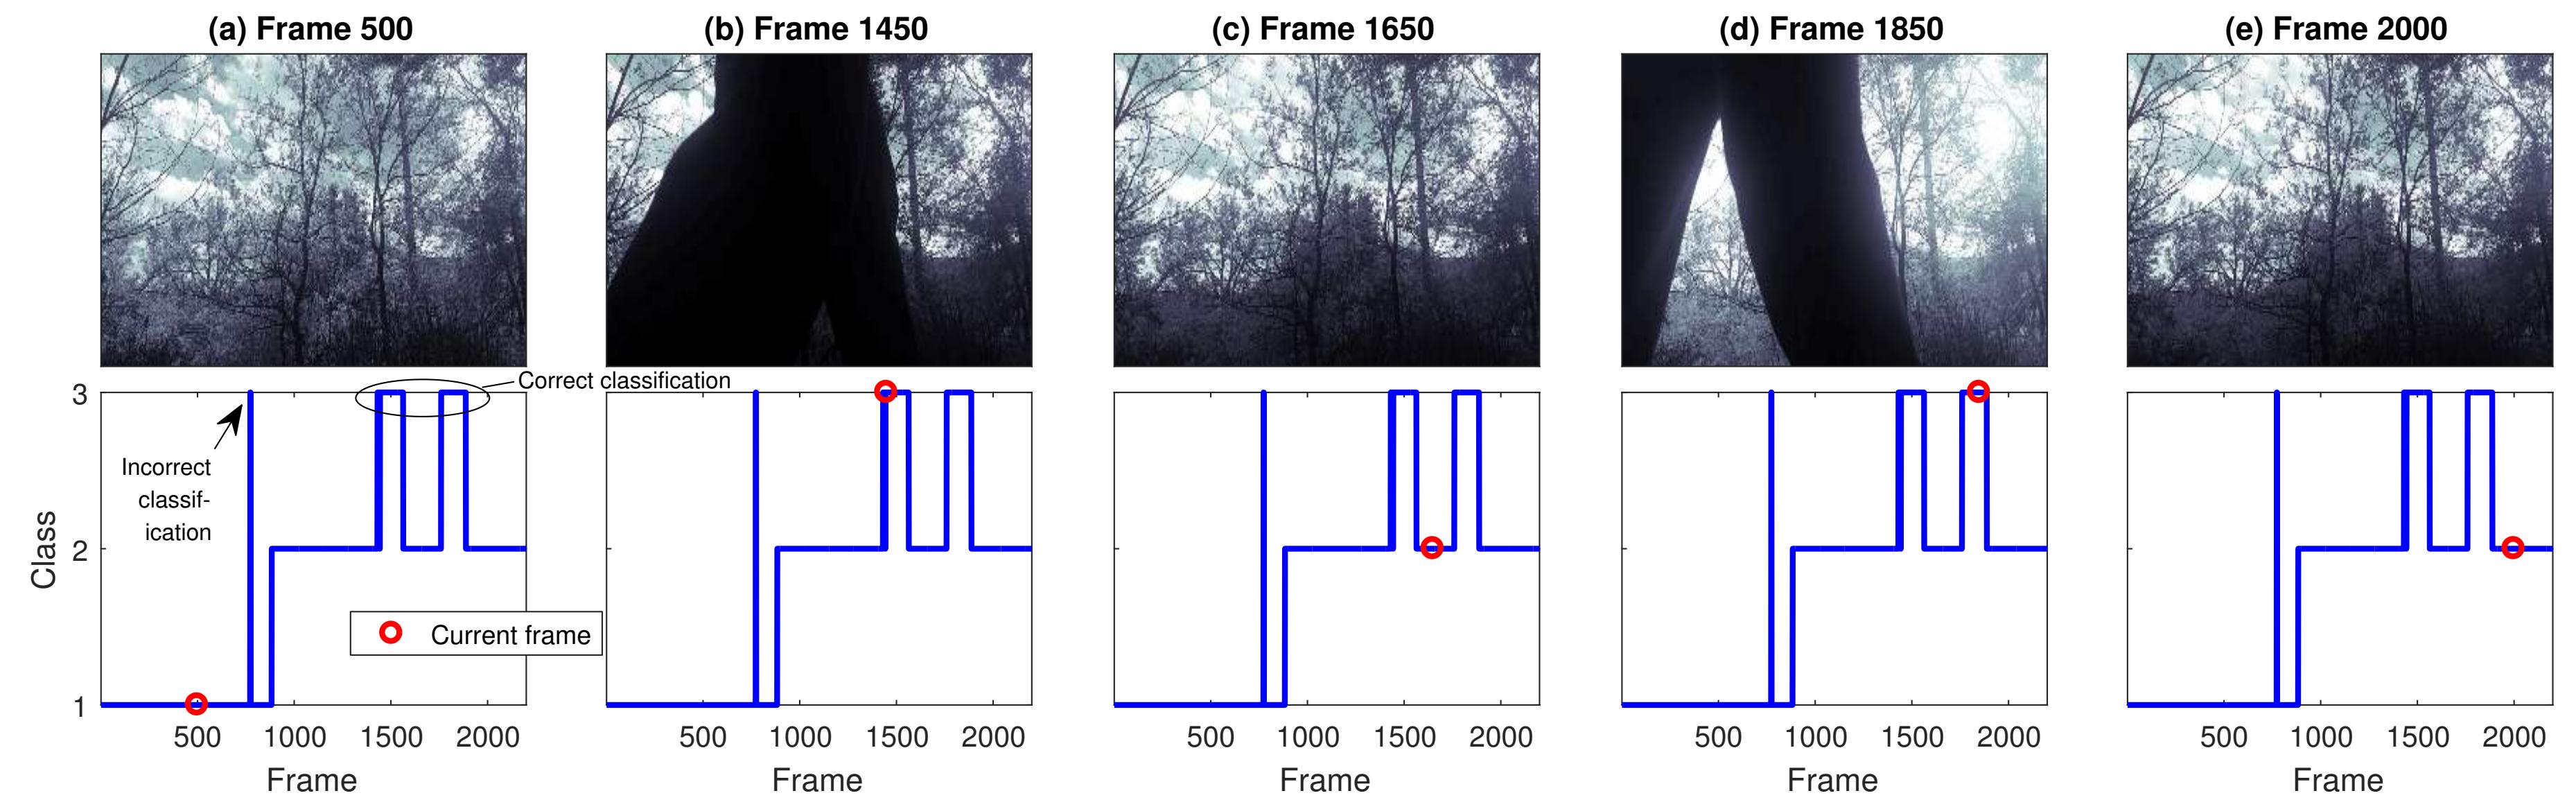


Figure 6: Five sample frames with their assigned classes. The frames (b) and (d) contain a disturbance.

38 GB video (can't all fit into RAM): 2,200 frames, each of size 1,080 by 1,980 pixels

Compute rank (10,10,10) factorization, 30 iterations max

Find 3 factor matrices — take the one corresponding to time and run k-means clustering

We load chunk-by-chunk and do the sketches (in one pass)

(used Tucker-TTMTS variant)

OUTLINE

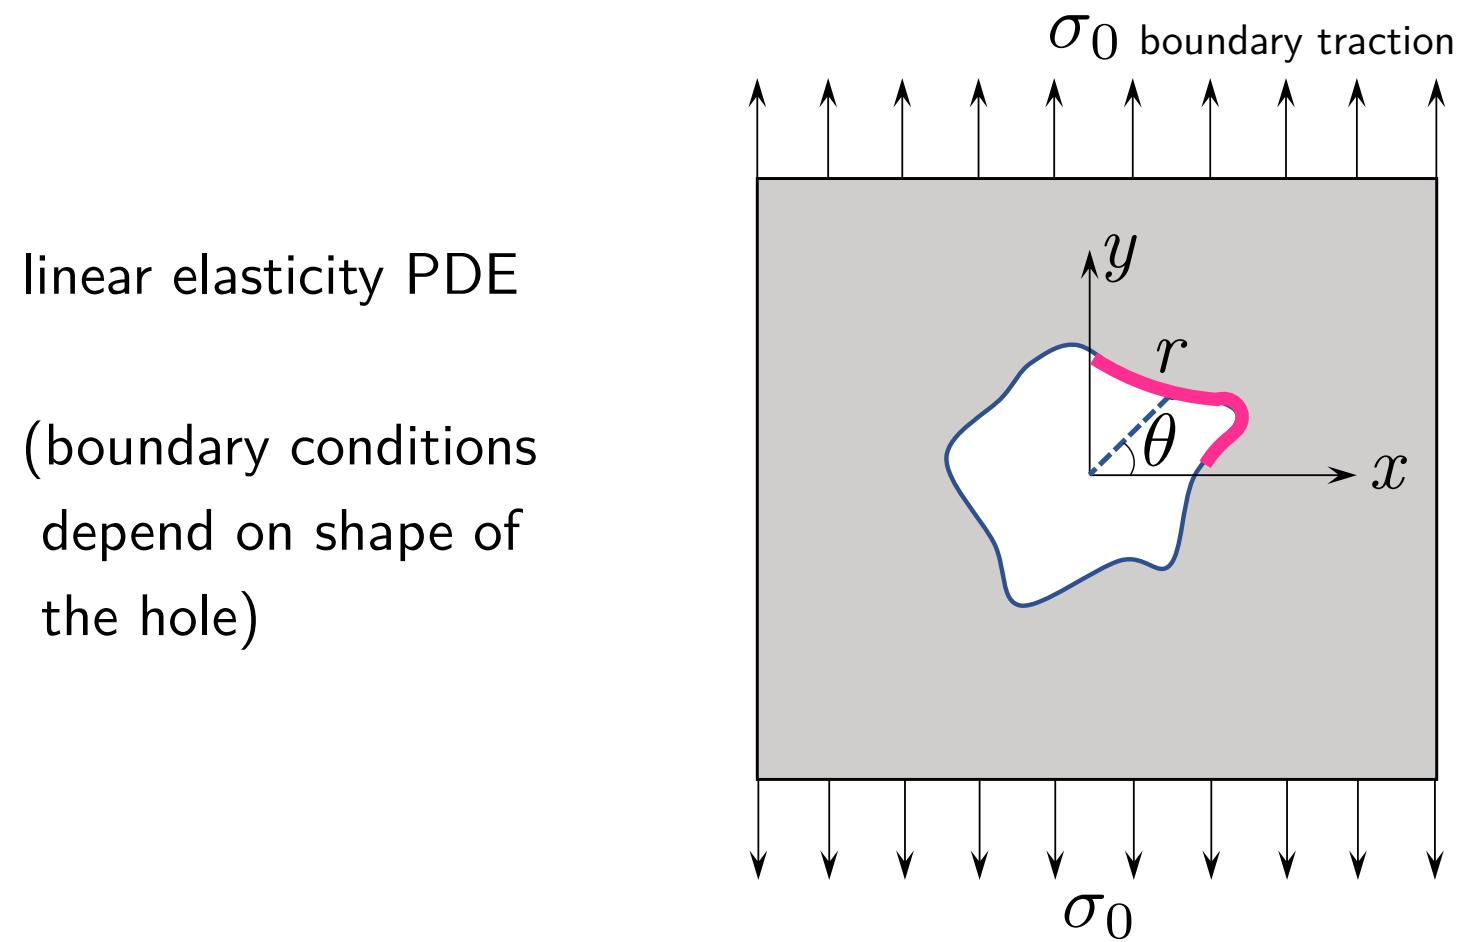
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization



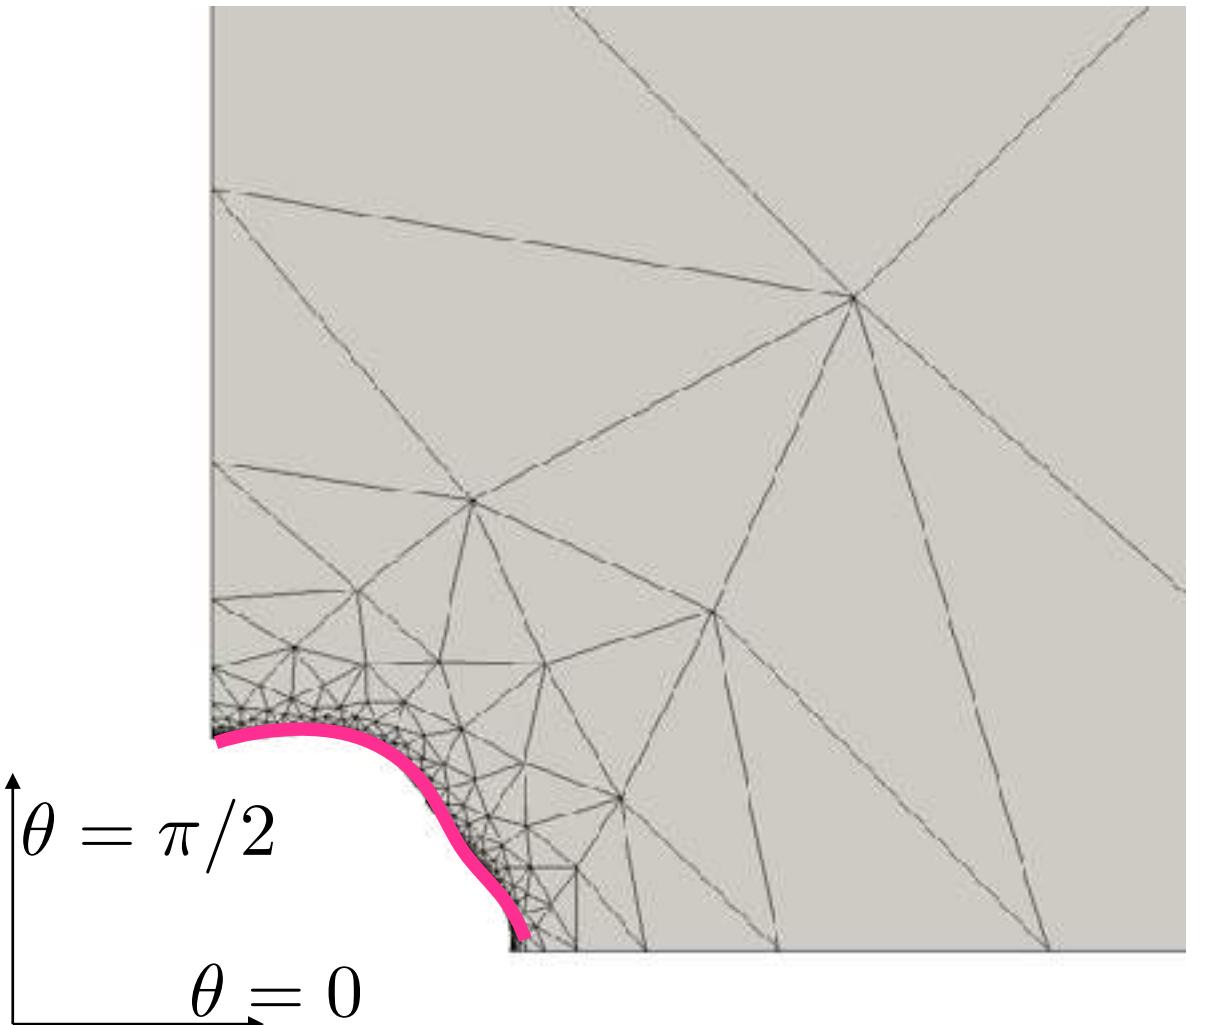
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Motivating Application: shape optimization

Forward problem: find vertical stress σ_y



conforming finite element mesh, used in FEniCS

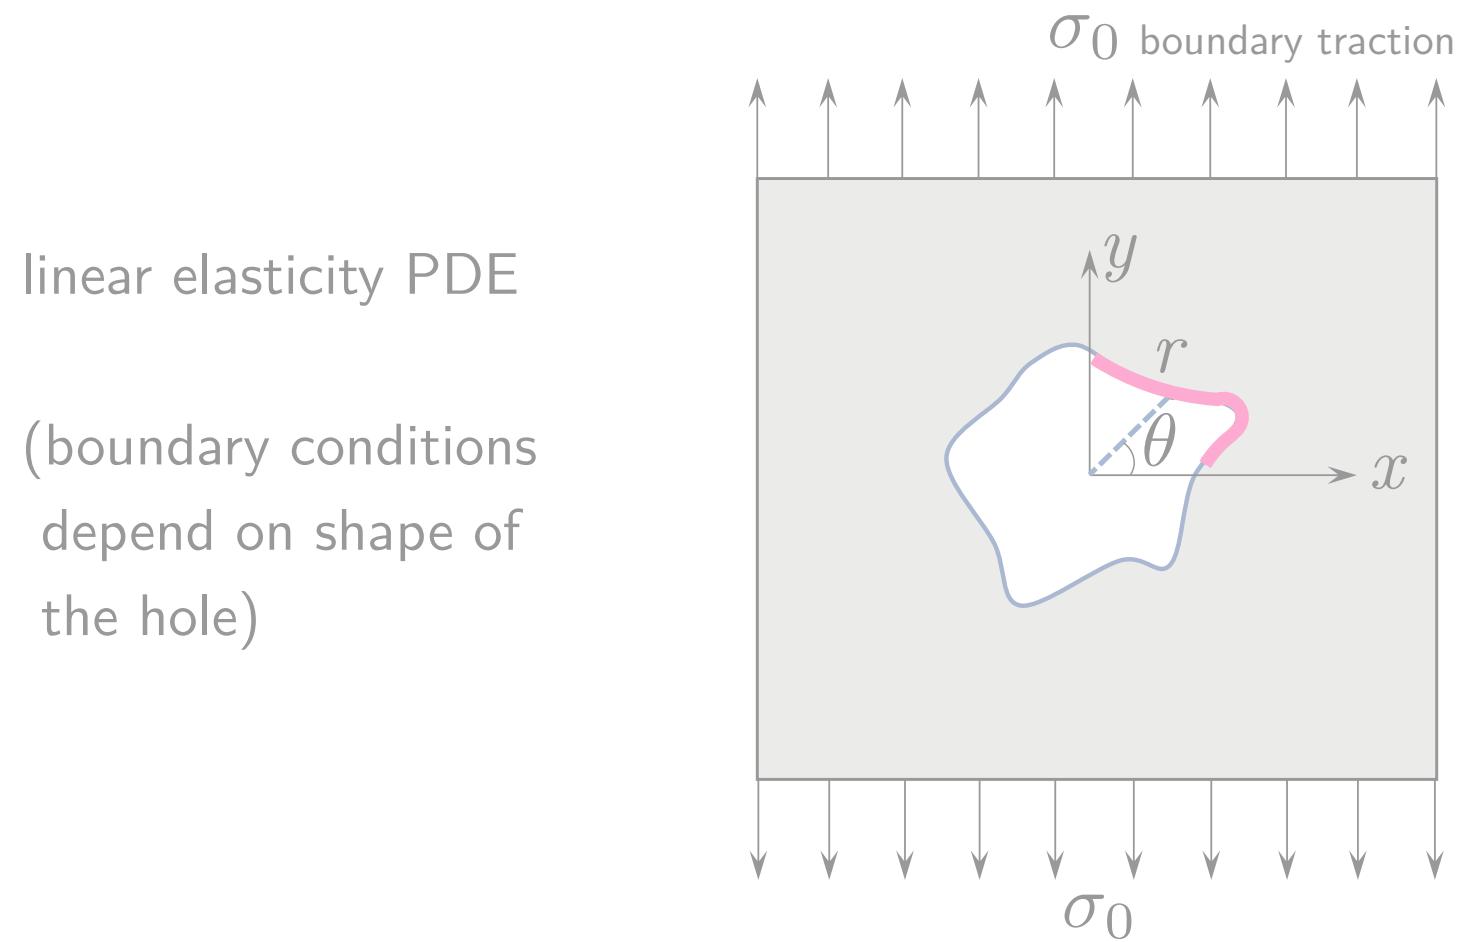


Change in notation temporarily

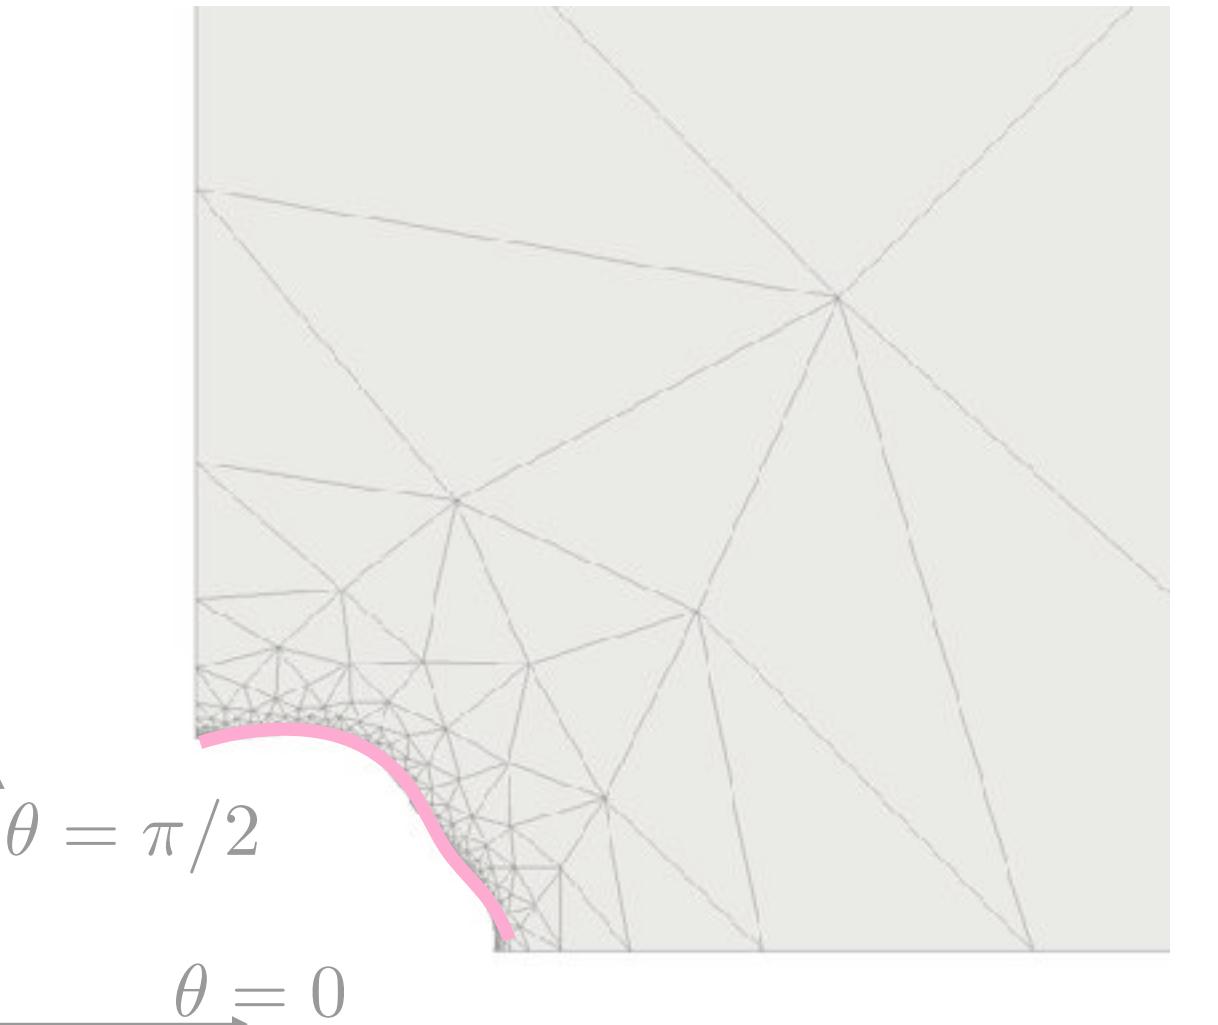
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. **Gradient-free optimization**

Motivating Application: shape optimization

Forward problem: find vertical stress σ_y



conforming finite element mesh, used in FEniCS



Inverse problem: what shape minimizes the vertical stress?

parameterize the shape of the hole as follows, which automatically enforces a constant area constraint

$$r(\theta) = \frac{1}{2\pi} + \delta \sum_{k=0}^{d/2-1} \frac{1}{\sqrt{2k+1}} (\textcolor{red}{x}_{2k+1} \sin((2k+1) \cdot \theta) + \textcolor{red}{x}_{2k+2} \cos((2k+1) \cdot \theta))$$

optimization variable:

$$\textcolor{red}{x} \in \mathbb{R}^d$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. **Gradient-free optimization**

Generic PDE-constrained optimization

implicitly saying that u solves the PDE

$$\boxed{\min_{u,x} \mathcal{L}(u) \quad \text{subject to} \quad \phi(u, x) = 0}$$

Examples:

$$\dot{u} = \Delta u, \quad u(0) = \textcolor{red}{h} \quad "x" \text{ is the initial condition}$$

$$\ddot{u} = \textcolor{red}{c}^2 \Delta u, \quad u(0) = h \quad "x" \text{ is a parameter}$$

$$\Delta u = 0, \quad u(\Gamma) = \textcolor{red}{h} \quad "x" \text{ is the boundary condition}$$

$\mathcal{L}(u)$ is the loss which penalizes something like:

- deviation from observations
- drag
- mass
- cost of materials
- compliance
- etc.

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. **Gradient-free optimization**

Generic PDE-constrained optimization

implicitly saying that u solves the PDE

$$\boxed{\min_{u,x} \mathcal{L}(u) \quad \text{subject to} \quad \phi(u, x) = 0}$$

$$\phi(u, x) = 0 \implies u = u(x)$$

Rewrite:

$$\boxed{\min_x f(x) \stackrel{\text{def}}{=} \mathcal{L}(u(x))}$$

... but finding the gradient is tricky:

$$\nabla f(x) = \frac{\partial \mathcal{L}}{\partial u} \cdot \frac{\partial u}{\partial x}$$

Why not just find gradients automatically?

✓ The **adjoint state method** and **reverse-mode automatic differentiation** can automatically calculate gradients in about the same time ($\sim 4x$) as a function evaluation

- ... so if we can evaluate $f(x)$ numerically, we can find the gradient
- (this applies if $f : \mathbb{R}^d \rightarrow \mathbb{R}$; $f : \mathbb{R}^d \rightarrow \mathbb{R}^q$ with $q \gg 1$ is another story)

Note: we're assuming derivative exists, just hard to actually calculate

This is *not* non-smooth optimization

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. **Gradient-free optimization**

Why not just find gradients automatically?



The **adjoint state method** and **reverse-mode automatic differentiation** can automatically calculate gradients in about the same time ($\sim 4x$) as a function evaluation



Requires specialized/restricted libraries/code (`dolfin-adjoint/FEniCS`, `autograd`)

`Jax`, `PyTorch`, `Tensorflow`, ...

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Why not just find gradients automatically?



The adjoint state method and reverse-mode automatic differentiation can automatically calculate gradients in about the same time ($\sim 4x$) as a function evaluation



Requires specialized/restricted libraries/code (`dolfin-adjoint/FEniCS, autograd`)

X Adjoint state method requires a method to solve adjoint PDE

(and have to parallelize for HPC)

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Why not just find gradients automatically?



The adjoint state method and reverse-mode automatic differentiation can automatically calculate gradients in about the same time ($\sim 4x$) as a function evaluation



Requires specialized/restricted libraries/code (`dolfin-adjoint/FEniCS, autograd`)

Adjoint state method requires a method to solve adjoint PDE

- difficult to maintain in large code bases, e.g., 4D-var for weather codes



Slow if used for intermediate calculations involving some

$$f : \mathbb{R}^d \rightarrow \mathbb{R}^q$$

- e.g., seismic inversion with many observations

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Why not just find gradients automatically?

✓ The adjoint state method and reverse-mode automatic differentiation can automatically calculate gradients in about the same time ($\sim 4x$) as a function evaluation

- ✗ Requires specialized/restricted libraries/code (`dolfin-adjoint/FEniCS, autograd`)
- ✗ Adjoint state method requires a method to solve adjoint PDE
 - difficult to maintain in large code bases, e.g., 4D-var for weather codes
- ✗ Slow if used for intermediate calculations involving some $f : \mathbb{R}^d \rightarrow \mathbb{R}^q$
 - e.g., seismic inversion with many observations
- ✗ Possible memory explosion
 - e.g., time-dependent problems. Check-pointing schemes somewhat helpful

Example: hyper-parameter optimization in deep learning

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Why not just find gradients automatically?

✓ The adjoint state method and reverse-mode automatic differentiation can automatically calculate gradients in about the same time ($\sim 4x$) as a function evaluation

- ✗ Requires specialized/restricted libraries/code (`dolfin-adjoint/FEniCS, autograd`)
- ✗ Adjoint state method requires a method to solve adjoint PDE
 - difficult to maintain in large code bases, e.g., 4D-var for weather codes
- ✗ Slow if used for intermediate calculations involving some $f : \mathbb{R}^d \rightarrow \mathbb{R}^q$
 - e.g., seismic inversion with many observations
- ✗ Possible memory explosion
 - e.g., time-dependent problems. Check-pointing schemes somewhat helpful
- ✗ Requires access to original source code

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Why not just find gradients automatically?

✓ The adjoint state method and reverse-mode automatic differentiation can automatically calculate gradients in about the same time ($\sim 4x$) as a function evaluation

- ✗ Requires specialized/restricted libraries/code (`dolfin-adjoint/FEniCS, autograd`)
- ✗ Adjoint state method requires a method to solve adjoint PDE
 - difficult to maintain in large code bases, e.g., 4D-var for weather codes
- ✗ Slow if used for intermediate calculations involving some $f : \mathbb{R}^d \rightarrow \mathbb{R}^q$
 - e.g., seismic inversion with many observations
- ✗ Possible memory explosion
 - e.g., time-dependent problems. Check-pointing schemes somewhat helpful
- ✗ Requires access to original source code
- ✗ Assumes a computational structure
 - inapplicable for physical observations (wind farms; *rollout* in AI)

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Baseline Algorithms (for comparison)

$$f : \mathbb{R}^d \rightarrow \mathbb{R}$$

These days, often called “zeroth order” optimization

Algorithm Gradient Descent via Finite Differences

```

1: for  $k = 1, 2, \dots$  do
2:   Estimate  $g_k \approx \nabla f(x_k)$             $\triangleright$  Use finite differences
3:    $x_{k+1} \leftarrow x_k - \eta_k g_k$         $\triangleright$  For appropriate step-size  $\eta_k$ 

```

ignoring finite-difference error, enjoys well-understood convergence

requires $d+1$ function evaluations per iter.

Why not use traditional Derivative Free Optimization (DFO) methods?

Answer: most classical DFO methods don't scale well with dimension

Note: there are many other DFO methods... but we won't discuss in this talk

From heuristics to theoretically based, from local to global

- Heuristics: Nelder-Mead, genetic algorithms, particle swarm optimization
 - e.g., CMA-ES, Covariance matrix adaptation evolution strategy
- Simulated Annealing (typical as heuristic)
- Bayesian Optimization
 - Model is Gaussian Process, with acquisition function (exploration/exploitation tradeoff)
- DFO-TR (Trust-region)
 - Model is polynomial

Baseline Algorithms (for comparison)

$$f : \mathbb{R}^d \rightarrow \mathbb{R}$$

Algorithm Gradient Descent via Finite Differences

```
1: for  $k = 1, 2, \dots$  do
2:   Estimate  $g_k \approx \nabla f(x_k)$        $\triangleright$  Use finite differences
3:    $x_{k+1} \leftarrow x_k - \eta_k g_k$      $\triangleright$  For appropriate step-size  $\eta_k$ 
```

✓ ignoring finite-difference error, enjoys well-understood convergence

✗ requires $d+1$ function evaluations per iter.

Algorithm Randomized Coordinate Descent (CD)

```
1: for  $k = 1, 2, \dots$  do
2:   Choose  $j \in \{1, 2, \dots, d\}$  at random
3:    $g_k = e_j e_j^T \nabla f(x_k)$ 
4:    $x_{k+1} \leftarrow x_k - \eta_k g_k$      $\triangleright$  For appropriate step-size  $\eta_k$  (or exact minimization... depends on structure)
```

✓ just 1 function evaluation per iteration

✗ poor convergence properties, slow rates

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. **Gradient-free optimization**

Stochastic Subspace Descent

directional derivative $qq^T \nabla f(x_k) = \left(\lim_{h \rightarrow 0} \frac{f(x_k + h \cdot q) - f(x_k)}{h} \right) q$

Assume we can compute this!

e.g.,

- 1) forward finite diff
- 2) forward-mode AD

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. **Gradient-free optimization**

Stochastic Subspace Descent

directional derivative $qq^T \nabla f(x_k) = \left(\lim_{h \rightarrow 0} \frac{f(x_k + h \cdot q) - f(x_k)}{h} \right) q$

$$Q = [q_1, q_2, \dots, q_\ell] \sim \text{Haar}(d \times \ell) \quad Q^T Q = I_{\ell \times \ell}, \quad \mathbb{E} \left(\frac{d}{\ell} Q Q^T \right) = I_{d \times d}$$

One benefit: in the limit $\ell = d$, $Q Q^T = I_{d \times d}$, and so we'll recover the full gradient
(for Gaussians, this is only true in expectation)

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. **Gradient-free optimization**

Stochastic Subspace Descent

directional derivative $qq^T \nabla f(x_k) = \left(\lim_{h \rightarrow 0} \frac{f(x_k + h \cdot q) - f(x_k)}{h} \right) q$

$$Q = [q_1, q_2, \dots, q_\ell] \sim \text{Haar}(d \times \ell) \quad Q^T Q = I_{\ell \times \ell}, \quad \mathbb{E} \left(\frac{d}{\ell} QQ^T \right) = I_{d \times d}$$

Algorithm “Stochastic Subspace Descent” (SSD)

```

1: for  $k = 1, 2, \dots$  do
2:   Draw  $Q \sim \text{Haar}(d \times \ell)$  or any generic SSD
3:    $x_{k+1} \leftarrow x_k - \eta_k \frac{d}{\ell} QQ^T \nabla f(x_k)$ 

```

Generic SSD

$$Q^T Q = I_{\ell \times \ell}, \quad \mathbb{E} \left(\frac{d}{\ell} QQ^T \right) = I_{d \times d}$$

Both **Haar** and **Coordinate Descent** methods are valid generic SSD

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. **Gradient-free optimization**

Stochastic Subspace Descent

directional derivative $qq^T \nabla f(x_k) = \left(\lim_{h \rightarrow 0} \frac{f(x_k + h \cdot q) - f(x_k)}{h} \right) q$

$$Q = [q_1, q_2, \dots, q_\ell] \sim \text{Haar}(d \times \ell) \quad Q^T Q = I_{\ell \times \ell}, \quad \mathbb{E} \left(\frac{d}{\ell} QQ^T \right) = I_{d \times d}$$

Algorithm “Stochastic Subspace Descent” (SSD)

```

1: for  $k = 1, 2, \dots$  do
2:   Draw  $Q \sim \text{Haar}(d \times \ell)$ 
3:    $x_{k+1} \leftarrow x_k - \eta_k \frac{d}{\ell} QQ^T \nabla f(x_k)$ 

```

We call Q a “Haar” distributed r.v. (i.e., the Haar measure over orthogonal matrices), but really care about QQ^T which is a projection matrix (onto $\text{col}(Q)$).

We get Q via Gram-Schmidt (or appropriately modified QR) on a Gaussian G , and note $\text{col}(Q) = \text{col}(G)$ w.p. 1, so our update is equivalent to

$$x_{k+1} \leftarrow x_k - \eta_k \frac{d}{\ell} \mathcal{P}_{\text{col}(G)}(\nabla f(x_k))$$

and hence the term “stochastic subspace”.

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

First theory results (for generic SSD)



Theorem (Kozak, Becker, Tenorio, Doostan '20)

Assume: minimizer attained, gradient Lipschitz, stepsize η_k chosen appropriately.

1. If f is **convex**,

$$\mathbb{E}f(x_k) - f^* \leq 2\frac{d}{\ell}\frac{L}{k}R^2 = \mathcal{O}(k^{-1})$$

$$\text{constant } L \quad \eta = \frac{\ell}{d} \frac{1}{L}$$

constant μ

2. If f is **not convex** but satisfies the **Polyak-Lojasiewicz** inequality,

$$\mathbb{E}f(x_k) - f^* \leq \rho^k(f(x_0) - f^*) = \mathcal{O}(\rho^k) \quad \text{and} \quad f(x_k) \xrightarrow{\text{a.s.}} f^*$$

3. If f is **strongly convex**, statements of 2 above hold, and also

$$x_k \xrightarrow{\text{a.s.}} \operatorname{argmin}_x f(x)$$

4. If f is **not convex** (nor PL),

$$\min_{k' \in \{0, \dots, k\}} \mathbb{E}\|\nabla f(x_{k'})\|^2 \leq \frac{d}{\ell} \frac{2L(f(x_0) - f^*)}{k + 1}$$

Generic SSD

$$Q^T Q = I_{\ell \times \ell}, \quad \mathbb{E} \left(\frac{d}{\ell} Q Q^T \right) = I_{d \times d}$$

$$f^* \stackrel{\text{def}}{=} \min_x f(x)$$

$$\rho = 1 - \frac{\mu}{L} \frac{\ell}{d}$$

d = ambient dimension

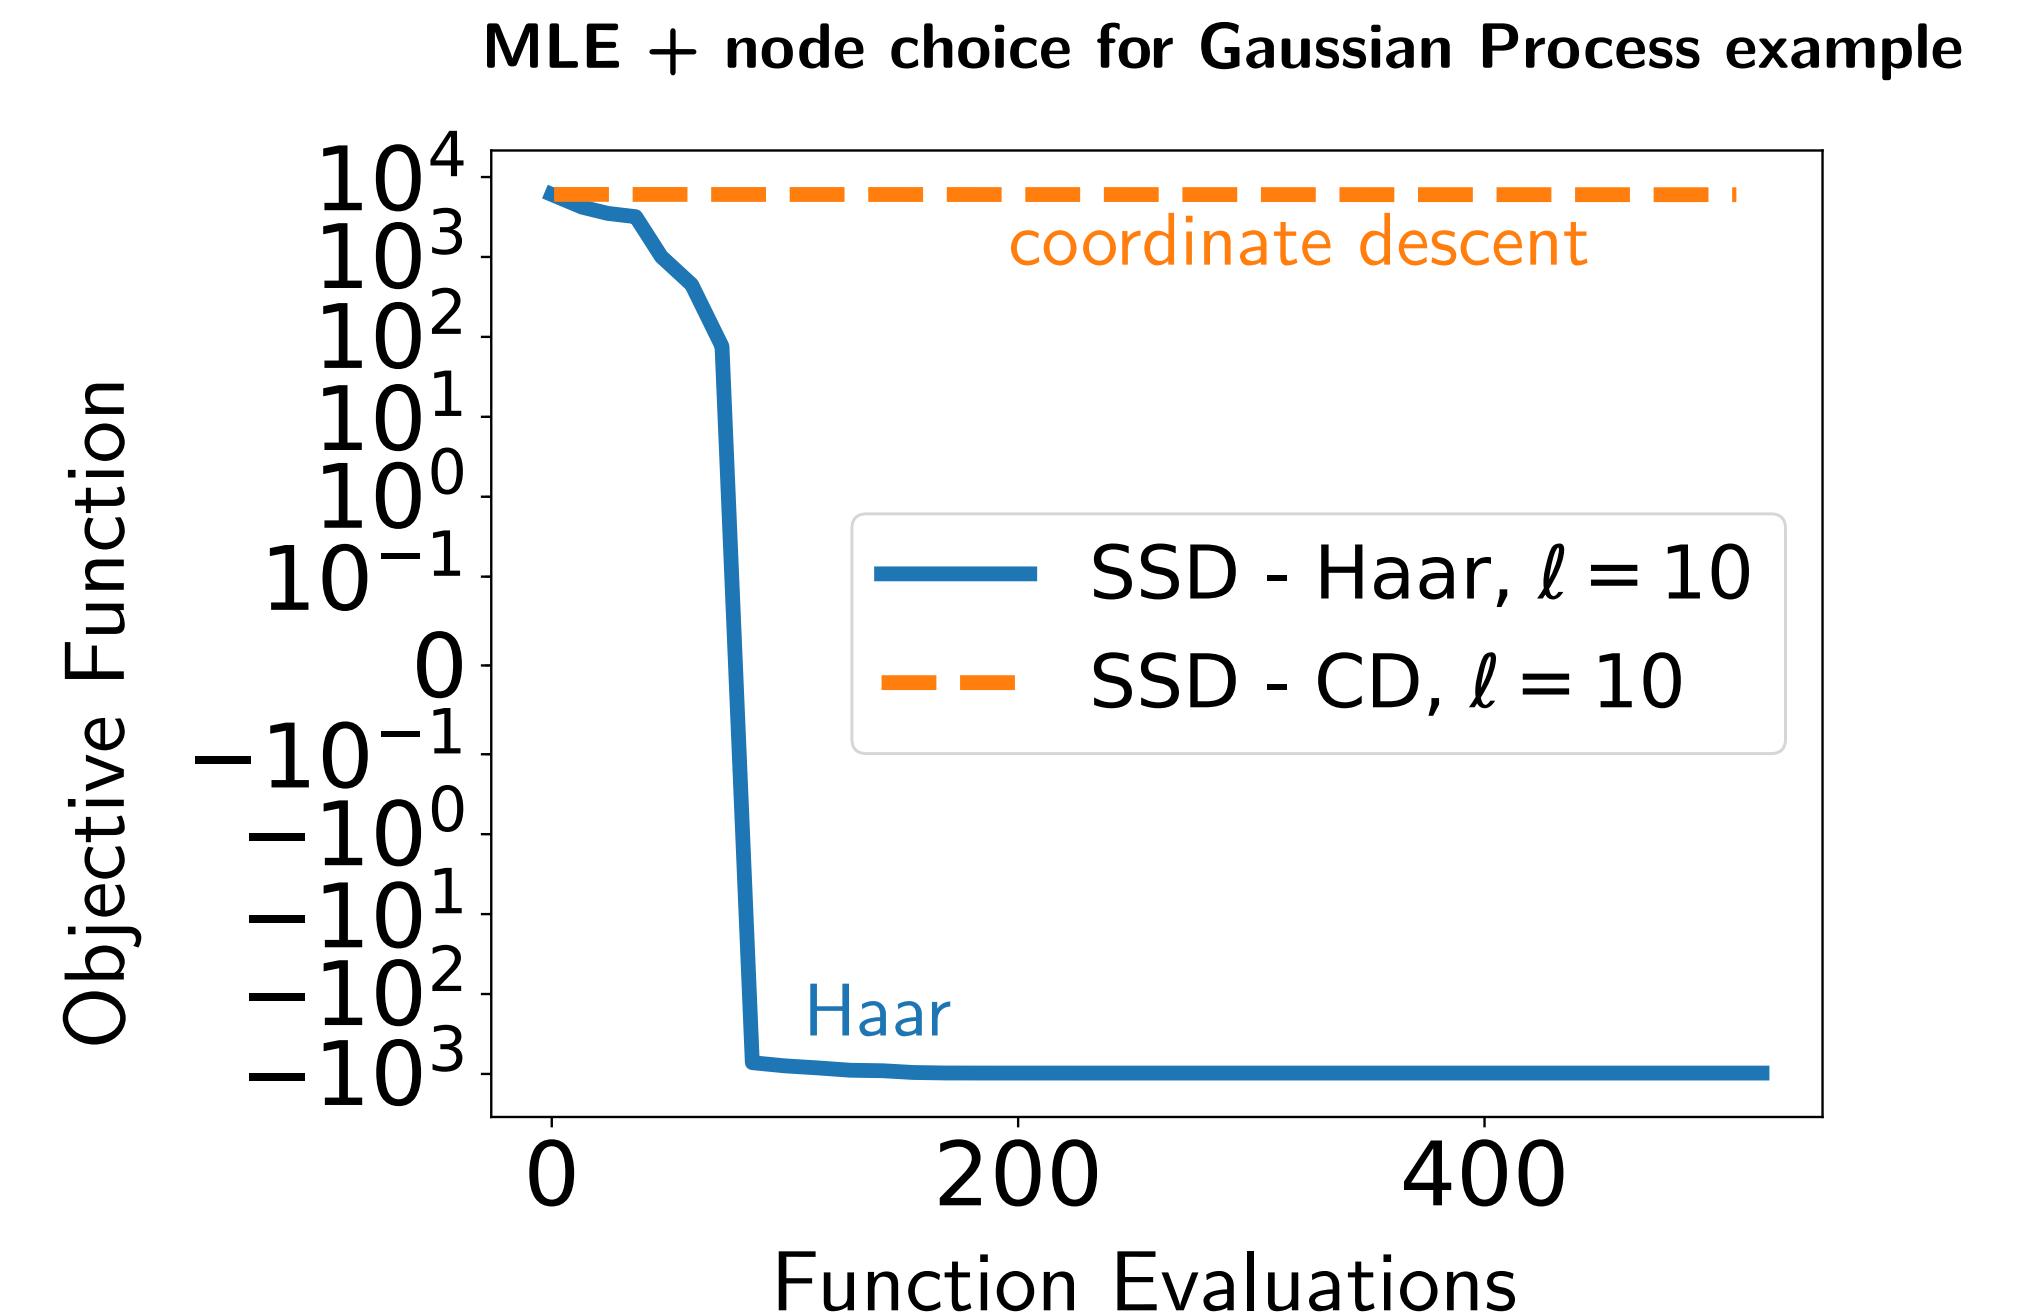
ℓ = # directional derivs

$\frac{d}{\ell} = 1$ is gradient descent

SSD is a **special type of SGD**,
but results are much better
than generic SGD analysis

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Numerical Results: better than expected



Observation: sometimes SSD (with Haar) drastically outperforms randomized coordinate descent (CD)

Generic SSD

$$Q^T Q = I_{\ell \times \ell}, \quad \mathbb{E} \left(\frac{d}{\ell} Q Q^T \right) = I_{d \times d}$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Numerical Results: better than expected

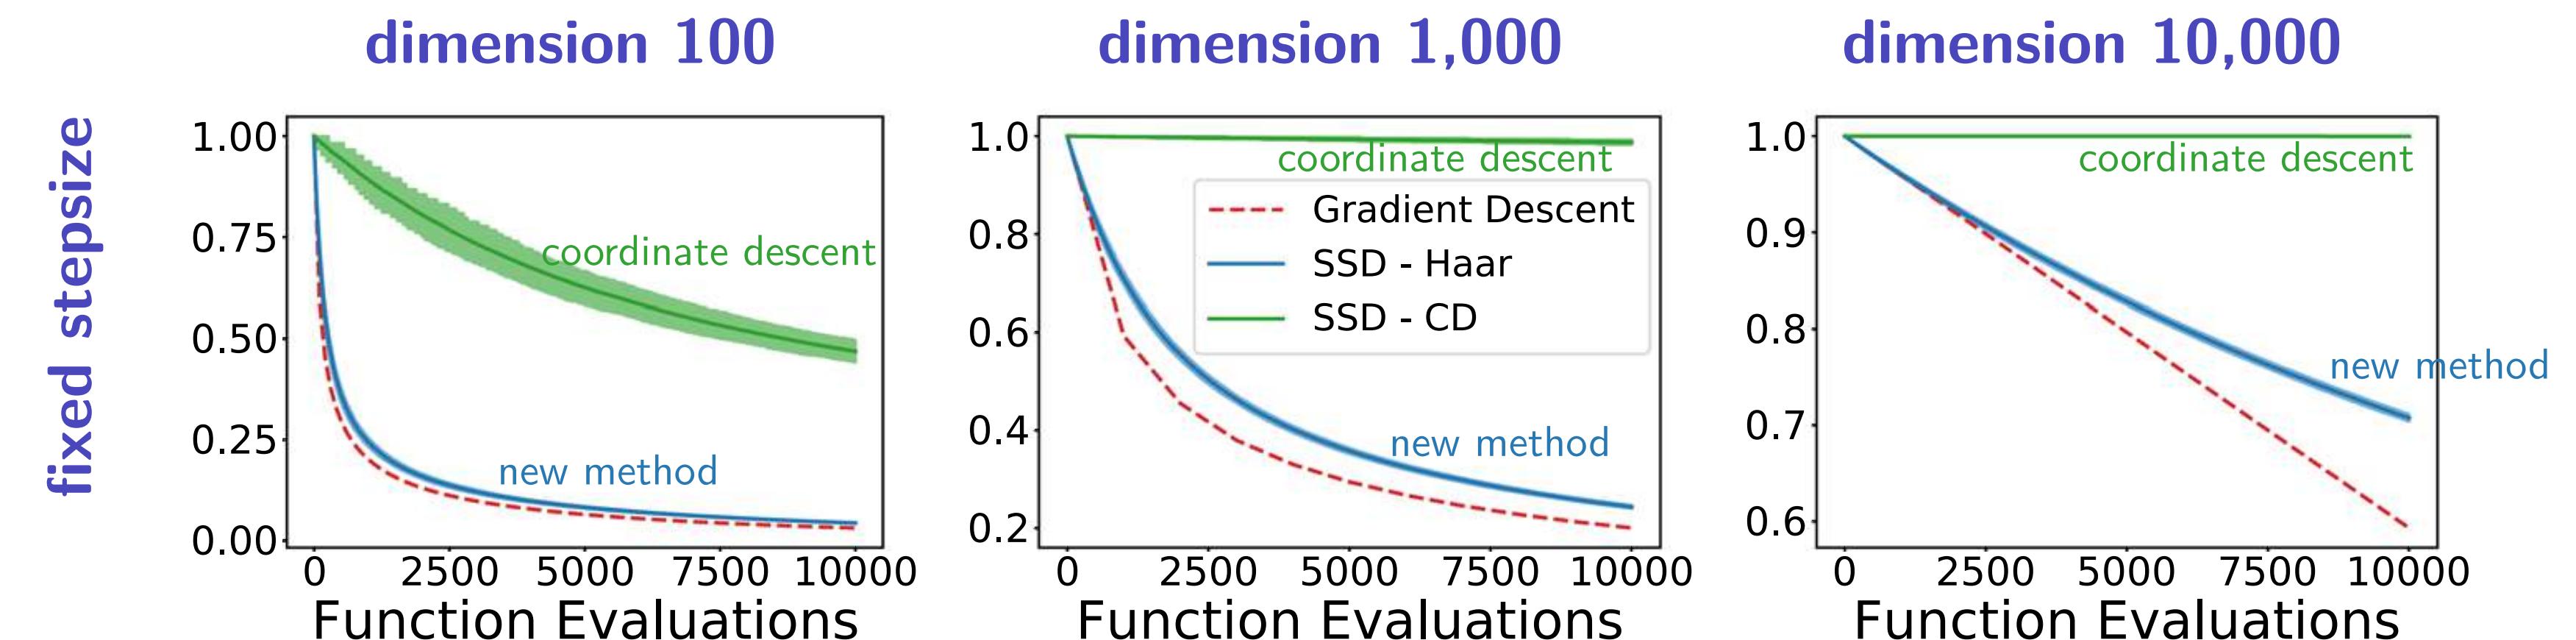
Haar SSD drastically outperforms randomized coordinate descent (CD)



We can force it to happen by making a problem with low “intrinsic” dimension, e.g., Nesterov’s “worst function in the world”

$$f_{\lambda,r}(\mathbf{x}) = \lambda((x_1^2 + \sum_{i=1}^{r-1} (x_i - x_{i+1})^2 + x_r^2)/2 - x_1)/4,$$

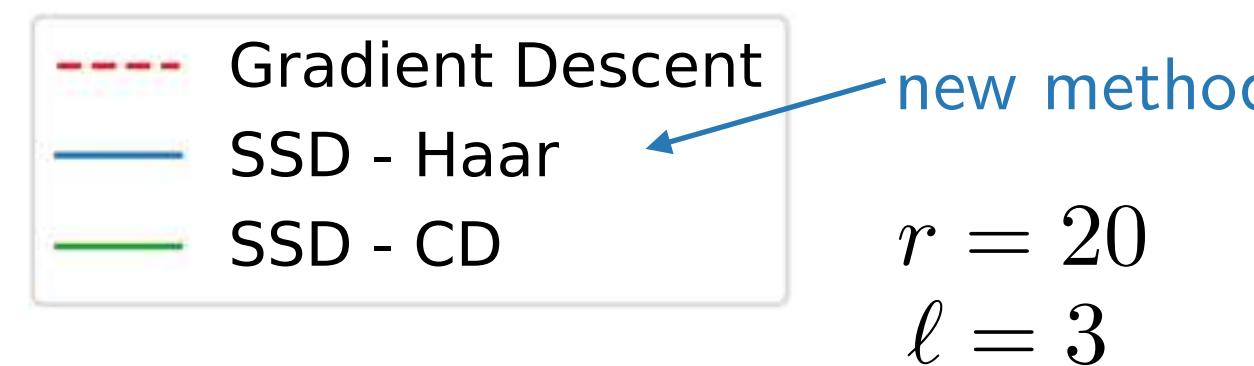
This has **intrinsic dimension** of r



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Numerical Results: better than expected

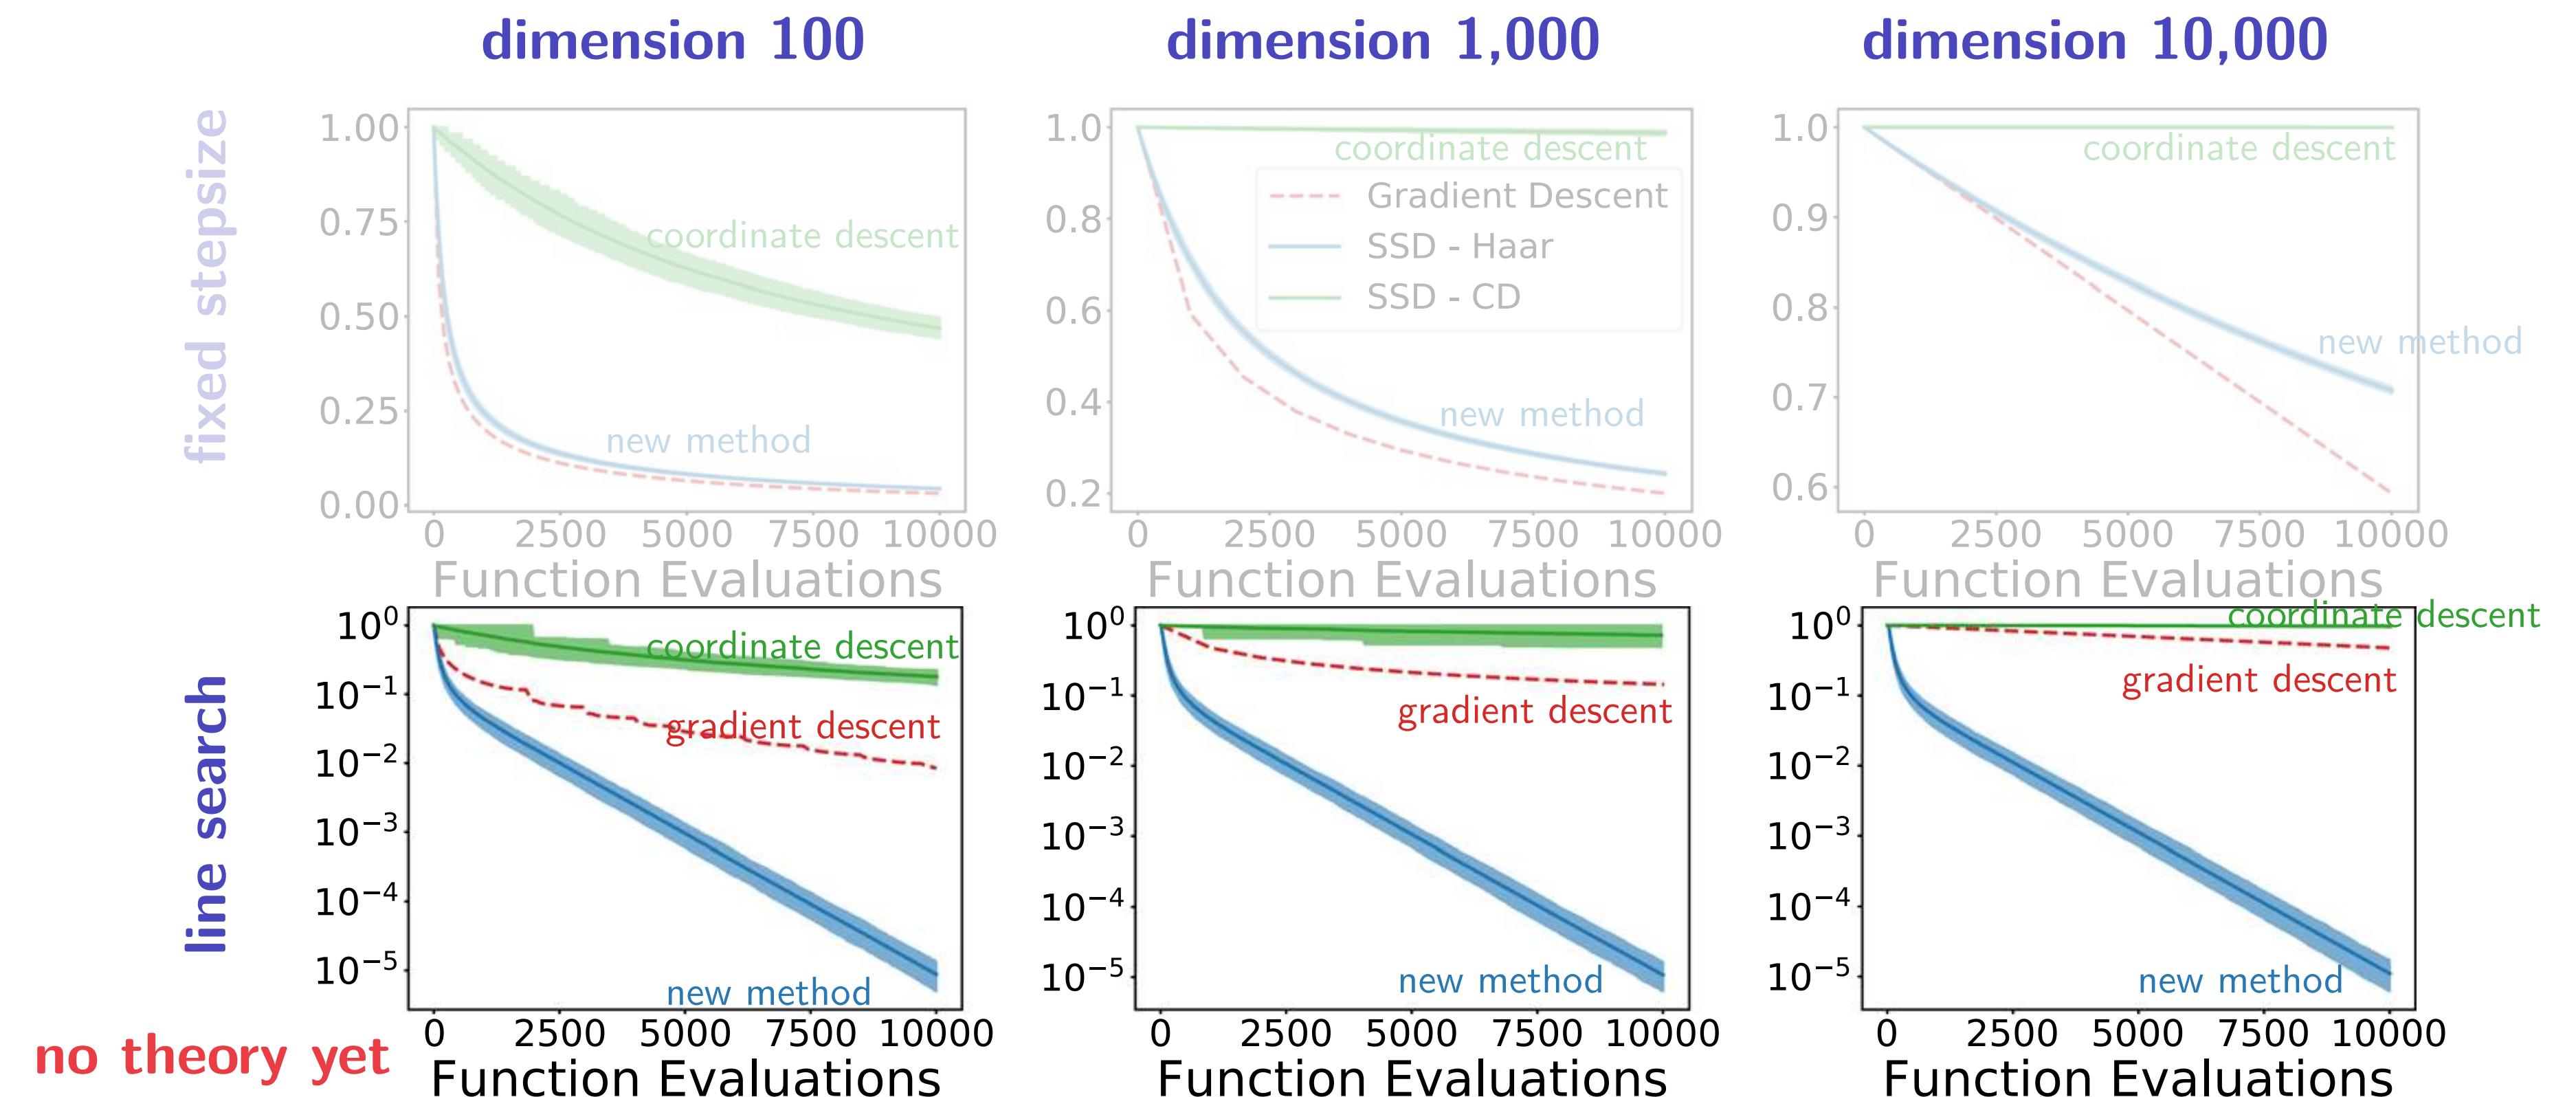
Haar SSD drastically outperforms randomized coordinate descent (CD)



We can force it to happen by making a problem with low "intrinsic" dimension, e.g., Nesterov's "worst function in the world"

$$f_{\lambda,r}(\mathbf{x}) = \lambda((x_1^2 + \sum_{i=1}^{r-1} (x_i - x_{i+1})^2 + x_r^2)/2 - x_1)/4,$$

This has **intrinsic dimension** of r



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Theory: explain better-than-expected results

Previous theorem didn't actually rely on properties of Haar distribution, just generic Q :

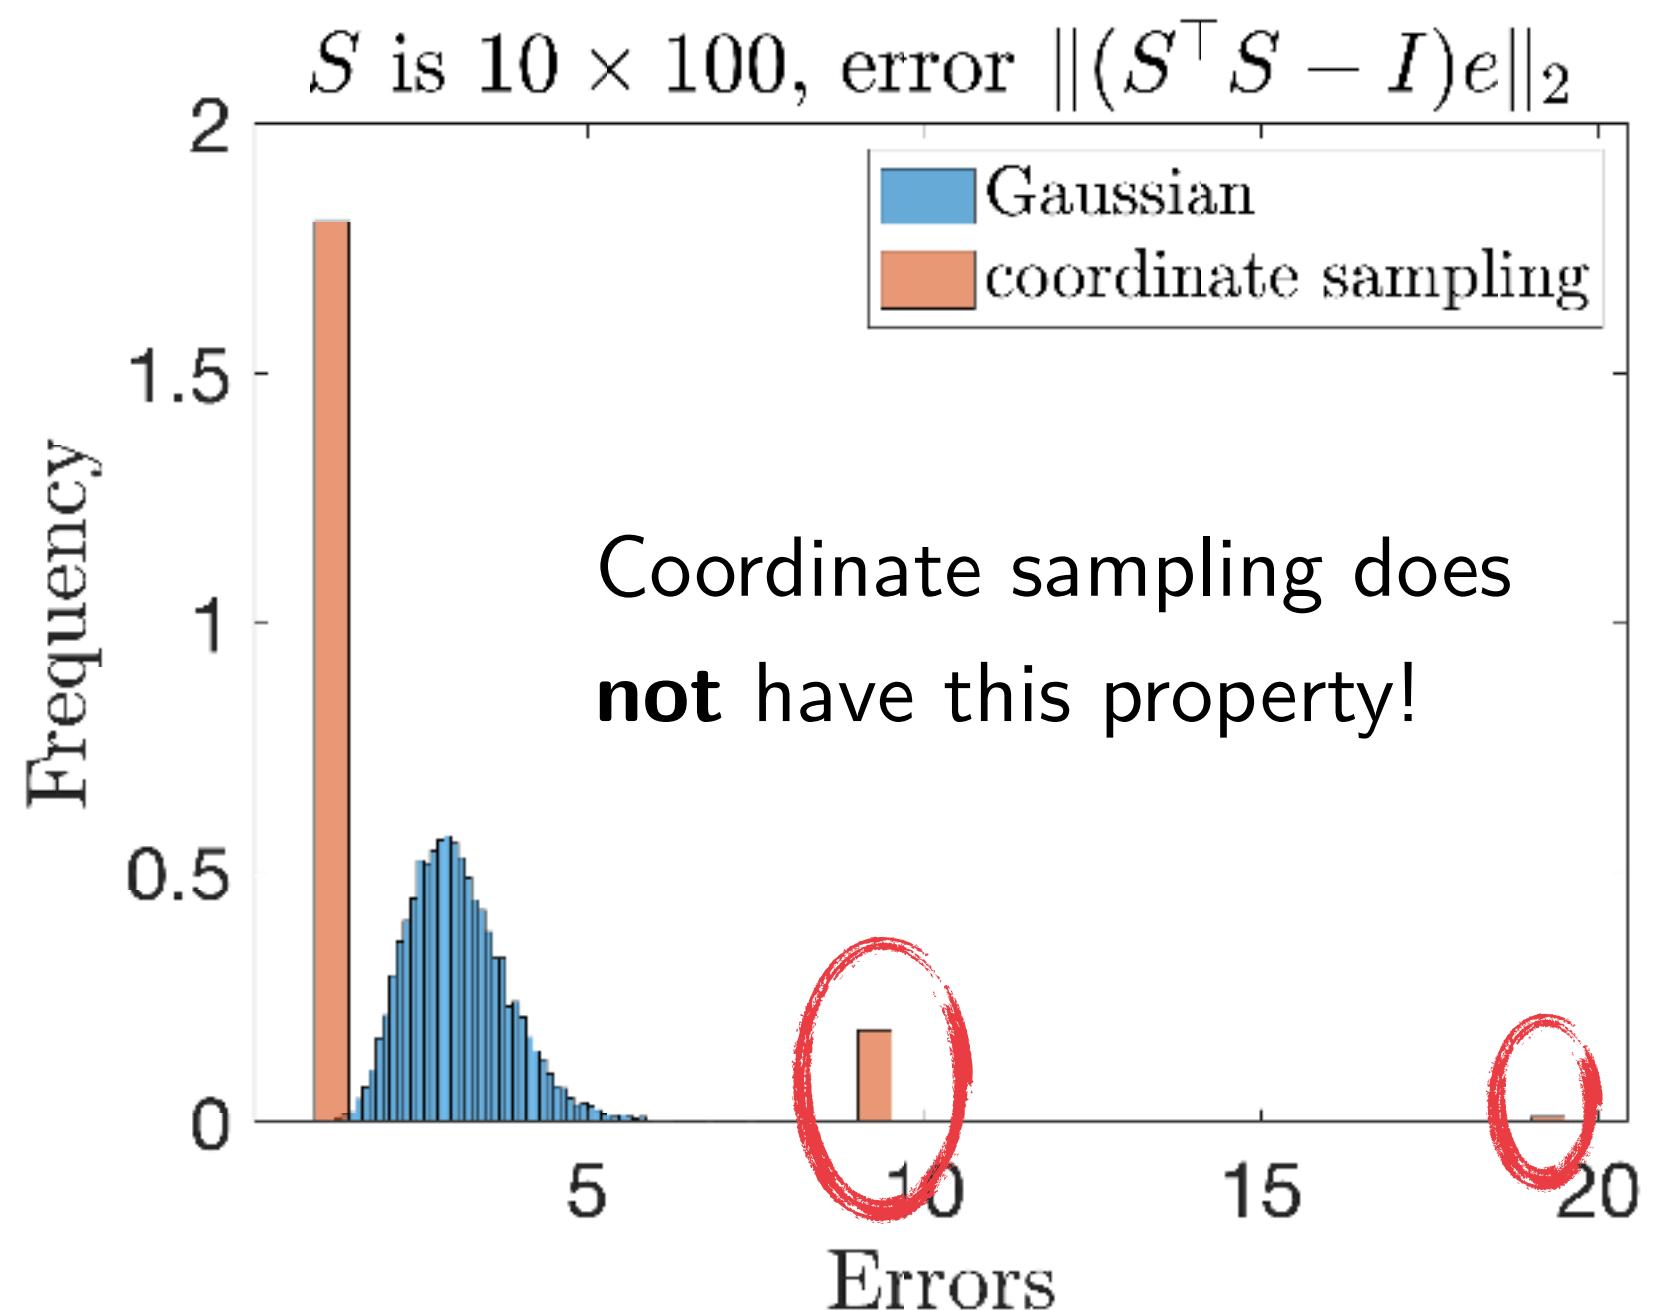
Tighter analysis using concentration-of-measure:

$$Q^T Q = I_{\ell \times \ell}, \quad \mathbb{E} \left(\frac{d}{\ell} Q Q^T \right) = I_{d \times d}$$

Lemma 2 (Johnson-Lindenstrauss style embedding, from Kozak, Becker, Tenorio '19, Lemma 1).

$\forall \epsilon \in (0, 1)$, if $\ell \gtrsim \epsilon^{-2}$, $Q \sim \text{Haar}(d \times \ell)$, then $\forall 0 \neq g \in \mathbb{R}^d$,

$$1 - \epsilon \leq \frac{d}{\ell} \frac{\|Q^T g\|^2}{\|g\|^2} \leq 1 + \epsilon \quad w/ \text{ prob. } \delta \geq 0.8$$



Recall...
 d = ambient dimension
 ℓ = # directional derivs

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Theory: explain better-than-expected results

Previous theorem didn't actually rely on properties of Haar distribution, just

$$Q^T Q = I_{\ell \times \ell}, \quad \mathbb{E} \left(\frac{d}{\ell} Q Q^T \right) = I_{d \times d}$$

Tighter analysis using concentration-of-measure:

Lemma 2 (Johnson-Lindenstrauss style embedding, from Kozak, Becker, Tenorio '19, Lemma 1).
 $\forall \epsilon \in (0, 1)$, if $\ell \gtrsim \epsilon^{-2}$, $Q \sim \text{Haar}(d \times \ell)$, then $\forall 0 \neq g \in \mathbb{R}^d$,

$$1 - \epsilon \leq \frac{d}{\ell} \frac{\|Q^T g\|^2}{\|g\|^2} \leq 1 + \epsilon \quad w/ \text{ prob. } \delta \geq 0.8$$

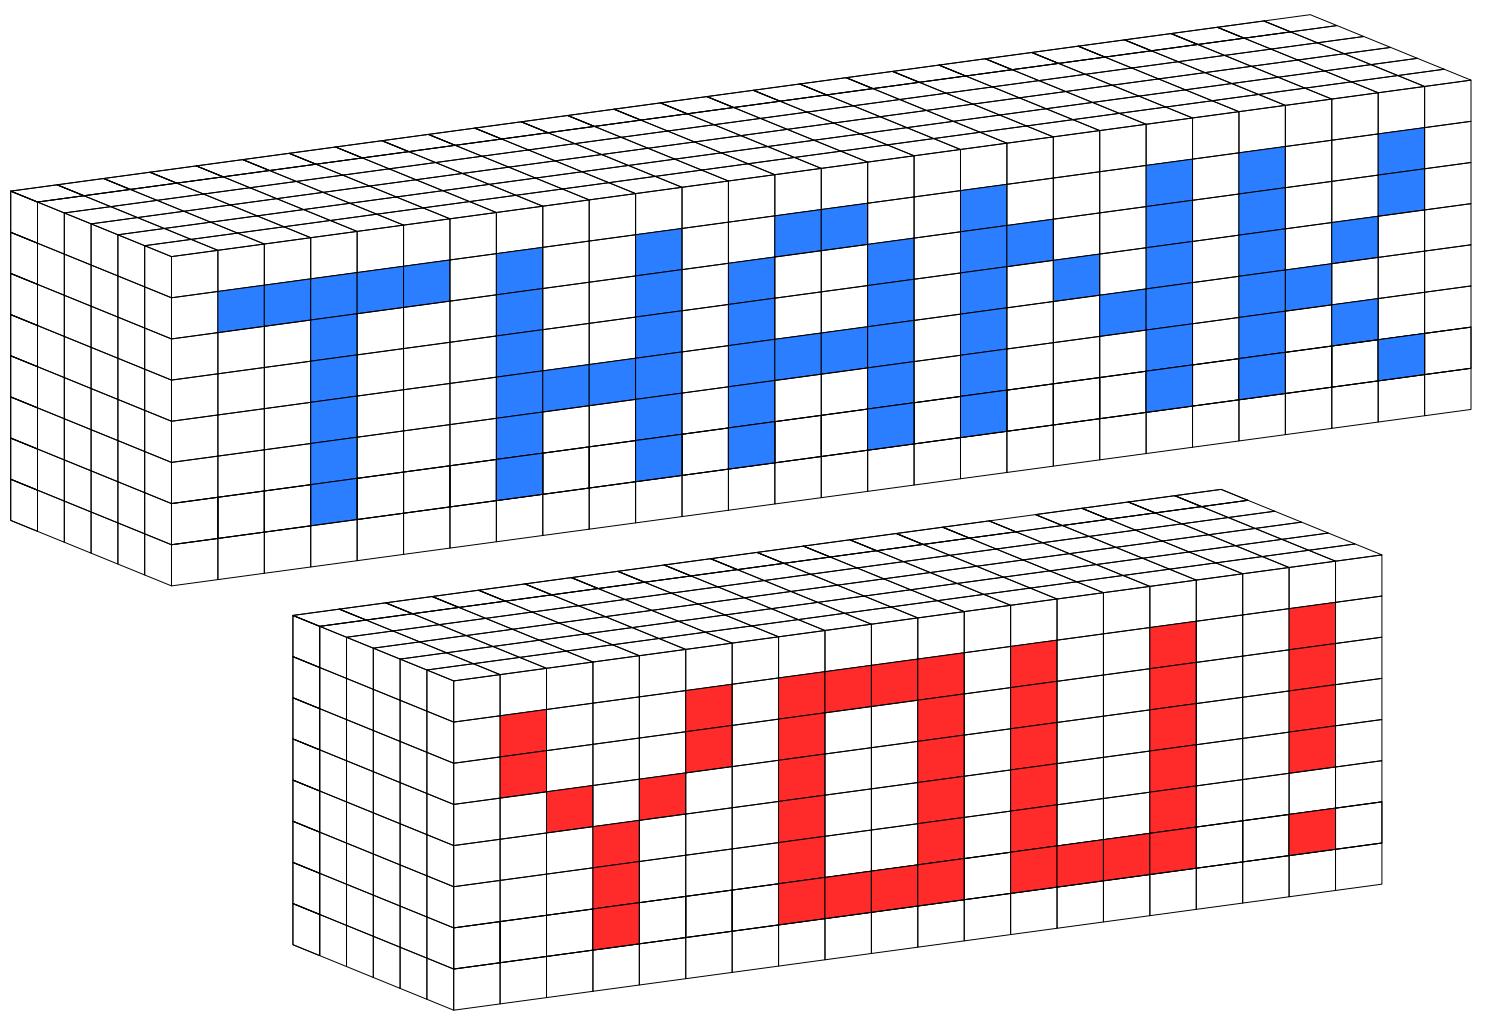
Theorem 3 (Kozak, Becker, Tenorio '19, Thm. 1). If f is strongly convex and ∇f is Lipschitz continuous, then for an appropriate stepsize η_k , the sequence (x_k) generated by SSD (with $Q \sim \text{Haar}$), for $k > 100$, satisfies

$$f(x_k) - f^* \leq (1 + (1 - \epsilon)\rho)^{k/2} (f(x_0) - f^*) \quad \text{with probability } \geq 0.998,$$

where $\rho < 1$ depends on ℓ , d and the Lipschitz and strong convexity parameters.

error in JL embedding

due to possibility of failure of JL



Papers and code available at <https://amath.colorado.edu/faculty/becker/>

1. Randomized "sketches"

a. Warmup: PCA

b. Classical sketches

c. Structured sketches

2. Applications

a. Warmup: linear algebra

b. K-means clustering

c. Tensor factorizations

d. Gradient-free optimization

Method 7: CountSketch

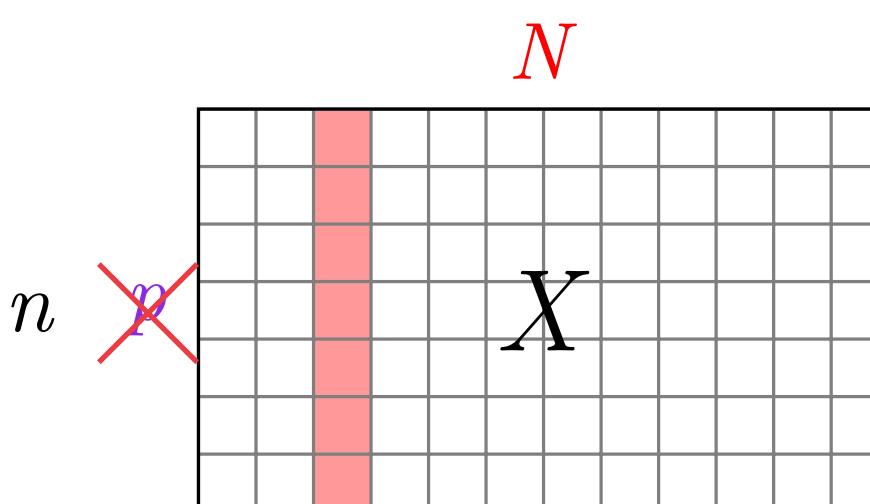
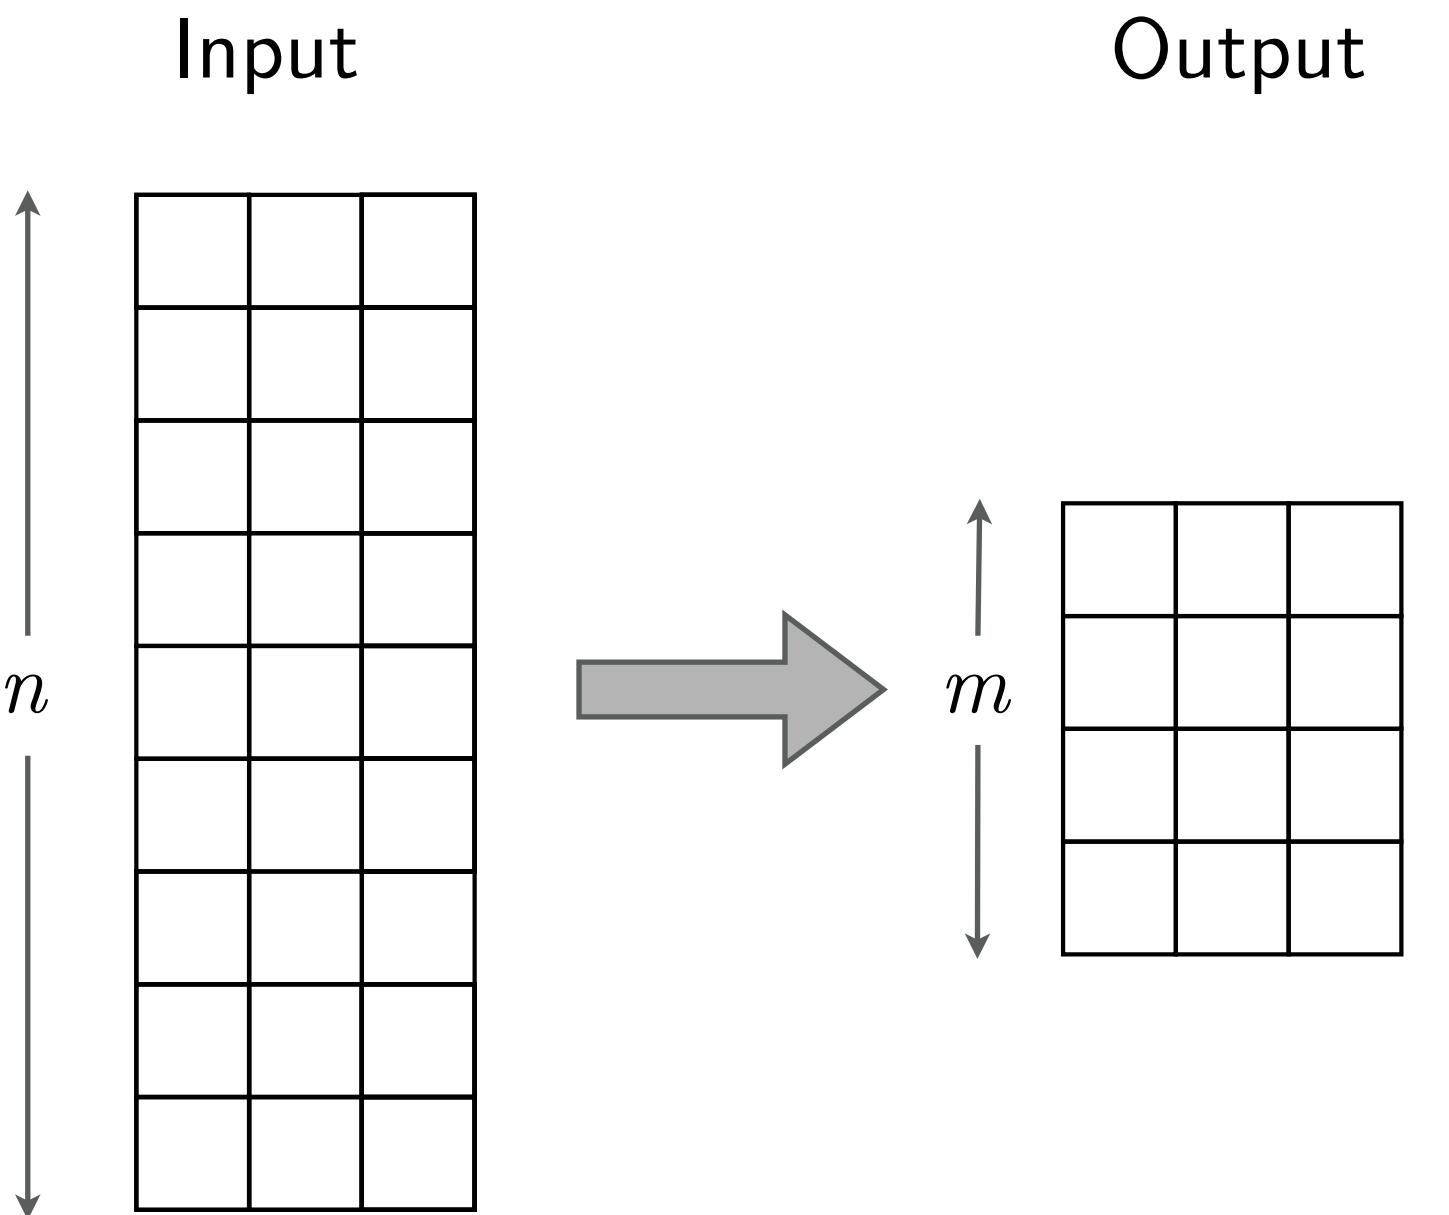
Introduced in [Charikar et al. \(2004\)](#), more analysis
in, e.g., [Clarkson and Woodruff \(2017\)](#)

$$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ linear operator}$$



Change in notation temporarily

$$\begin{aligned}\Phi &\rightarrow \mathcal{S} \\ p &\rightarrow n \\ p_{\text{small}} &\rightarrow m\end{aligned}$$



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ linear operator

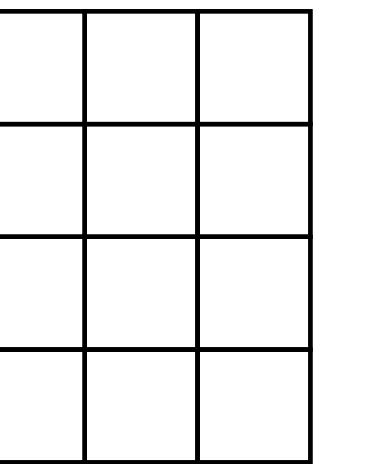
Step 1:

Multiply by random sign*

$s : [n] \rightarrow \{\pm 1\}$
(uniform)

+1	+1	+1
+1	+1	+1
-1	-1	-1
+1	+1	+1
-1	-1	-1
-1	-1	-1
-1	-1	-1
+1	+1	+1
-1	-1	-1

Notation: $[n] = (0, 1, \dots, n - 1)$
(use 0-based indexing)



* technically doesn't have to be fully random, but must be 2-wise independent

$$s(i) \perp s(j) \quad \text{if } i \neq j$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

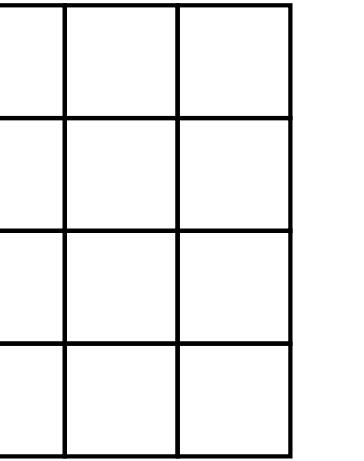
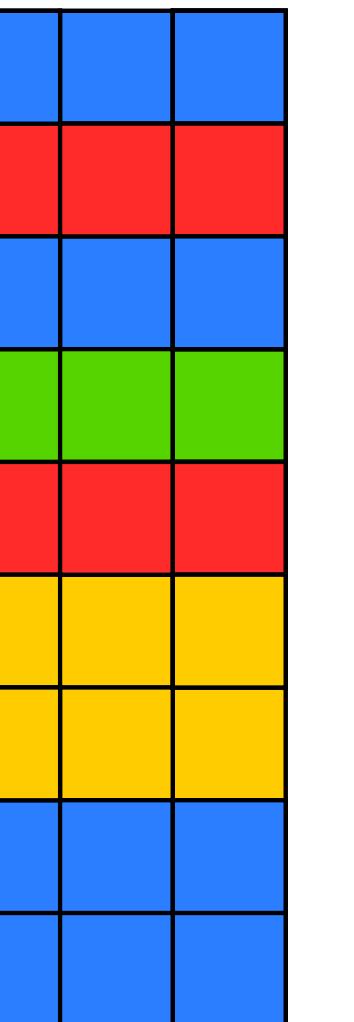
$$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ linear operator}$$

Step 2:

Assign an output row to
every input row (randomly
or with a hash function)

$$h : [n] \rightarrow [m]$$

(uniform)



Again, don't need *all* of $h[i]$ to be independent, only pairwise independent
(but easy enough to make them all independent)

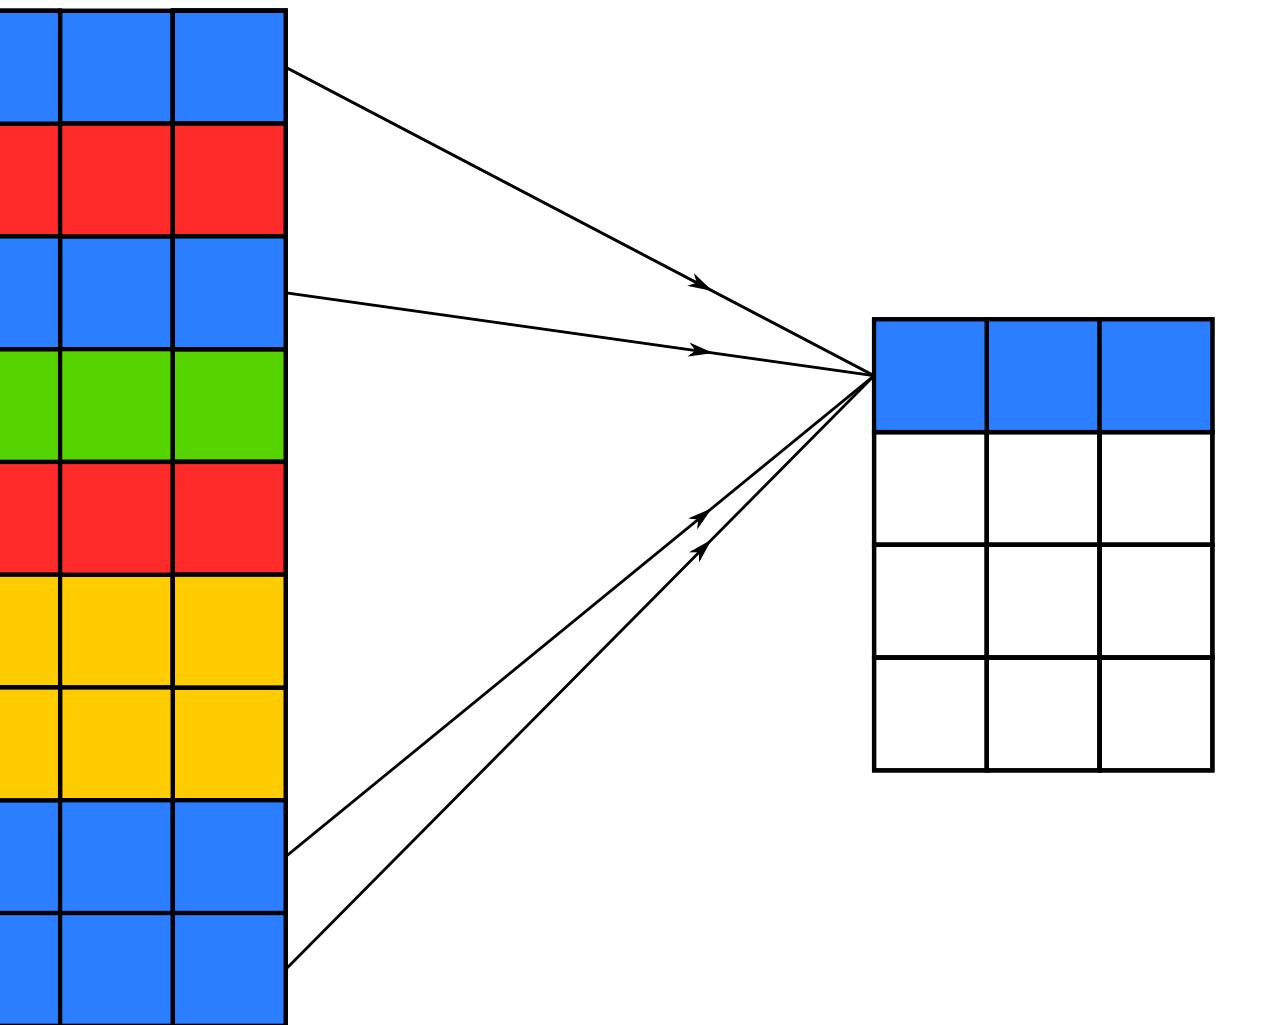
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

$$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ linear operator}$$

Step 3:

Each output row is the sum
of corresponding input rows



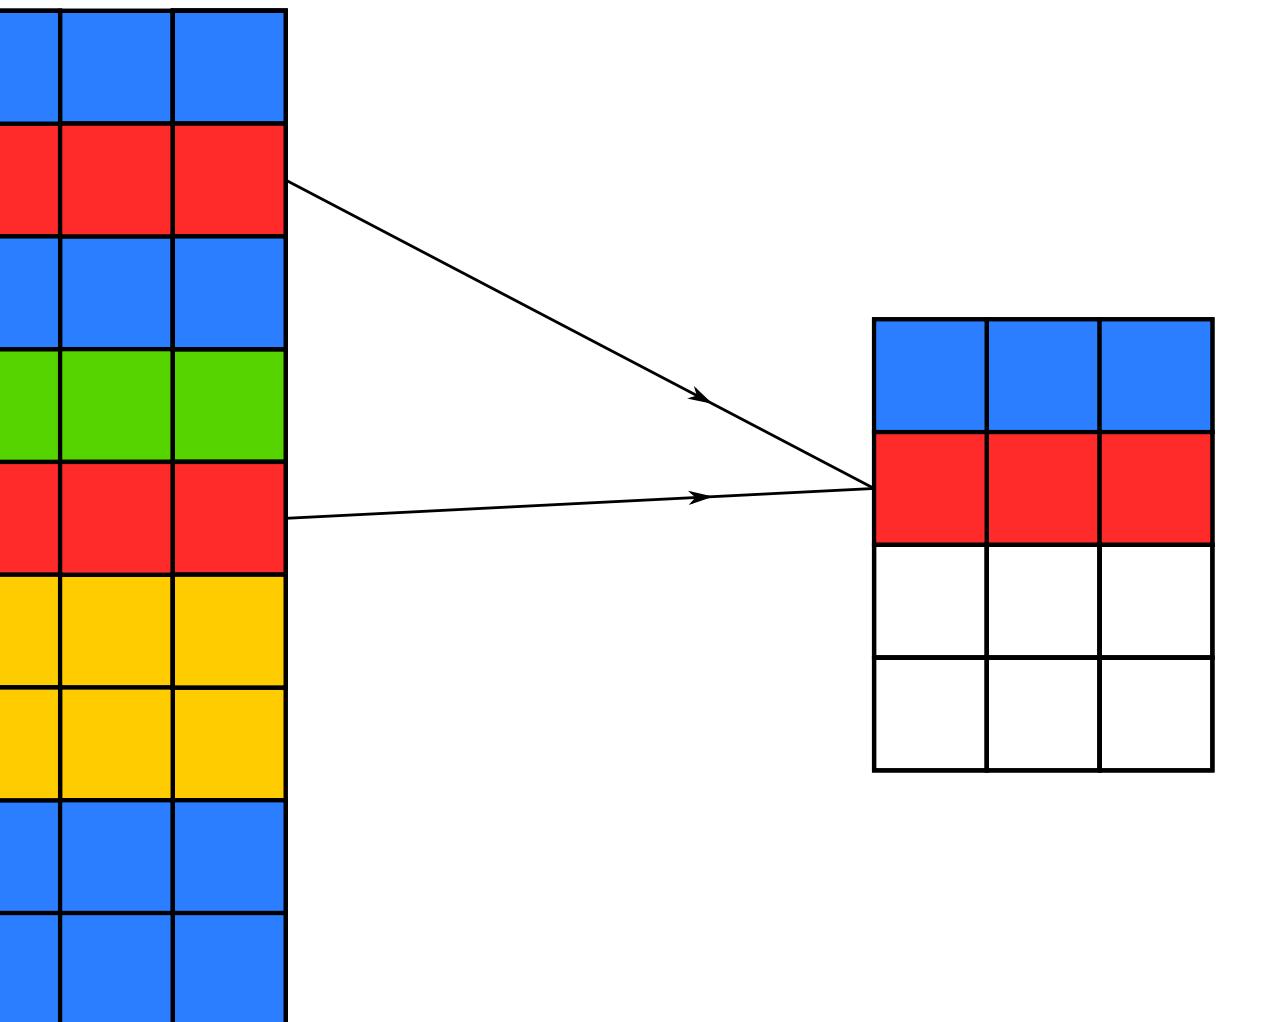
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

$$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ linear operator}$$

Step 3:

Each output row is the sum
of corresponding input rows



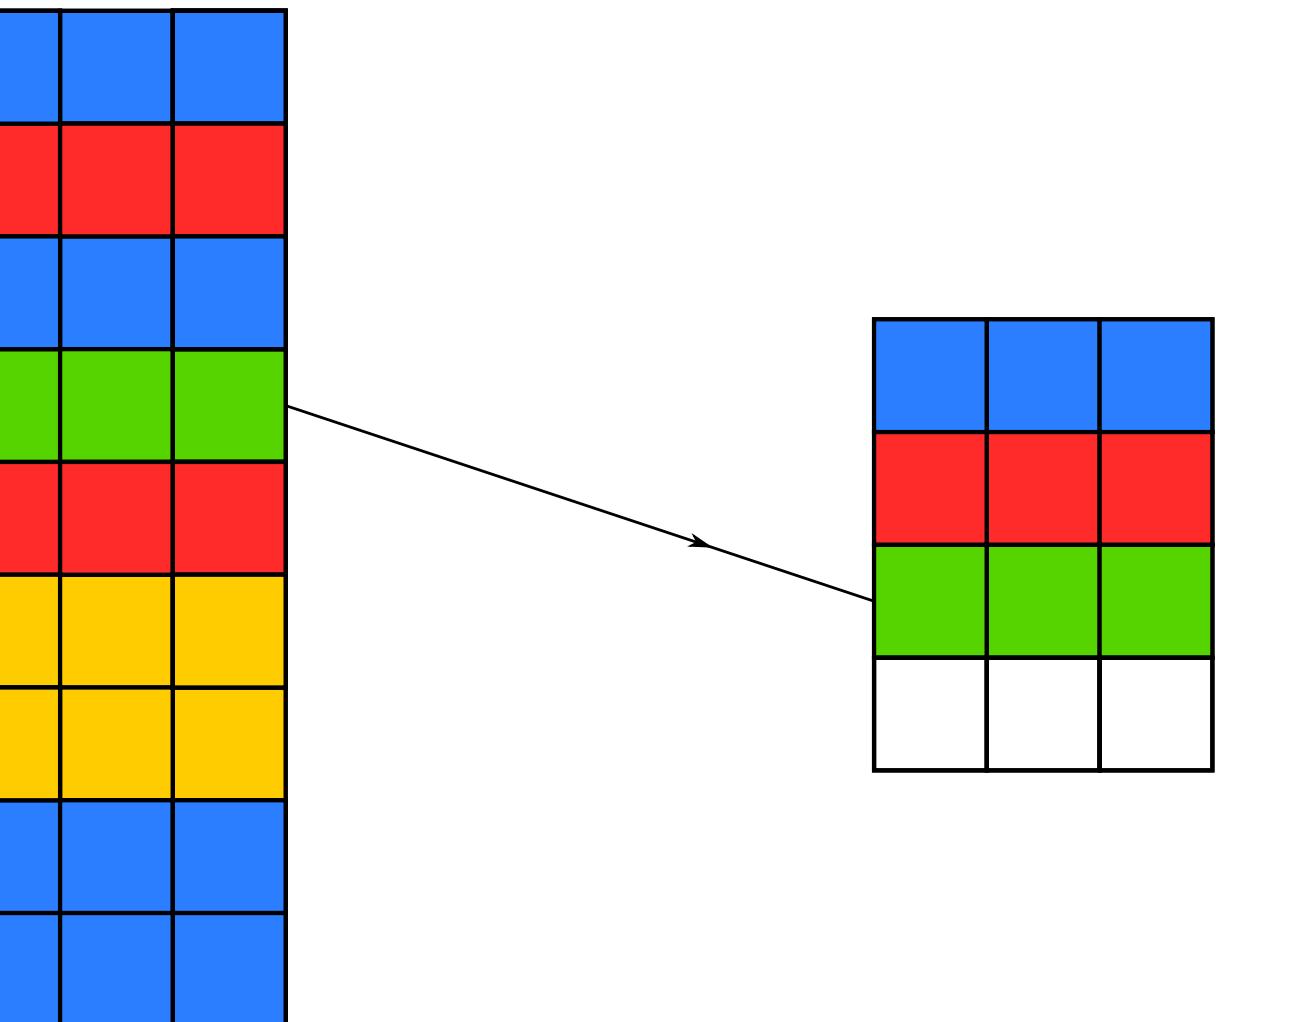
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

$$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ linear operator}$$

Step 3:

Each output row is the sum
of corresponding input rows



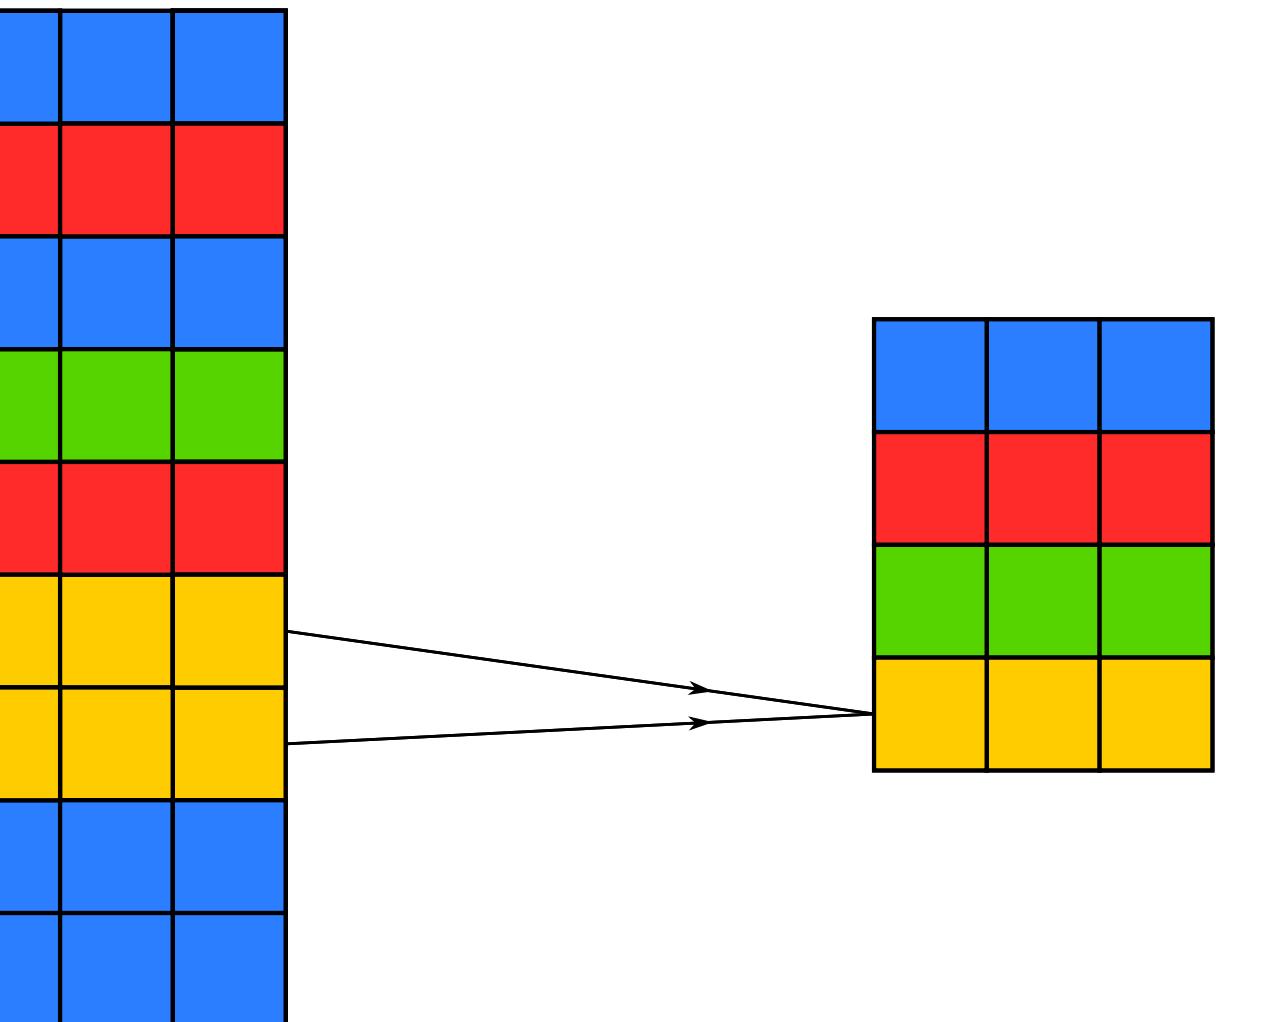
1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

$$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ linear operator}$$

Step 3:

Each output row is the sum
of corresponding input rows



1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

$$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m \text{ linear operator}$$

Complexity analysis: every input element is touched once
... so linear complexity
(and can exploit sparsity)

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

Method 7: CountSketch

$\mathcal{S} : \mathbb{R}^n \rightarrow \mathbb{R}^m$ linear operator

A formula for the output:

$$\mathbf{u} = \mathcal{S}(\mathbf{v}), \quad u_i = \sum_{j|h(j)=i} s(j)v_j$$

$$\begin{aligned} s &: [n] \rightarrow \{\pm 1\} \\ h &: [n] \rightarrow [m] \quad \text{hash} \end{aligned}$$

... and note that the following polynomial has the output as its coefficients:

$$p(\mathbf{x}) \stackrel{\text{def}}{=} \sum_{j \in [n]} s(j) \cdot v_j \cdot \mathbf{x}^{h(j)} = \sum_{i \in [m]} \mathbf{x}^i \underbrace{\left(\sum_{j|h(j)=i} s(j)v_j \right)}_{u_i} = \sum_{i \in [m]} u_i \cdot \mathbf{x}^i$$

(this will be key shortly...)

Method 8: TensorSketch

Introduced in Pagh (2013), more analysis
in, e.g., Diao, Zong, Sun, Woodruff (2018)

TensorSketch is just CountSketch when the input can be written as a tensor product
(for a special choice of the hash and sign functions)



Change in notation temporarily

$$\Phi \rightarrow \mathcal{S}$$

$$p \rightarrow n$$

$$p_{\text{small}} \rightarrow m$$

$$\mathcal{T} : \mathbb{R}^n \rightarrow \mathbb{R}^m$$

$$\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)} \quad \text{where size is} \quad n = n^{(1)} \cdot n^{(2)}$$

Not (yet!) related to tensors

TensorSketch

Introduced in Pagh (2013), more analysis
in, e.g., Diao, Zong, Sun, Woodruff (2018)

TensorSketch is just CountSketch when the input can be written as a tensor product
(for a special choice of the hash and sign functions)

$$\mathcal{T} : \mathbb{R}^n \rightarrow \mathbb{R}^m$$

$$\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)} \quad \text{where size is} \quad n = n^{(1)} \cdot n^{(2)}$$

Kronecker/tensor product

$$\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \\ a_3 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix} \quad \mathbf{a} \otimes \mathbf{b} = \begin{bmatrix} a_1 \mathbf{b} \\ a_2 \mathbf{b} \\ a_3 \mathbf{b} \end{bmatrix} = \begin{bmatrix} a_1 b_1 \\ a_1 b_2 \\ a_1 b_3 \\ a_2 b_1 \\ a_2 b_2 \\ a_2 b_3 \\ a_3 b_1 \\ a_3 b_2 \\ a_3 b_3 \end{bmatrix}$$

TensorSketch

Introduced in Pagh (2013), more analysis
in, e.g., Diao, Zong, Sun, Woodruff (2018)

TensorSketch is just CountSketch when the input can be written as a tensor product
(for a special choice of the hash and sign functions)

$$\mathcal{T} : \mathbb{R}^n \rightarrow \mathbb{R}^m$$

$$\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)} \quad \text{where size is} \quad n = n^{(1)} \cdot n^{(2)}$$

Kronecker/tensor product

$$\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \\ a_3 \end{bmatrix}, \quad \mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

$$\mathbf{a} \otimes \mathbf{b} = \begin{bmatrix} \mathbf{a}_1 \mathbf{b} \\ \mathbf{a}_2 \mathbf{b} \\ \mathbf{a}_3 \mathbf{b} \end{bmatrix} = \begin{bmatrix} a_1 b_1 \\ a_1 b_2 \\ a_1 b_3 \\ a_2 b_1 \\ a_2 b_2 \\ a_2 b_3 \\ a_3 b_1 \\ a_3 b_2 \\ a_3 b_3 \end{bmatrix}$$

another equivalent definition:

$$\mathbf{a} \otimes \mathbf{b} = \text{vec}_{\text{col}} (\mathbf{b} \mathbf{a}^T)$$

$$\mathbf{b} \mathbf{a}^T = \begin{bmatrix} b_1 a_1 & b_1 a_2 & b_1 a_3 \\ b_2 a_1 & b_2 a_2 & b_2 a_3 \\ b_3 a_1 & b_3 a_2 & b_3 a_3 \end{bmatrix} = \begin{bmatrix} a_1 b_1 & a_2 b_1 & a_3 b_1 \\ a_1 b_2 & a_2 b_2 & a_3 b_2 \\ a_1 b_3 & a_2 b_3 & a_3 b_3 \end{bmatrix}$$

$$\begin{aligned} \text{vec} : [n^{(1)}] \times [n^{(2)}] &\rightarrow [n^{(1)} \cdot n^{(2)}] \\ \text{vec}^{-1} = \text{mat} : [n^{(1)} \cdot n^{(2)}] &\rightarrow [n^{(1)}] \times [n^{(2)}] \end{aligned}$$

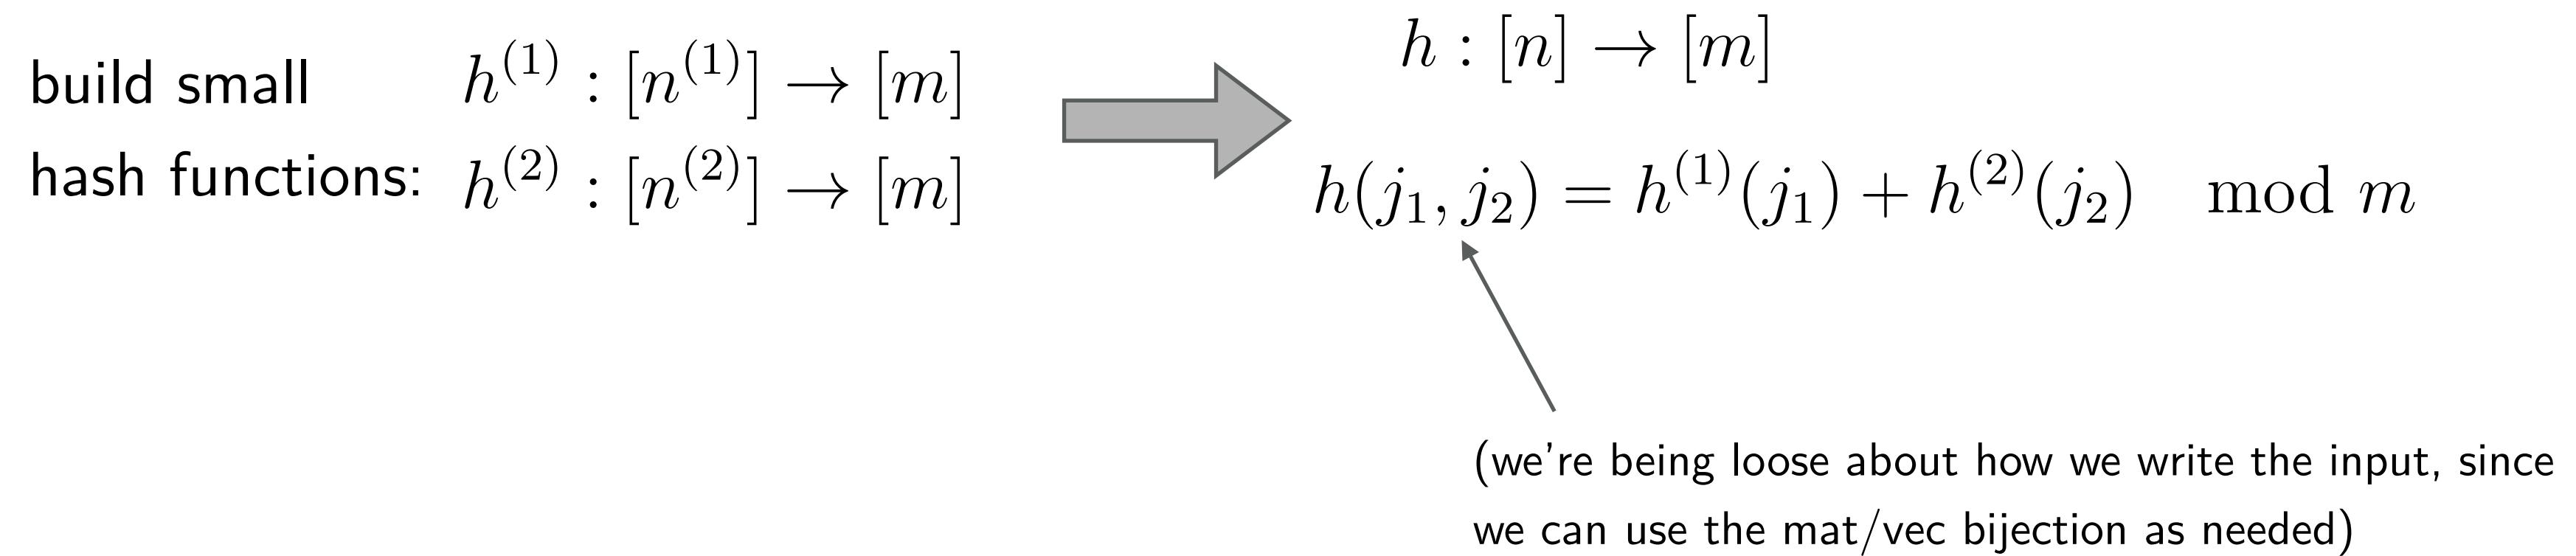
more generally,

$$(\mathbf{A} \otimes \mathbf{B}) \text{vec}_{\text{col}} (\mathbf{X}) = \text{vec}_{\text{col}} (\mathbf{B} \mathbf{X} \mathbf{A}^T)$$

TensorSketch

Pick hash functions in a decomposable way:

Input has structure $\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)}$ with sizes $n = n^{(1)} \cdot n^{(2)}$



Fact: as long as $h^{(i)}$ is fully independent (or at least 3-wise independent), then h is 2-wise independent

TensorSketch

Pick hash functions in a decomposable way:

Input has structure $\mathbf{v} = \mathbf{v}^{(1)} \otimes \mathbf{v}^{(2)}$ with sizes $n = n^{(1)} \cdot n^{(2)}$

build small hash functions: $h^{(1)} : [n^{(1)}] \rightarrow [m]$

$h^{(2)} : [n^{(2)}] \rightarrow [m]$



$h : [n] \rightarrow [m]$

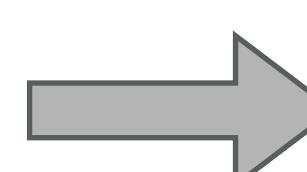
$h(j_1, j_2) = h^{(1)}(j_1) + h^{(2)}(j_2) \pmod m$

Fact: h is 2-wise independent

(we're being loose about how we write the input, since we can use the mat/vec bijection as needed)

same trick for sign functions: $s^{(1)} : [n^{(1)}] \rightarrow \{\pm 1\}$

$s^{(2)} : [n^{(2)}] \rightarrow \{\pm 1\}$



$s : [n] \rightarrow \{\pm 1\}$

$s(j_1, j_2) = s^{(1)}(j_1) \cdot s^{(2)}(j_2)$

still pairwise independent, as needed for theory!

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

what's the point? Huge computational speedup.

$$h^{(1)} : [n^{(1)}] \rightarrow [m]$$

$$h^{(2)} : [n^{(2)}] \rightarrow [m]$$

$$p^{(1)}(\textcolor{red}{x}) \stackrel{\text{def}}{=} \sum_{j \in [n^{(1)}]} s^{(1)}(j) \cdot v_j^{(1)} \cdot \textcolor{red}{x}^{h^{(1)}(j)} = \sum_{i \in [m]} u_i^{(1)} \cdot \textcolor{red}{x}^i$$

$$p^{(2)}(\textcolor{red}{x}) \stackrel{\text{def}}{=} \sum_{j \in [n^{(2)}]} s^{(2)}(j) \cdot v_j^{(2)} \cdot \textcolor{red}{x}^{h^{(2)}(j)} = \sum_{i \in [m]} u_i^{(2)} \cdot \textcolor{red}{x}^i$$

Let's compute this: $p^{(1)}(\textcolor{red}{x}) \cdot p^{(2)}(\textcolor{red}{x}) \pmod{x^m - 1}$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

what's the point? Huge computational speedup.

$$p^{(1)}(\textcolor{red}{x}) \cdot p^{(2)}(\textcolor{red}{x}) = \left(\sum_{j_1 \in [n^{(1)}]} s^{(1)}(j_1) \cdot v_{j_1}^{(1)} \cdot \textcolor{red}{x}^{h^{(1)}(j_1)} \right) \cdot \left(\sum_{j_2 \in [n^{(2)}]} s^{(2)}(j_2) \cdot v_{j_2}^{(2)} \cdot \textcolor{red}{x}^{h^{(2)}(j_2)} \right)$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

what's the point? Huge computational speedup.

$$\begin{aligned}
 p^{(1)}(\textcolor{red}{x}) \cdot p^{(2)}(\textcolor{red}{x}) &= \left(\sum_{j_1 \in [n^{(1)}]} s^{(1)}(j_1) \cdot v_{j_1}^{(1)} \cdot \textcolor{red}{x}^{h^{(1)}(j_1)} \right) \cdot \left(\sum_{j_2 \in [n^{(2)}]} s^{(2)}(j_2) \cdot v_{j_2}^{(2)} \cdot \textcolor{red}{x}^{h^{(2)}(j_2)} \right) \\
 &= \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} s^{(1)}(j_1) \cdot s^{(2)}(j_2) \cdot v_{j_1}^{(1)} \cdot v_{j_2}^{(2)} \cdot \textcolor{red}{x}^{h^{(1)}(j_1)} \cdot \textcolor{red}{x}^{h^{(2)}(j_2)}
 \end{aligned}$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

what's the point? Huge computational speedup.

$$\begin{aligned}
 p^{(1)}(\mathbf{x}) \cdot p^{(2)}(\mathbf{x}) &= \left(\sum_{j_1 \in [n^{(1)}]} s^{(1)}(j_1) \cdot v_{j_1}^{(1)} \cdot \mathbf{x}^{h^{(1)}(j_1)} \right) \cdot \left(\sum_{j_2 \in [n^{(2)}]} s^{(2)}(j_2) \cdot v_{j_2}^{(2)} \cdot \mathbf{x}^{h^{(2)}(j_2)} \right) \\
 &= \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} s^{(1)}(j_1) \cdot s^{(2)}(j_2) \cdot v_{j_1}^{(1)} \cdot v_{j_2}^{(2)} \cdot \mathbf{x}^{h^{(1)}(j_1)} \cdot \mathbf{x}^{h^{(2)}(j_2)} \\
 &= \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} \mathcal{S}(j_1, j_2) \cdot v_{j_1 j_2} \cdot \mathbf{x}^{h^{(1)}(j_1) h^{(2)}(j_2)}
 \end{aligned}$$

because of how we defined our sketch

because input has tensor product structure

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

what's the point? Huge computational speedup.

$$\begin{aligned}
 p^{(1)}(\textcolor{red}{x}) \cdot p^{(2)}(\textcolor{red}{x}) &= \left(\sum_{j_1 \in [n^{(1)}]} s^{(1)}(j_1) \cdot v_{j_1}^{(1)} \cdot \textcolor{red}{x}^{h^{(1)}(j_1)} \right) \cdot \left(\sum_{j_2 \in [n^{(2)}]} s^{(2)}(j_2) \cdot v_{j_2}^{(2)} \cdot \textcolor{red}{x}^{h^{(2)}(j_2)} \right) \\
 &= \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} s^{(1)}(j_1) \cdot s^{(2)}(j_2) \cdot v_{j_1}^{(1)} \cdot v_{j_2}^{(2)} \cdot \textcolor{red}{x}^{h^{(1)}(j_1)} \cdot \textcolor{red}{x}^{h^{(2)}(j_2)} \\
 &= \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} \mathcal{S}(j_1, j_2) \cdot v_{j_1 j_2} \cdot \textcolor{red}{x}^{h^{(1)}(j_1)h^{(2)}(j_2)} \\
 &\equiv \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} \mathcal{S}(j_1, j_2) \cdot v_{j_1 j_2} \cdot \textcolor{red}{x}^{h(j_1, j_2)} \mod \textcolor{red}{x}^m - 1
 \end{aligned}$$



$$\begin{aligned}
 \textcolor{red}{x}^{k\ell} &= \textcolor{red}{x}^{dm + (k\ell \mod m)} = \textcolor{red}{x}^{dm + (k\ell \mod m)} = (\textcolor{red}{x}^m)^d \textcolor{red}{x}^{k\ell \mod m} \equiv \textcolor{red}{x}^{k\ell \mod m} \mod \textcolor{red}{x}^m - 1 \\
 &\qquad\qquad\qquad \text{Recall } \textcolor{red}{x}^m \equiv 1 \mod \textcolor{red}{x}^m - 1
 \end{aligned}$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

what's the point? Huge computational speedup.

$$\begin{aligned}
 p^{(1)}(\mathbf{x}) \cdot p^{(2)}(\mathbf{x}) &= \left(\sum_{j_1 \in [n^{(1)}]} s^{(1)}(j_1) \cdot v_{j_1}^{(1)} \cdot \mathbf{x}^{h^{(1)}(j_1)} \right) \cdot \left(\sum_{j_2 \in [n^{(2)}]} s^{(2)}(j_2) \cdot v_{j_2}^{(2)} \cdot \mathbf{x}^{h^{(2)}(j_2)} \right) \\
 &= \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} s^{(1)}(j_1) \cdot s^{(2)}(j_2) \cdot v_{j_1}^{(1)} \cdot v_{j_2}^{(2)} \cdot \mathbf{x}^{h^{(1)}(j_1)} \cdot \mathbf{x}^{h^{(2)}(j_2)} \\
 &= \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} \mathcal{S}(j_1, j_2) \cdot v_{j_1 j_2} \cdot \mathbf{x}^{h^{(1)}(j_1) h^{(2)}(j_2)} \\
 &\equiv \sum_{j_1 \in [n^{(1)}], j_2 \in [n^{(2)}]} \mathcal{S}(j_1, j_2) \cdot v_{j_1 j_2} \cdot \mathbf{x}^{h(j_1, j_2)} \mod \mathbf{x}^m - 1 \\
 &= p(\mathbf{x})
 \end{aligned}$$

coefficients are the output
of the CountSketch!



1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

How to multiply polynomials?

$$(2\cancel{x}^2 + 3\cancel{x} + 4) \cdot (\cancel{x}^2 - \cancel{x} + 5)$$

$p(x) \qquad q(x)$

$$p(x) \quad \begin{matrix} \cancel{x}^0 & \cancel{x}^1 & \cancel{x}^2 \\ \boxed{4} & 3 & 2 \end{matrix} \quad \begin{matrix} 1 & -1 & 5 \\ \cancel{x}^2 & \cancel{x}^1 & \cancel{x}^0 \end{matrix} \quad q(x)$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

How to multiply polynomials?

$$(2x^2 + 3x + 4) \cdot (x^2 - x + 5)$$

$$\begin{array}{ccc}
 x^0 & x^1 & x^2 \\
 \boxed{4} & 3 & \boxed{2} \\
 & \boxed{1} & -1 & 5 \\
 x^2 & x^1 & x^0
 \end{array}
 = 2x^4$$

$$2 \cdot 1$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

How to multiply polynomials?

$$(2x^2 + 3x + 4) \cdot (x^2 - x + 5)$$

$$\begin{array}{ccc}
 x^0 & x^1 & x^2 \\
 \boxed{4} & \boxed{3} & \boxed{2} \\
 \hline
 & \boxed{1} & \boxed{-1} \\
 x^2 & x^1 & x^0
 \end{array}
 = 2x^4 + x^3$$

$$3 \cdot 1 + 2 \cdot (-1) = 1$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

How to multiply polynomials?

$$(2x^2 + 3x + 4) \cdot (x^2 - x + 5)$$

$$\begin{array}{ccc}
 x^0 & x^1 & x^2 \\
 \boxed{\begin{array}{ccc} 4 & 3 & 2 \\ 1 & -1 & 5 \end{array}} \\
 x^2 & x^1 & x^0
 \end{array}
 = 2x^4 + x^3 + 11x^2$$

$$4 \cdot 1 + 3 \cdot (-1) + 2 \cdot 5 = 11$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

How to multiply polynomials?

$$(2x^2 + 3x + 4) \cdot (x^2 - x + 5)$$

$$\begin{array}{ccc}
 x^0 & x^1 & x^2 \\
 \boxed{\begin{array}{cc} 4 & 3 \\ 1 & -1 \end{array}} & 2 \\
 \hline
 1 & -1 & 5 \\
 x^2 & x^1 & x^0
 \end{array}
 = 2x^4 + x^3 + 11x^2 + 11x$$

$$4 \cdot (-1) + 3 \cdot 5 = 11$$

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

How to multiply polynomials?

$$(2x^2 + 3x + 4) \cdot (x^2 - x + 5)$$

$$\begin{array}{ccc}
 & x^0 & x^1 & x^2 \\
 & | & | & | \\
 \begin{matrix} 4 \\ 3 \\ 2 \end{matrix} & \boxed{4} & 3 & 2 \\
 \hline
 1 & -1 & \boxed{5} \\
 \hline
 x^2 & x^1 & x^0
 \end{array}
 = 2x^4 + x^3 + 11x^2 + 11x + 20$$

$$4 \cdot 5 = 20$$

1. Randomized “sketches”
 - a. Warmup: PCA
 - b. Classical sketches
 - c. **Structured sketches**
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch

How to multiply polynomials?

$$(2x^2 + 3x + 4) \cdot (x^2 - x + 5)$$

x^0	x^1	x^2			
4	3	2			
1	-1	5			
x^2	x^1	x^0			

$= 2x^4 + x^3 + 11x^2 + 11x + 20$
 $\equiv 11x^2 + 13x + 21 \pmod{x^3 - 1}$

$$4 \cdot 5 = 20$$

... this is circular convolution!

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

TensorSketch: complexity analysis

Convolution Theorem (for circular convolutions $*$)

pointwise multiplication

$$\mathcal{F}(g * h) = \mathcal{F}(g)\mathcal{F}(h) \quad \text{i.e.} \quad g * h = \mathcal{F}^{-1}(\mathcal{F}(g)\mathcal{F}(h))$$

Discrete Fourier Transform (implemented via FFT)

1. Randomized "sketches"
 - a. Warmup: PCA
 - b. Classical sketches
 - c. Structured sketches
2. Applications
 - a. Warmup: linear algebra
 - b. K-means clustering
 - c. Tensor factorizations
 - d. Gradient-free optimization

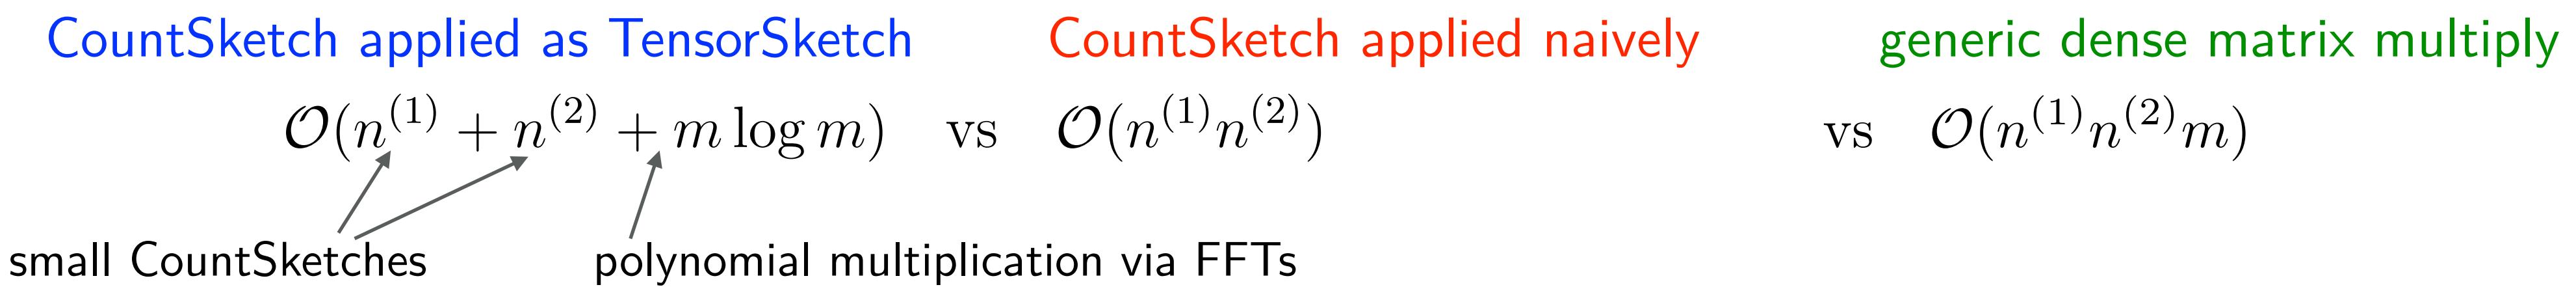
TensorSketch: complexity analysis

Convolution Theorem (for circular convolutions $*$)

$$\mathcal{F}(g * h) = \mathcal{F}(g)\mathcal{F}(h) \quad \text{i.e.} \quad g * h = \mathcal{F}^{-1}(\mathcal{F}(g)\mathcal{F}(h))$$

Complexity:

$$\begin{aligned}\mathcal{T} : \mathbb{R}^n &\rightarrow \mathbb{R}^m \\ n &= n^{(1)} \cdot n^{(2)}\end{aligned}$$



Savings grow as we have more tensor products

If $n = n^{(1)}n^{(2)} \dots n^{(q)}$

$$\mathcal{O}(n^{(1)} + n^{(2)} + \dots + n^{(q)} + m \log m) \quad \text{vs} \quad \mathcal{O}(n^{(1)}n^{(2)} \dots n^{(q)})$$